


Lecture 2:

High-Level Synthesis

Sitao Huang, sitaoh@uci.edu

October 11, 2023



AMD Xilinx FPGA Design Tools

Vitis Unified Software Platform

- **Option 1:** UCI Engineering Remote Labs:
 - Remote Windows machines
 - <https://laptops.eng.uci.edu/computer-labs/remote-labs>
- **Option 2:** UCI EECS Instructional Servers
 - Servers: laguna.eecs.uci.edu, bondi.eecs.uci.edu, crystalcove.eecs.uci.edu
 - Setup X11 display server (Xming for Windows and XQuartz for macOS)
 - SSH to any of these servers with your UCI credentials
 - `ssh -X yourUCInetID@laguna.uci.edu`
 - Xilinx tools installed under: `/ecelib/eceware/xilinx_2022.2/`
 - Initialization: `source /ecelib/eceware/xilinx_2022.2/Vitis/2022.2/settings64.csh`
 - Start Vitis HLS: `vitis_hls &`
- **Option 3:** Install Vitis tools on your own computer (>100 GB)
 - <https://www.xilinx.com/support/download.html>

High-Level Synthesis

- **High-level synthesis (HLS)**, also referred to as C synthesis, electronic system-level (ESL) synthesis, algorithmic synthesis, or behavior synthesis
- HLS takes an abstract behavioral specification of a digital system, and finds a register-transfer level structure that realizes the given behavior

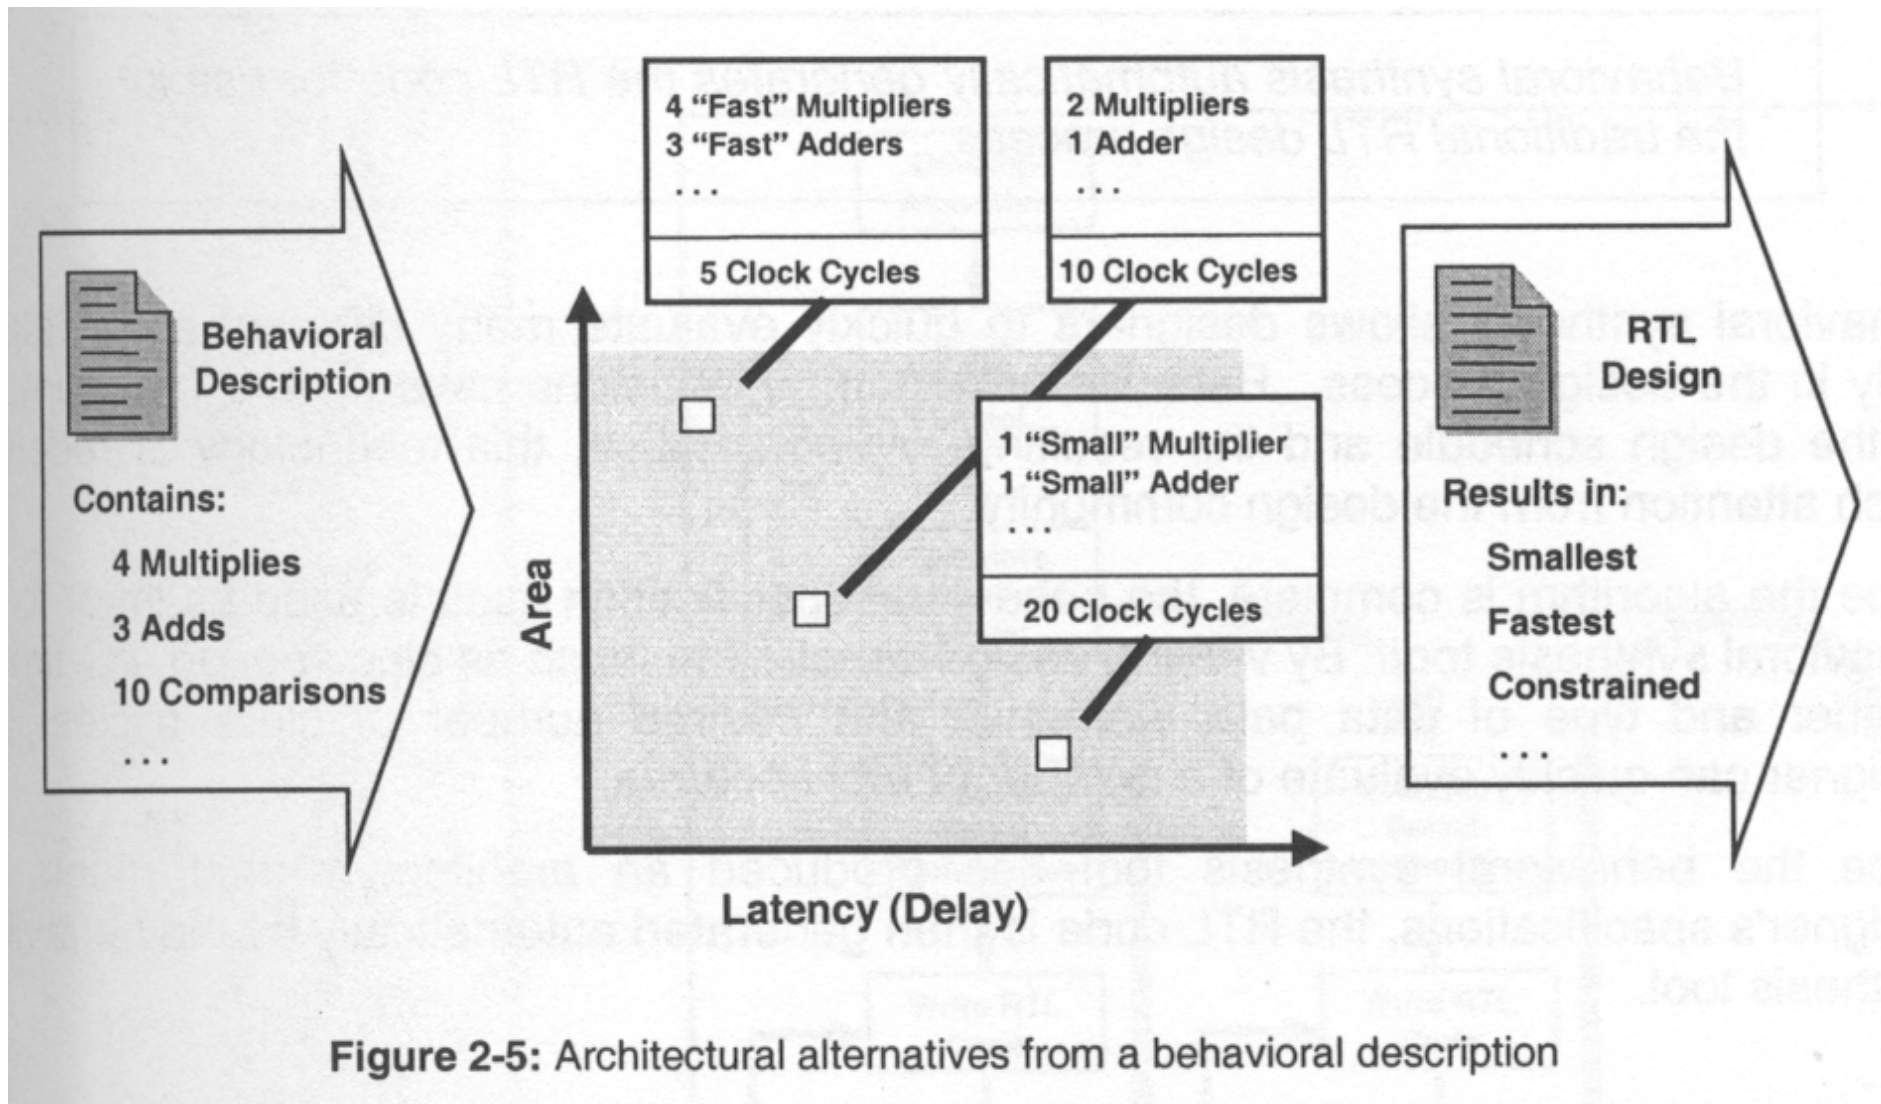
INPUT:

- A high-level, algorithmic description
 - Control structures (if/else, loop, subroutines)
 - Concurrent and sequential semantics
 - Abstract data types
 - Logical and arithmetic operators
- A set of constraints
 - Speed, power, area, interconnect style
 - A library of pre-specified components


OUTPUT:

- A register-transfer level description for further synthesis and optimization

High-Level Synthesis




Design Space Exploration



Add the Power tradeoff and this becomes a 3-dimensional graph

Design Space Exploration: RTL vs. Behavioral Synthesis



High-Level Synthesis Process

Resource Allocation:

- Allocating resources (library components) to each of the operations, buses, muxes, and registers for storage

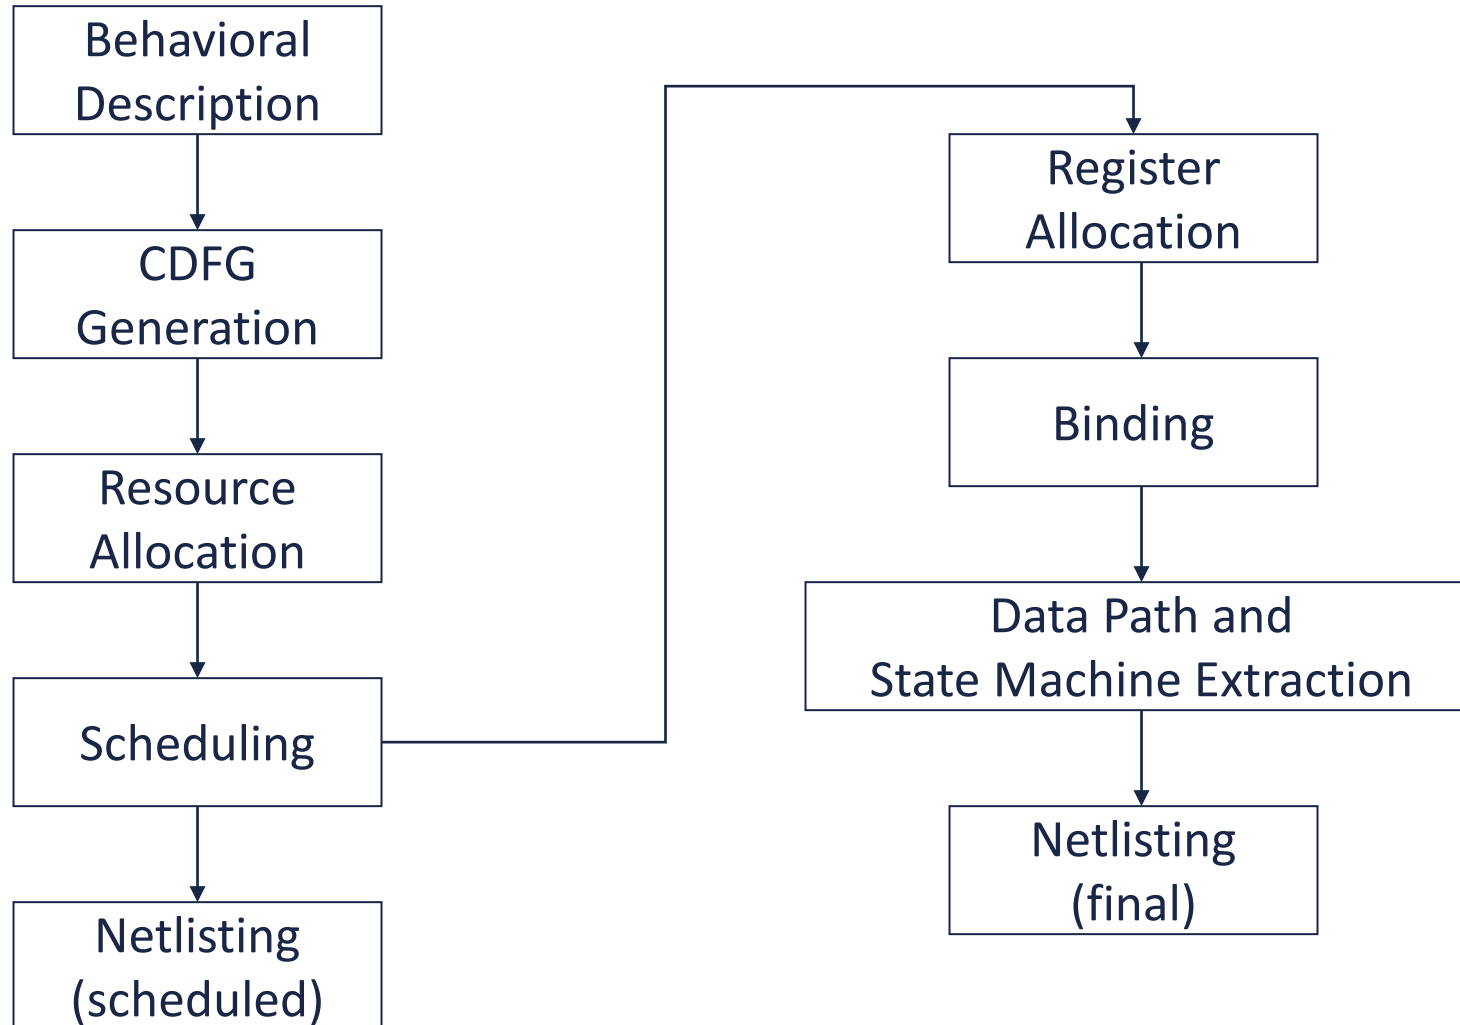
Scheduling:

- Scheduling the operations in the CDFG to minimize area, time and/or power

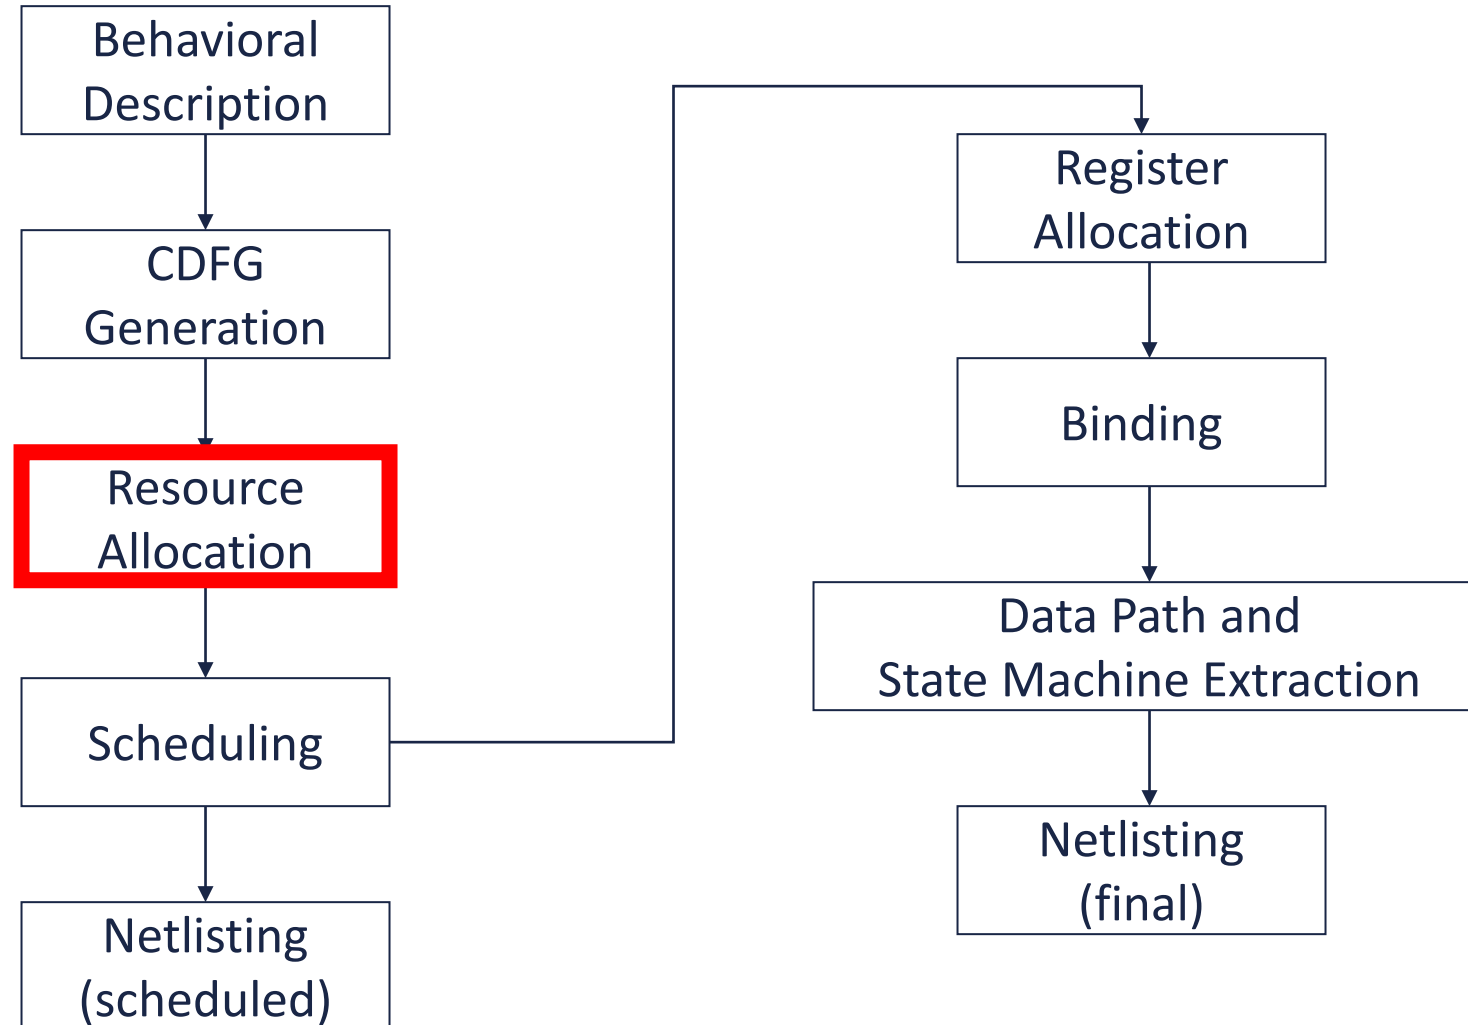
Binding:

- Determining the time of use of each component
- e.g., which registers to use and when

High-Level Synthesis Steps



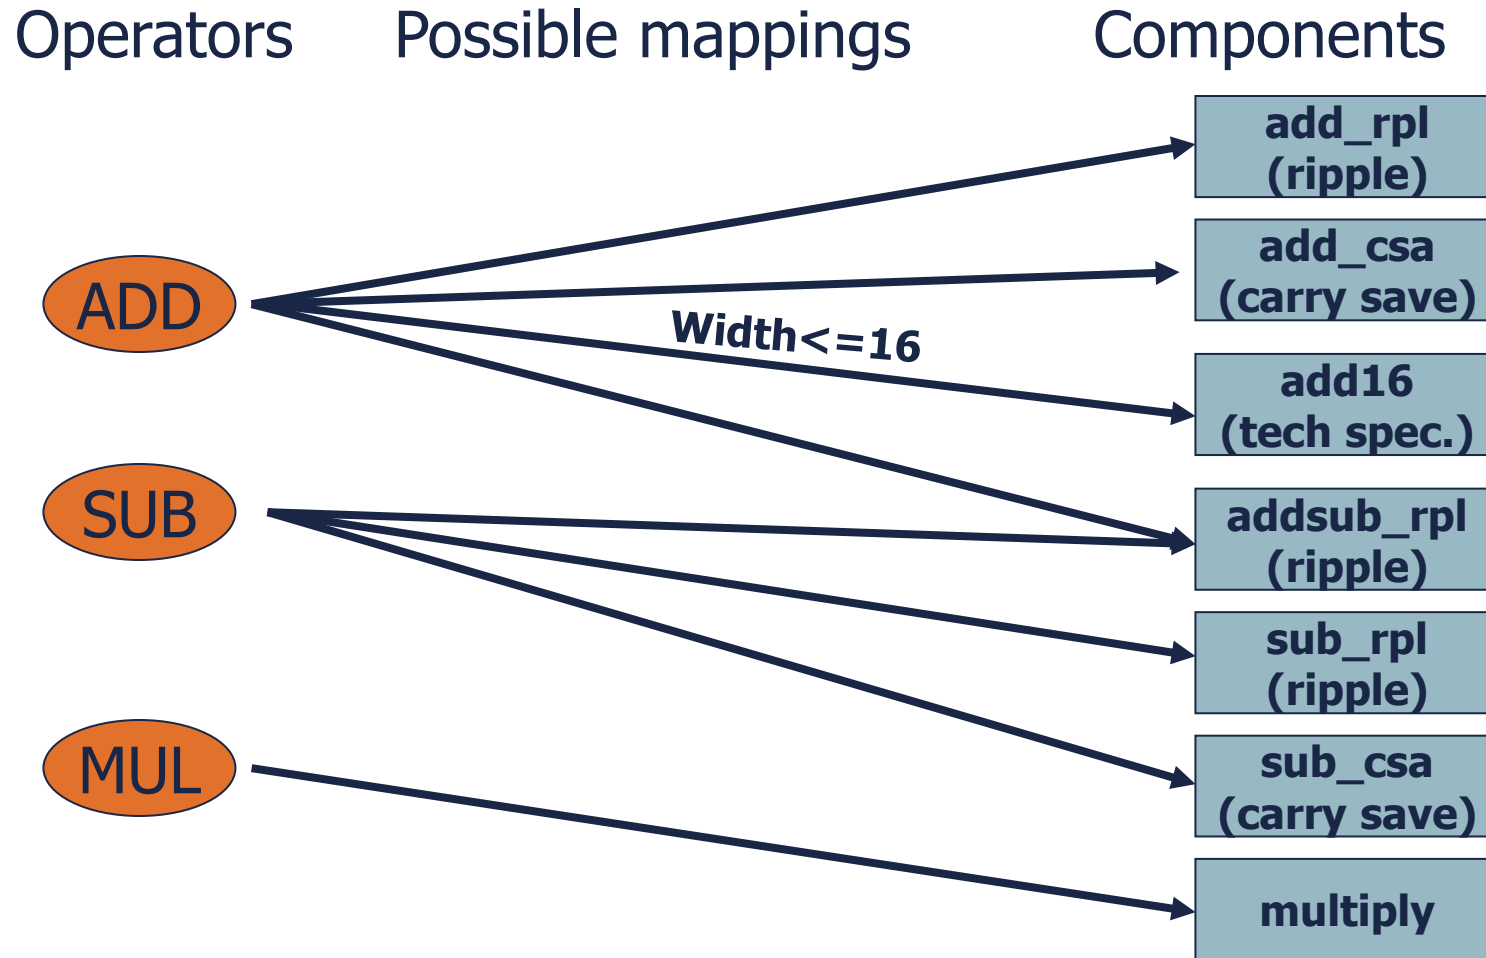
High-Level Synthesis Steps



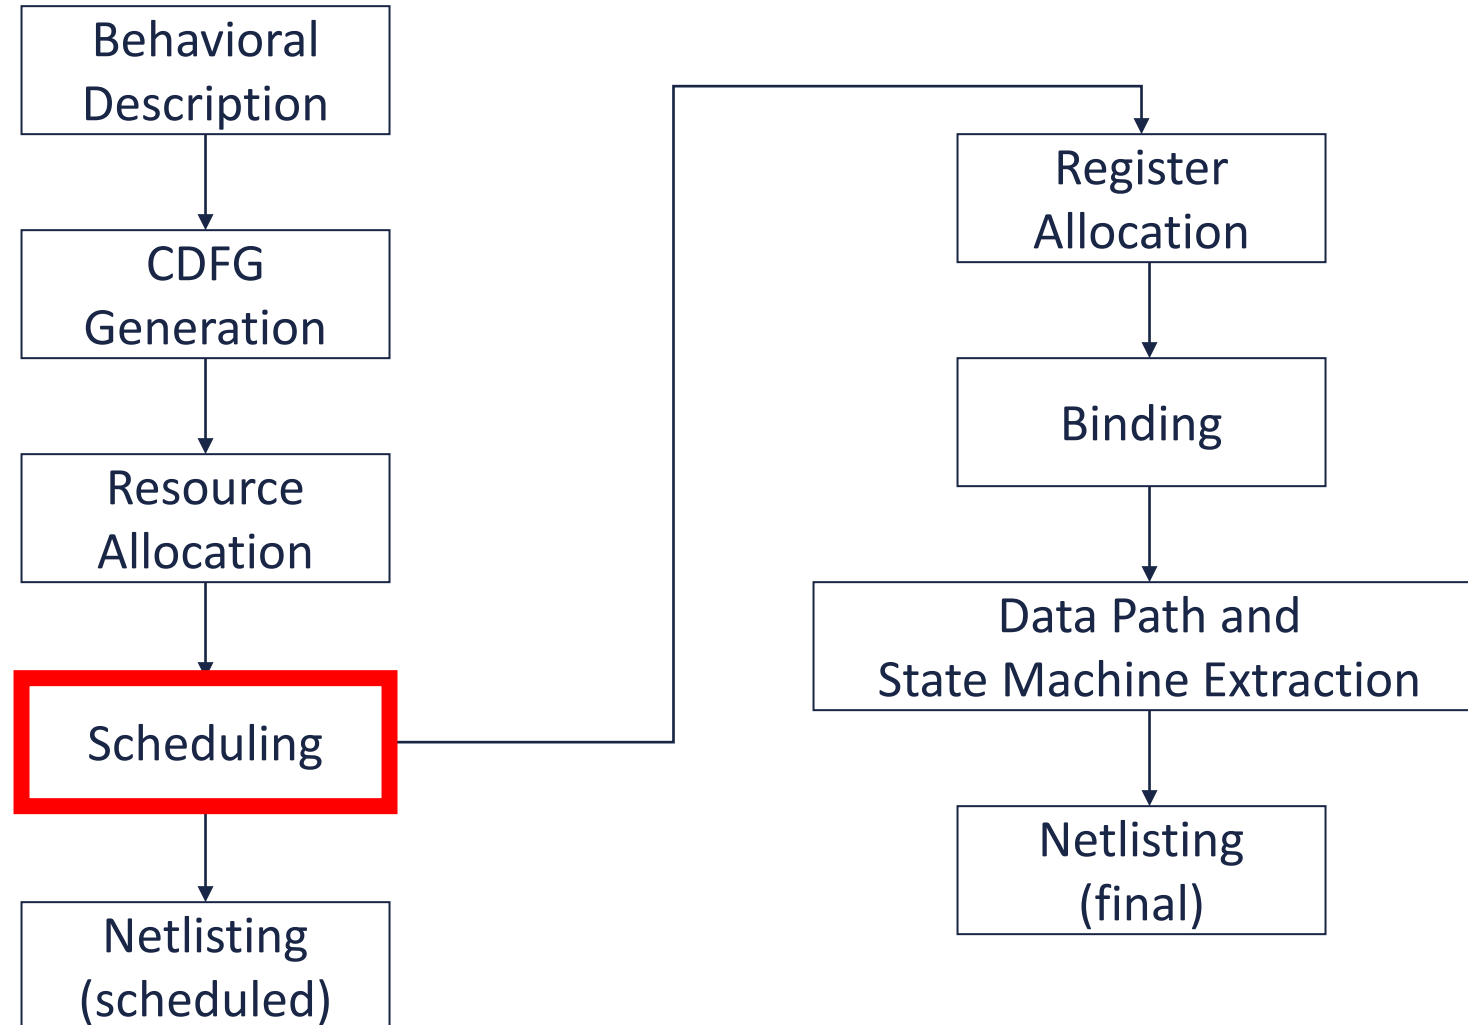
HLS Step: Resource Allocation

- Deciding how many and which kinds of resources will be used in a given implementation
- This has a major impact on final design
 - Number of operation units (multiple adders?) set the maximum parallelism that the architecture can provide
 - Reuse of overloaded operators (e.g., an adder/subtractor unit) provides smallest designs
 - Choice of buses or muxes provides parallelism vs. size
 - Choice of registers, multi-ported register files or RAM also limits parallelism in data movement

Example: Mapping of Operators to Components



High-Level Synthesis Steps



HLS Step: Scheduling

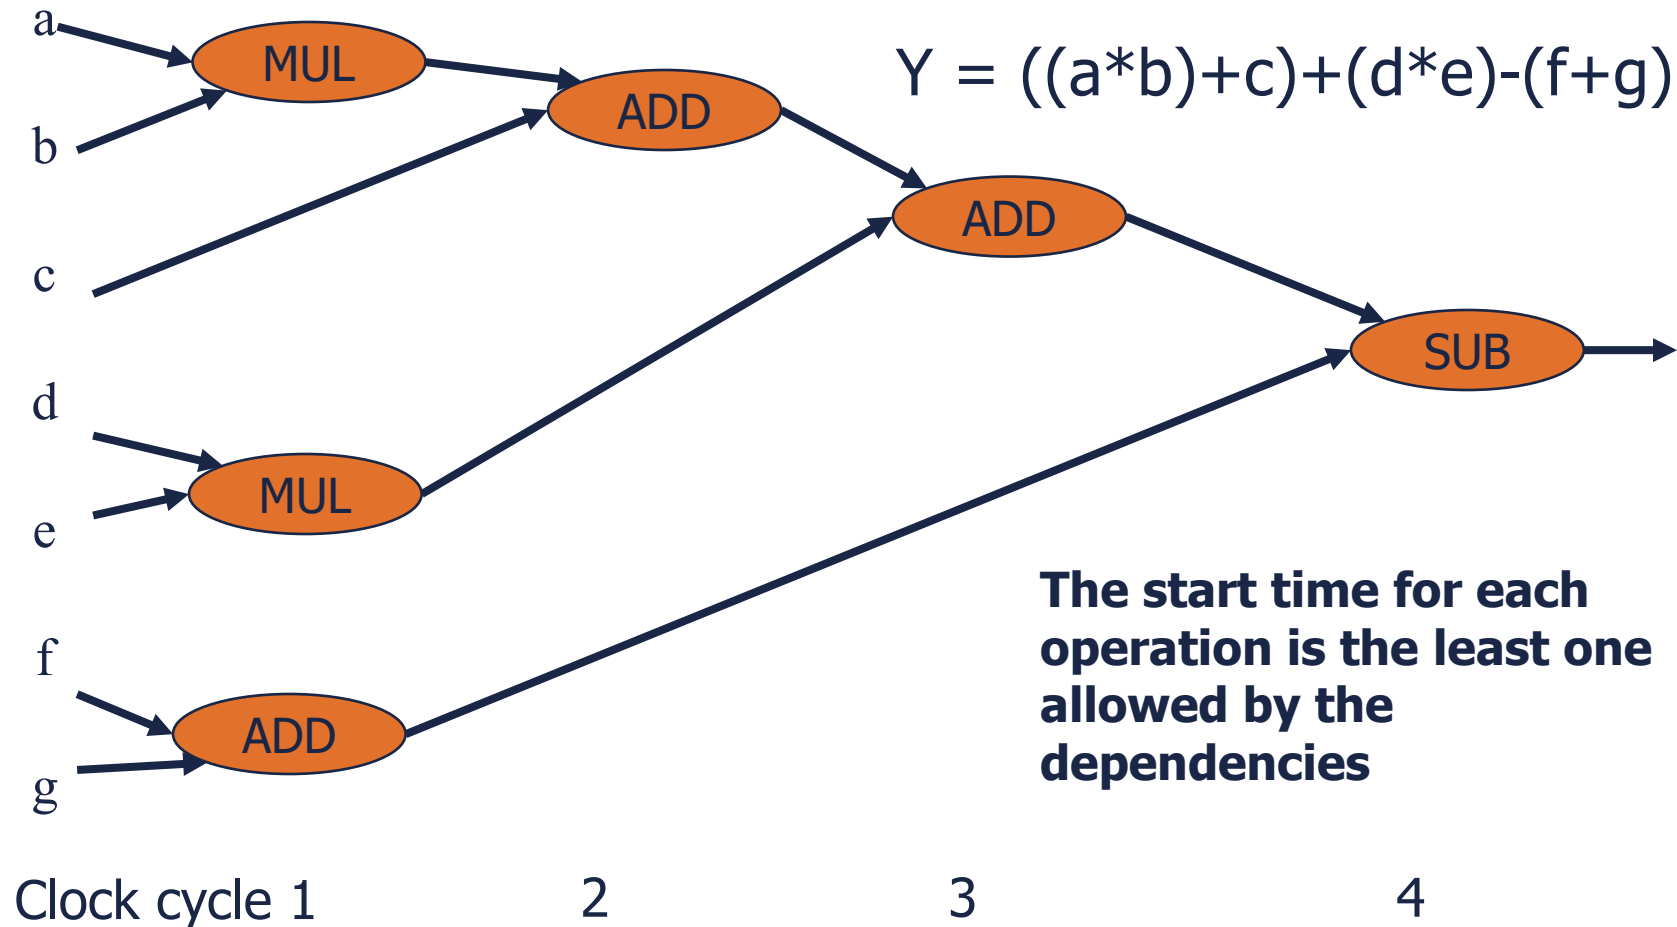
Scheduling:

- Mapping of operations to time slots (cycles)

Basic Algorithms:

- ASAP: As Soon As Possible
- ALAP: As Late As Possible
- List Scheduling

ASAP Schedule (unconstrained)




Can we do better than this in terms of the DFG itself? What is the resource usage?


ASAP Algorithm

```
ASAP ( $G(V, E)$ ) {  
    Schedule all the nodes driven only by PIs to cycle 1,  
    for all such  $v_i$  nodes,  $t_i$  (starting time) = 1;  
    Repeat {  
        Select a vertex  $v_i$  whose predecessors are all scheduled;  
        Schedule  $v_i$  by setting  $t_i = \text{MAX}(t_j + d_j); (v_j, v_i) \in E$   
    }  
    Until all the nodes are scheduled;  
    Return the schedule in a vector;  
}
```

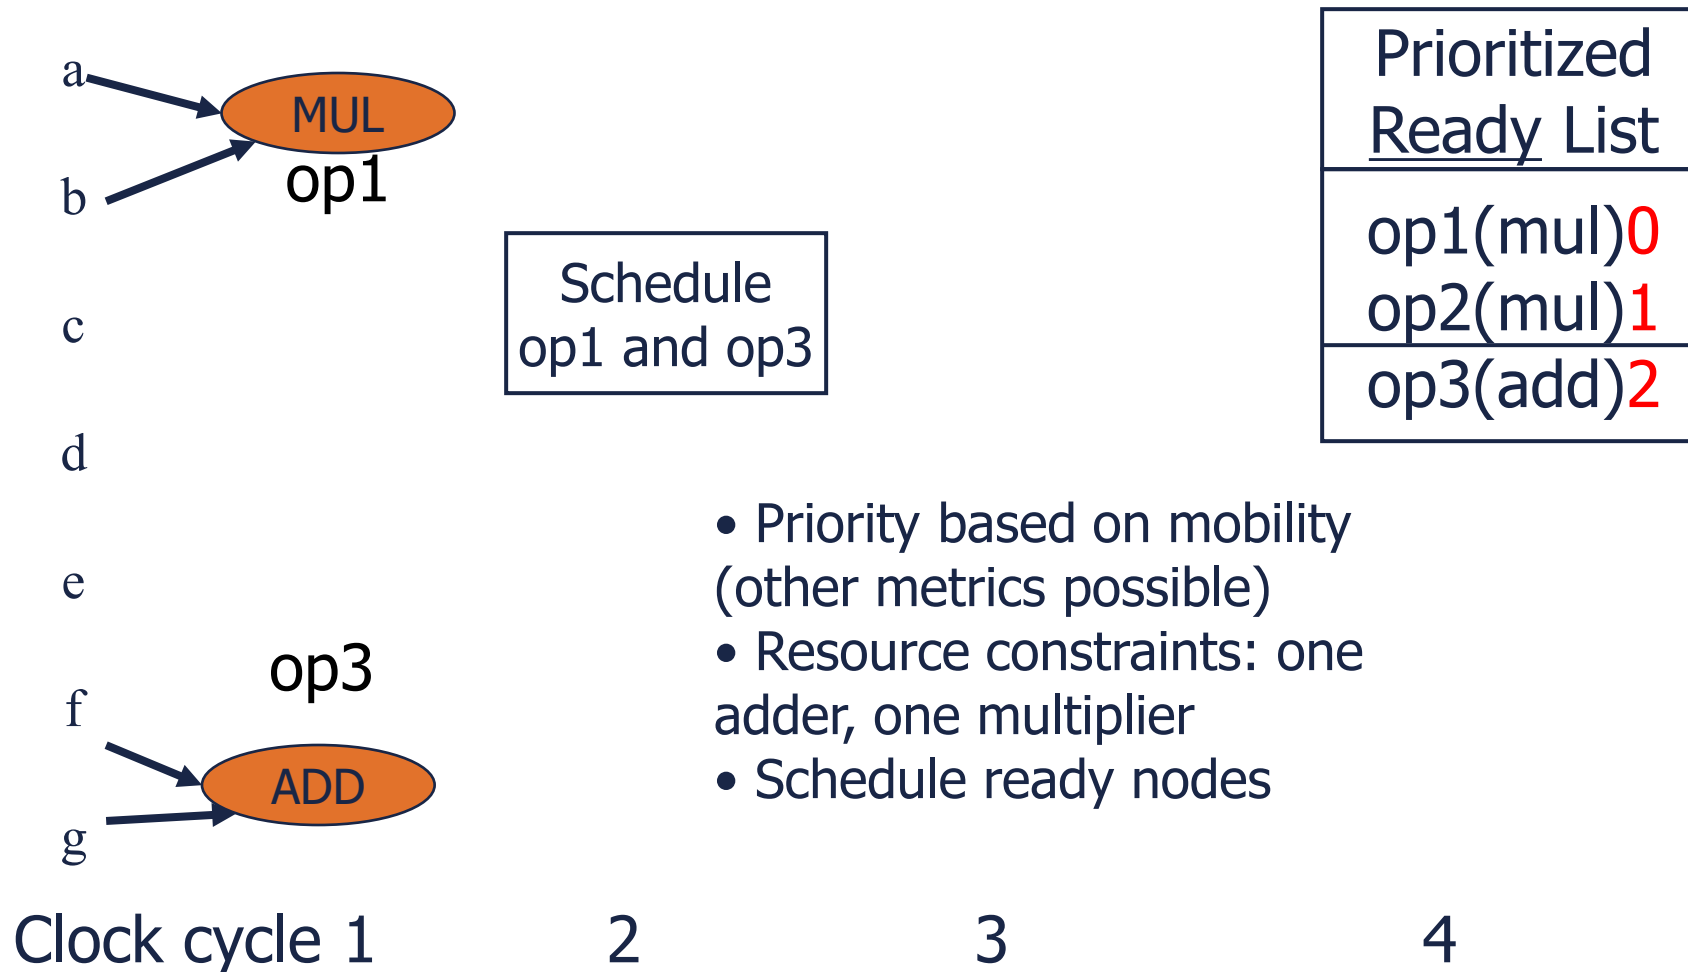

ALAP Schedule



Mobility (or Slack)

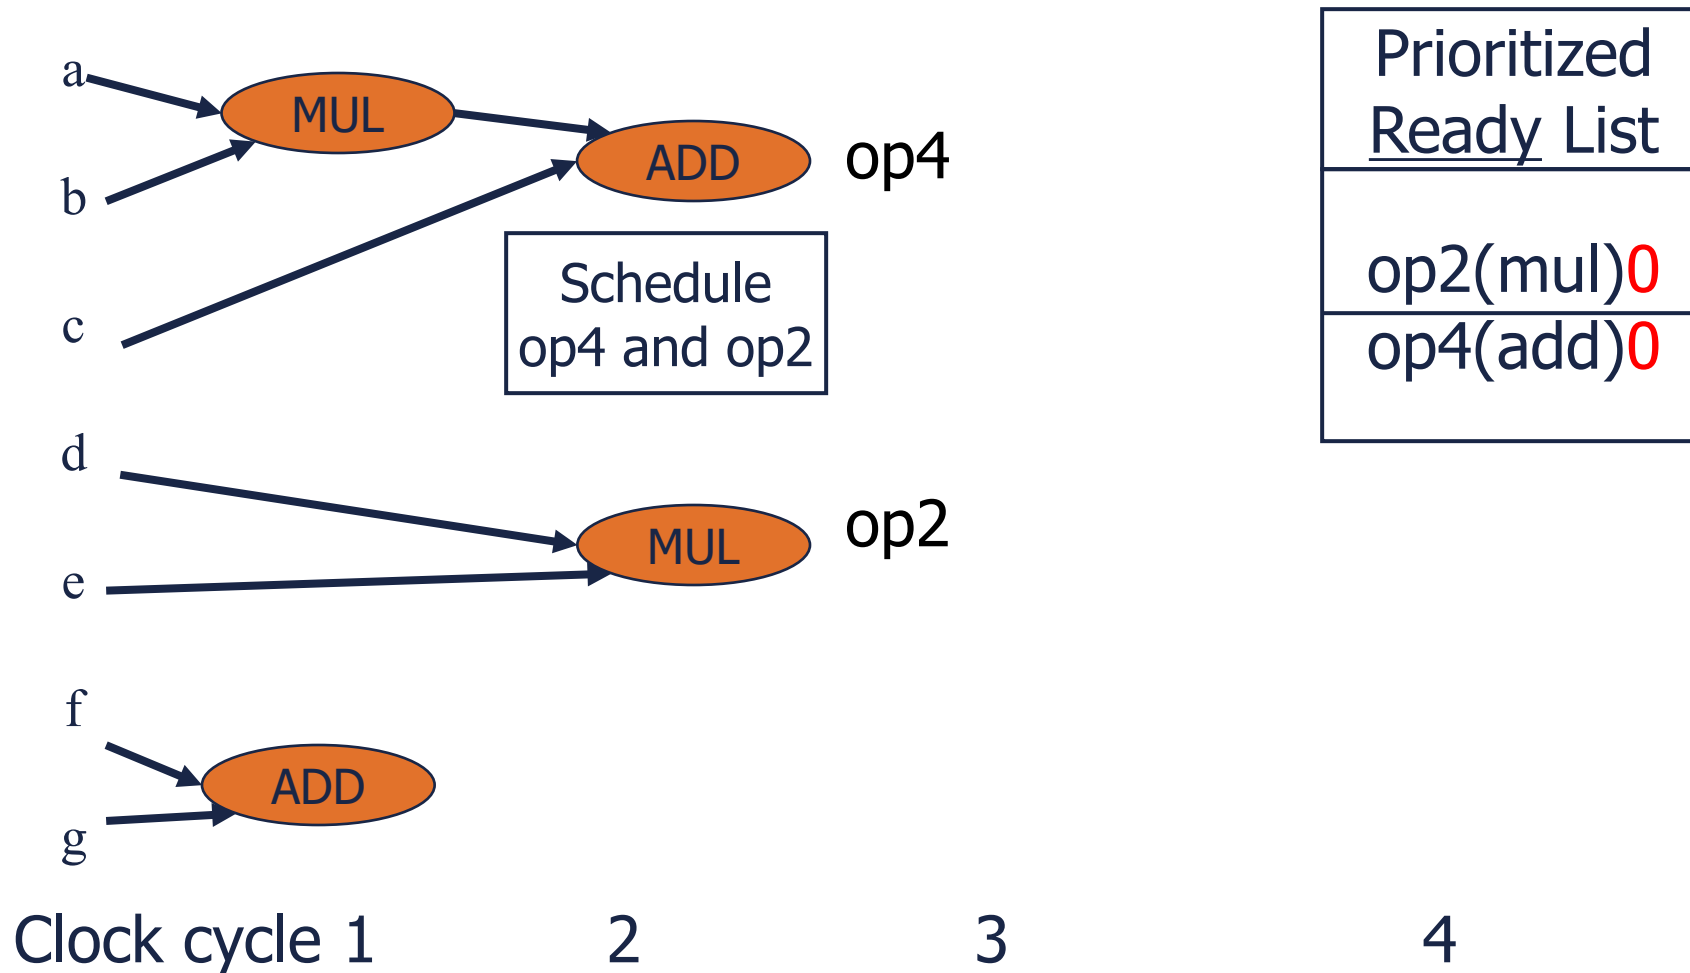


List Scheduling (1)

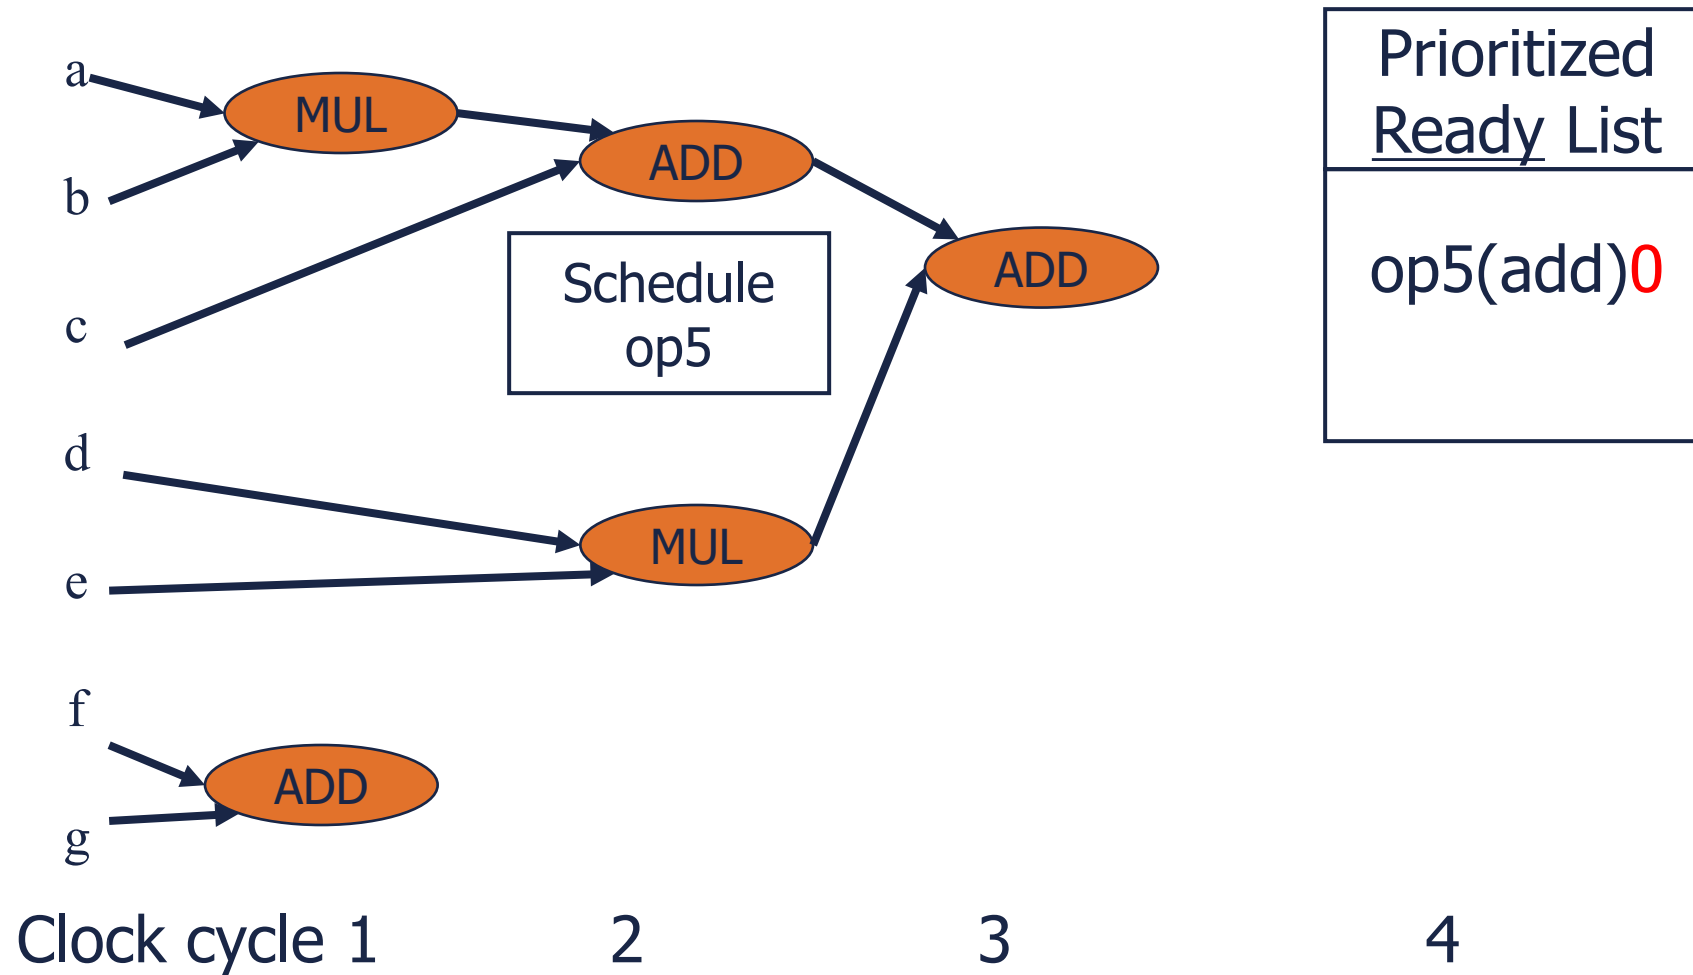


A popular scheduling algorithm

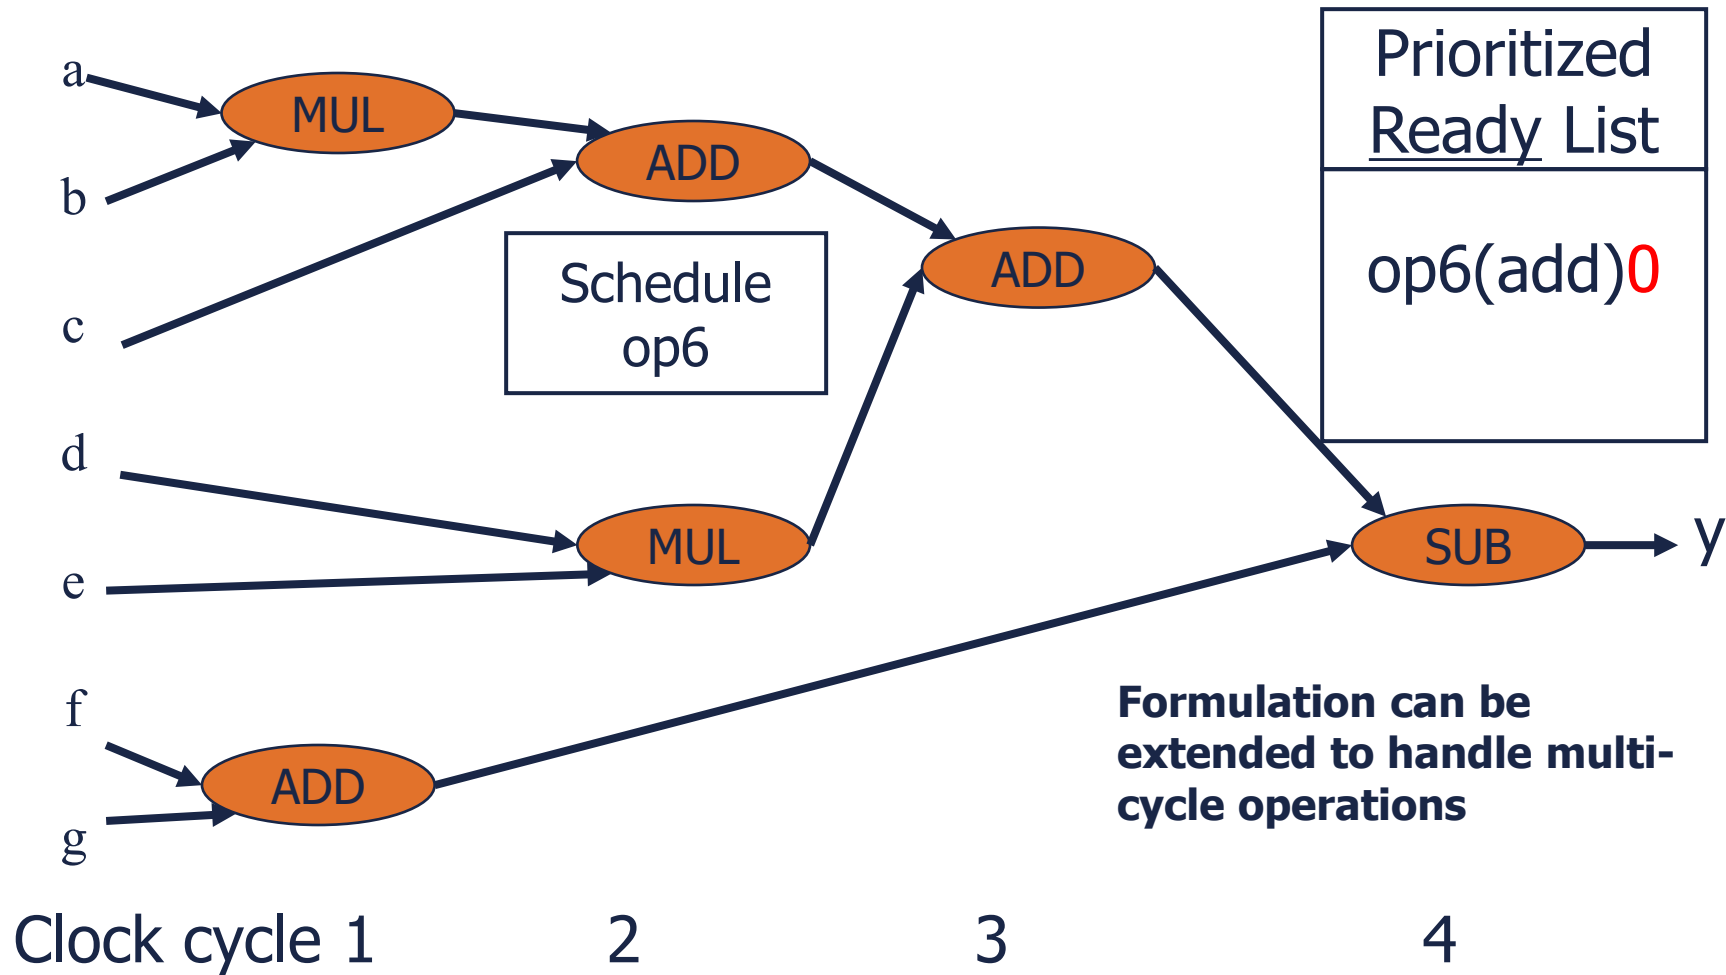
List Scheduling (2)



List Scheduling (3)




List Scheduling (4)




List Scheduling Algorithm

```
LIST_L( $G(V, E)$ ,  $a$ ) {  
     $l = 1$ ;  
    repeat {  
        for each resource type  $k = 1, 2, \dots, n_{res}$  {  
            Determine candidate operations  $U_{l,k}$ ;  
            Determine unfinished operations  $T_{l,k}$ ;  
            Select  $S_k \subseteq U_{l,k}$  vertices, such that  $|S_k| + |T_{l,k}| \leq a_k$ ;  
            Schedule the  $S_k$  operations at step  $l$  by setting  
                 $t_i = l$  for all  $i : v_i \in S_k$ ;  
        }  
         $l = l + 1$ ;  
    }  
    until (all nodes are scheduled);  
    return ( $t$ );  
}
```

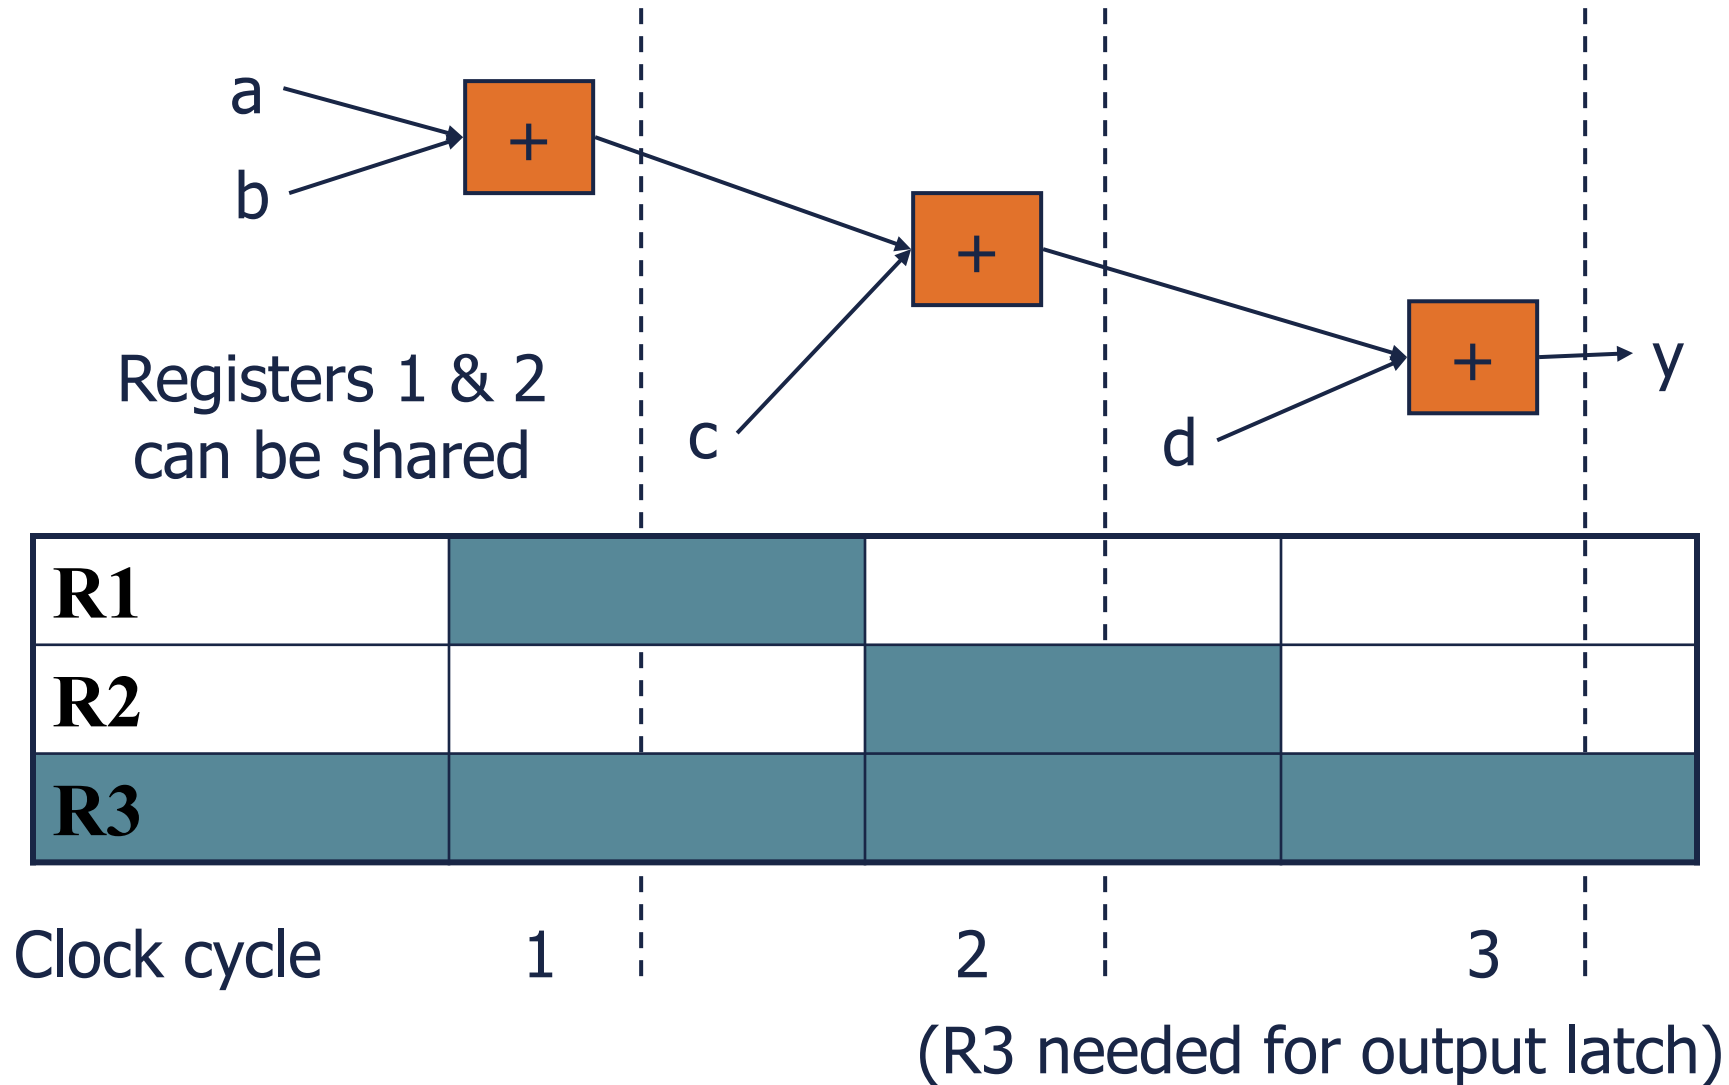

High-Level Synthesis Steps



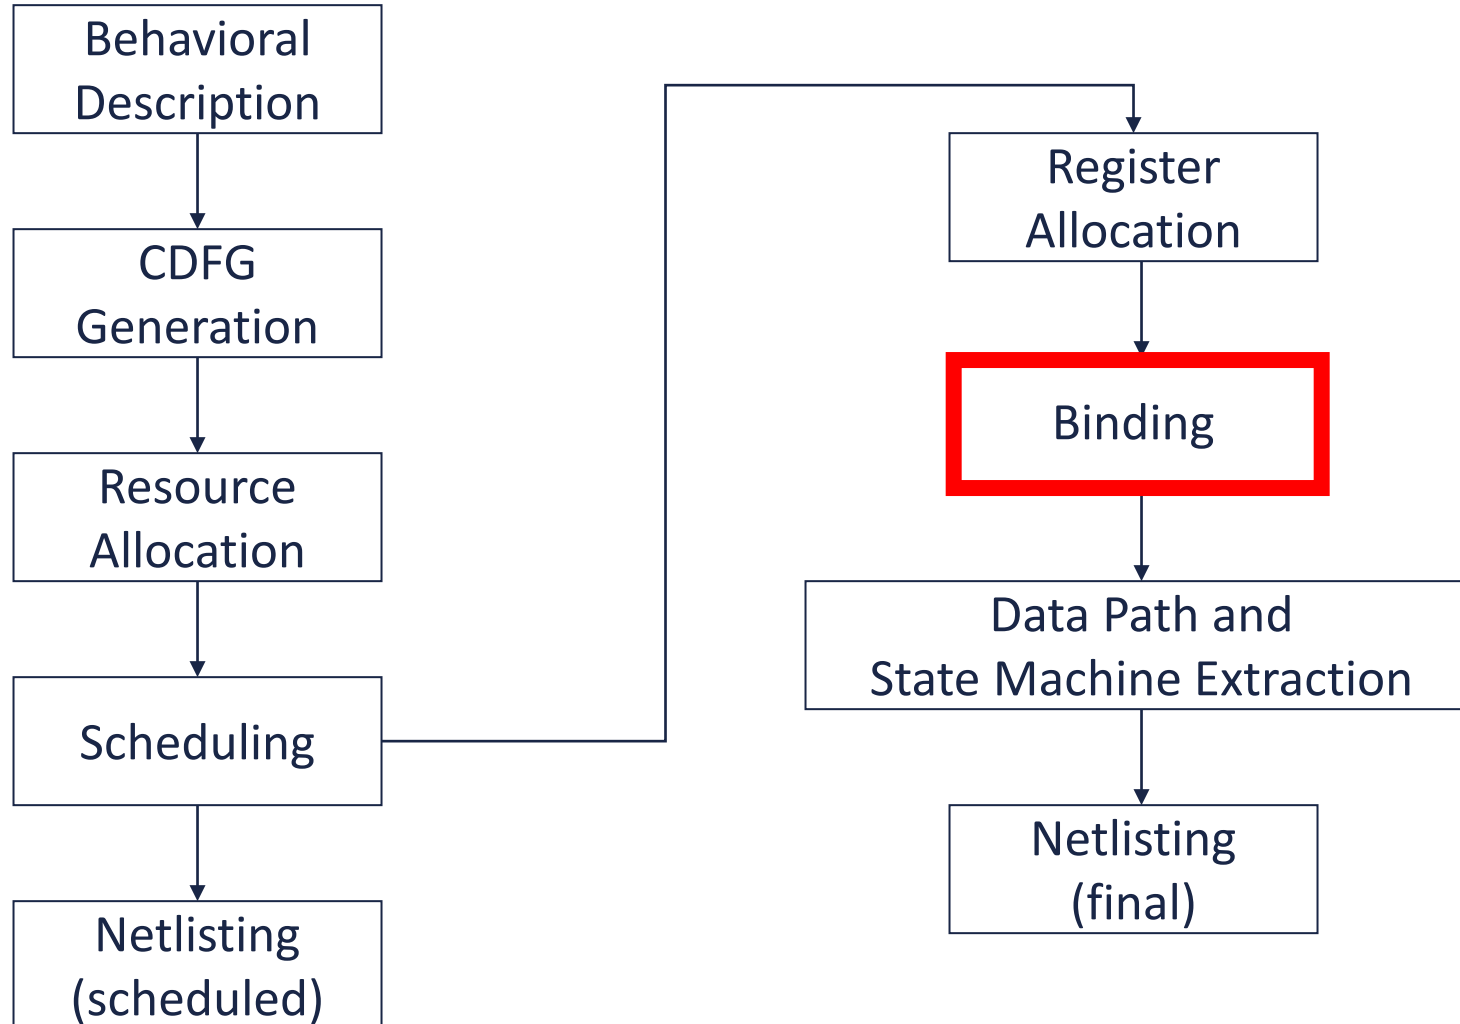
Register Allocation



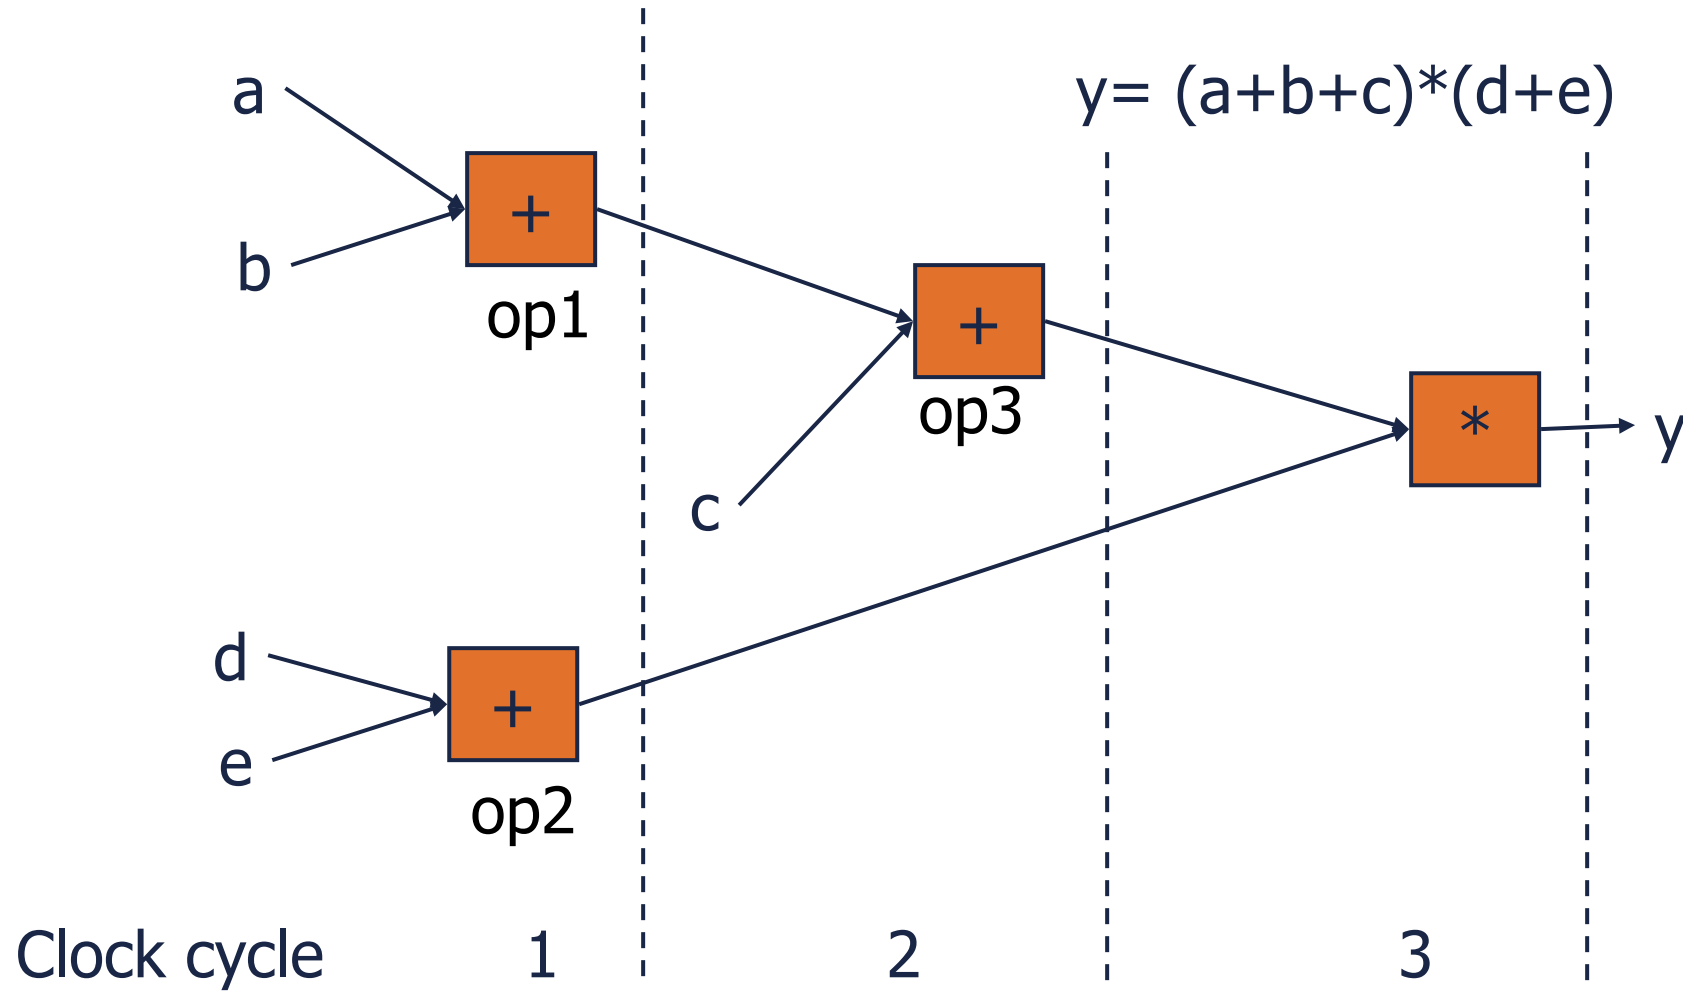
Lifetime Analysis



High-Level Synthesis Steps



Binding



Binding Choices

Binding 1

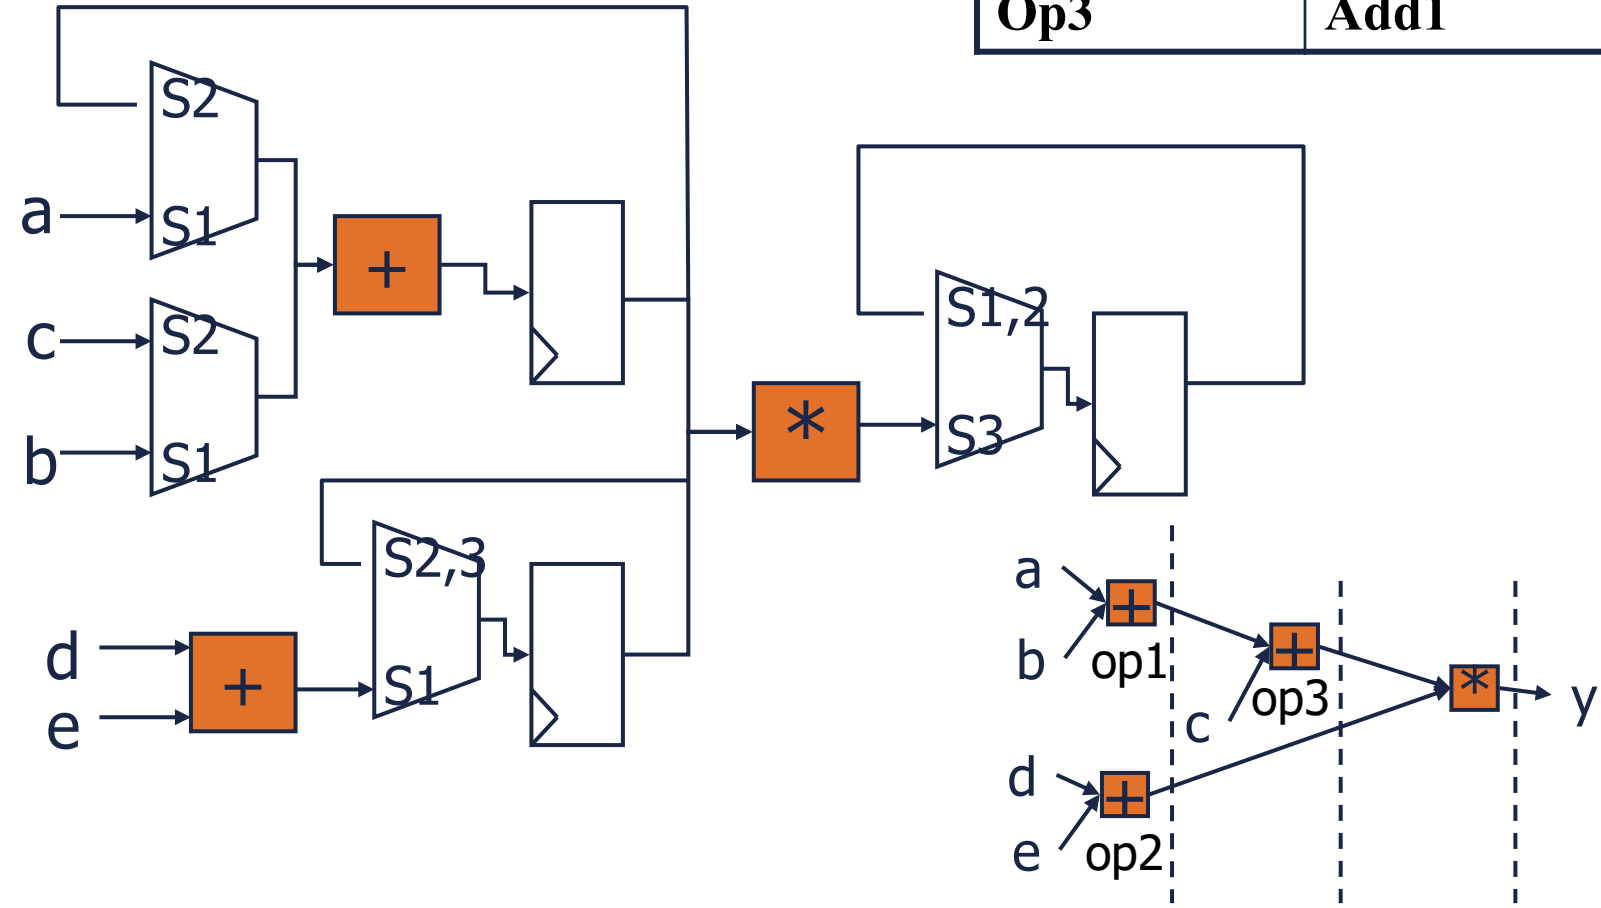
Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add1

Binding 2

Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add2

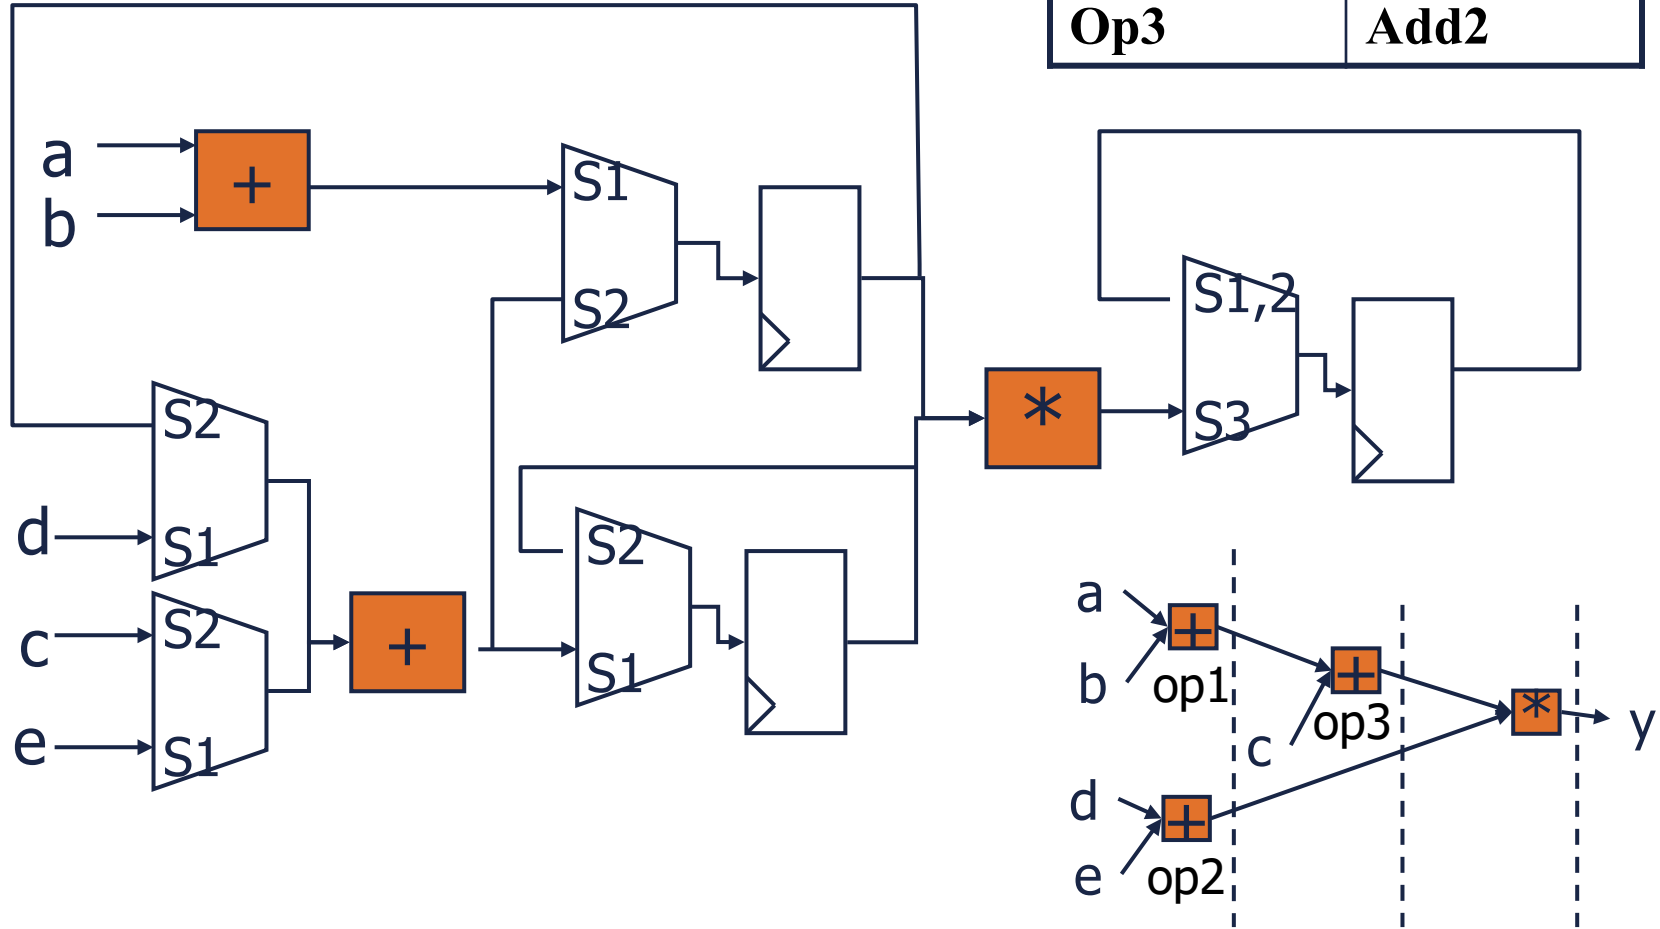
Binding 1 Results

Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add1



Binding 2 Results

Operation	Binding
Op1	Add1
Op2	Add2
Op3	Add2



Linear Programming

- Objective function: a linear function to be maximized or minimized, e.g.,
 - Maximize $c_1x_1 + c_2x_2$
- Constraints: e.g.,
 - $a_{11}x_1 + a_{12}x_2 \leq b_1$
 - $a_{21}x_1 + a_{22}x_2 \leq b_2$
 - $a_{31}x_1 + a_{32}x_2 \leq b_3$
- Non-negative variables, e.g.,
 - $x_1 \geq 0$
 - $x_2 \geq 0$
- The problem is usually expressed in *matrix form*, and then becomes:
 - Maximize $\mathbf{c}^T \mathbf{x}$
 - Subject to $\mathbf{Ax} \leq \mathbf{b}, \mathbf{x} \geq \mathbf{0}$

Integer Linear Programming (ILP)

- If the variables are all required to be integers, then the problem is called an integer programming (IP) or integer linear programming (ILP) problem.
- In contrast to linear programming, which can be solved efficiently in the worst case, integer programming problems are in the worst case NP-hard.
- 0-1 integer programming is the special case of integer programming where variables are required to be 0 or 1 (rather than arbitrary integers). This method is also classified as NP-hard, and in fact the decision version was one of Karp's 21 NP-complete problems.

ILP-Based Scheduling

- Use binary decision variables
 - $X = \{x_{il}; i = 1, \dots, n; l = 1, \dots, L\}$
 - $x_{il} \in \{0, 1\}$
- First the start time of each operation is unique

$$\sum_l x_{il} = 1, i = 1, 2, \dots, n$$

The start time of any operation v_i :

$$t_i = \sum_l l \cdot x_{il}$$

ILP-Based Scheduling: Constraints

- Given the CDFG represented by $G(V, E)$ d_j : latency of operation v_j
 $t_i \geq t_j + d_j \quad \forall i, j : (v_j, v_i) \in E$ implies

$$\sum_l l \cdot x_{il} \geq \sum_l l \cdot x_{jl} + d_j \quad i, j = 1, 2, \dots, n : (v_j, v_i) \in E$$

- Resource bounds must be met at every schedule time step.

$$\sum_{i: T(v_i)=k} \sum_{m=l-d_i+1}^l x_{im} \leq a_k, \quad k = 1, 2, \dots, n_{res}, \quad l = 1, 2, \dots, L$$

ILP-Based Scheduling: Formulation

- Denote t the vector whose entries are the start times
 - Minimize $C^T t$ such that

$$\sum_l x_{il} = 1, i = 1, 2, \dots, n$$

$$\sum_l l \cdot x_{il} - \sum_l l \cdot x_{jl} - d_j \geq 0, i, j = 1, 2, \dots, n : (v_j, v_i) \in E$$

$$\sum_{i: T(v_i)=k} \sum_{m=l-d_i+1}^l x_{im} \leq a_k, k = 1, 2, \dots, n_{res}, l = 1, 2, \dots, L$$

$$x_{il} \in \{0, 1\}, i = 1, 2, \dots, n, l = 1, 2, \dots, L$$

ILP-Based Scheduling: Vector c

- Minimize $c^T t$

The start time of any operation v_i :

$$t_i = \sum_l l \cdot x_{il}$$

- $c = [0, \dots, 0, 1]^T$
 - Minimize the latency of the schedule
- $c = [1, \dots, 1, 1]^T$
 - Finding the earliest start times of all operations

What designs fit HLS well?

- Data path, uni-directional data streaming pipeline
- Flexible in latency requirement, loose timing relationship between interfaces
- Computational kernels
 - E.g., image/video processing/arithmetic computation
 - Dataflow design is better than control flow design
 - HLS tool may provide IPs for arithmetic and other typical computational operations
- No inter-iteration dependencies
 - Preferably no dependencies between two loop iterations
 - Advanced HLS tool can handle dependencies to certain extent
- Independent kernels and module level designs
 - Not top level/system level designs entirely for HLS

Challenges for HLS

- May not be a good fit for very high speed design
 - Designers can push until the last FF and create deep pipelines with RTL
- Cycle accurate design, strong timing relationship on interfaces
- Design with feedback loops and strong timing requirement
 - E.g., accumulators whose current result affects next cycle accumulation immediately
- Control-intensive logic -- but is getting better
- HLS may need designer's manual intervention in order to generate hierarchical designs with high quality
- Coarse-grain pipelining
 - Handling multiple kernels at the top level can be a challenge
 - Advanced HLS tools can work adequately in this area (an example to follow)
- Complicated protocol handling
- Analog/mixed signal portions of the design

Where are these challenges from?

- HLS works with a design entry at a higher abstraction level
 - C/C++ has no idea of “time”
 - It might be do able to capture some timing concept, but the efficiency drops greatly, and we lose the benefit not worth it.
- RTL, as well as SystemC, has “time” built in
 - Fits design with accurate timing requirement best

Coding with HLS

- Write code with sub functions corresponding to modules
 - With hardware implementation concepts in mind
 - Think about the inter-function communication scheme
- Generate hardware for each function using HLS, leave the arguments as interface/ports
- Generate the glue logic/communication and top-level design
 - Integrate the hardware block generated from HLS into a system (e.g., a virtual platform)
 - If the hardware is hierarchical, to achieve better QoR, manual work may be required

Software vs. Hardware Compilation

C Construct	Example	Software Compilation	High-level Synthesis
Function	<code>void go() { ... }</code>	groups of instructions	hardware module
Function args and return	<code>int det(int a[16][16]);</code>	pushed onto call stack	input and output ports
Function call	<code>go();</code>	call instructions	submodule
Operations	<code>prod = a.x * b.x + a.y * b.y;</code>	computation instructions	functional units
Local variables	<code>unsigned index; int sum;</code>	architectural registers	physical registers
Arrays	<code>int A[16][16];</code>	stack, heap, static memory	memory blocks
Control flow	<code>for (i = 0; i < 16; i++) { ... }</code>	branch instructions	state machines
source code	all of the above	machine instructions	hardware modules

C/C++ Code Synthesizability for HLS

- In HLS, your code is the hardware specification

Unsynthesizable C Construct	Example	High-level Synthesis?	Synthesizable alternative
Variable memory allocation	<code>A = malloc(rows * cols * sizeof(A[0][0]));</code>	resizable RAM?	<code>int A[max_rows][max_cols];</code>
Recursion	<code>int bin_search(Array array) { ... bin_search(array); ... }</code>	infinitely nested modules?	use iterative form of binary search
Indirect calls /function pointers	<code>typedef void Callback(); void go(Callback callback) { ... callback(); ... }</code>	dynamically defined submodule?	inline or create “go” function for each callback
Operating system calls	<code>file = fopen(filename, “r”);</code>	I/O? Interrupt?	use input ports for streaming in data

C/C++ Code Synthesizability for HLS (cont'd)

Construct	Example	Synthesizable?
Debug output	<code>printf("sum: %x\n", sum);</code>	Can be translated to unsynthesizable RTL debug functions.
Time functions	<code>clock_gettime(CLOCK_REALTIME, &cur_time);</code>	Not synthesizable. HLS tools often have profiling support built-in.
C++ memory allocation	<code>Node * node = new Node();</code>	Equivalent to malloc; not synthesizable.
C++ template library	<code>std::vector<int> a; ... std::sort(a.begin(), a.end());</code>	Uses dynamic memory allocation and recursion; not synthesizable.
Floating point operations	<code>float a, b, c; ... float d = a * b + c;</code>	Supported by some HLS engines. Can be expensive, especially on FPGAs.

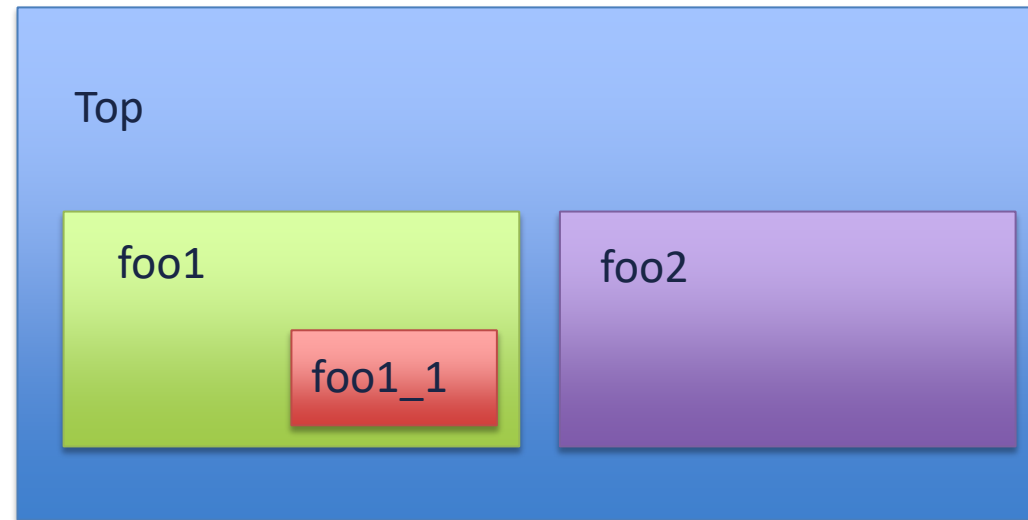
Variables

- Customize bit-widths (arbitrary bit-widths) instead of standard C data types
 - E.g., use 24-bit width instead of 32 bits if 24 bits precision is enough
 - Most HLS tools support arbitrary bit-width data types
 - E.g., `ap_int`, `ap_fixed` data types in Vivado HLS
 - Improves efficiency, performance, and saves area

Functions

- C functions are synthesized to RTL modules
 - By default, sub functions become sub modules
 - Usually, every function call is an instance of the RTL module
 - Inline function is an exception – it flattens the design

```
void top(int a, int b) {  
    foo1(a);  
    foo2(b);  
}  
void foo1(int a) {  
    foo1_1(a);  
}
```



Functions: Splitting

- Use subfunctions to improve parallelism
 - Use function to maintain hierarchy
 - Most tools support dataflow/block-level pipelining and parallelism
 - May not be feasible due to dependencies

```
void top (int a[100], int b[100]) {  
    for(i=0; i<100; i++) {  
        a[i] = i;  
    }  
    for(j = 0; j < 100; j++) {  
        z += b[j];  
    }  
}
```



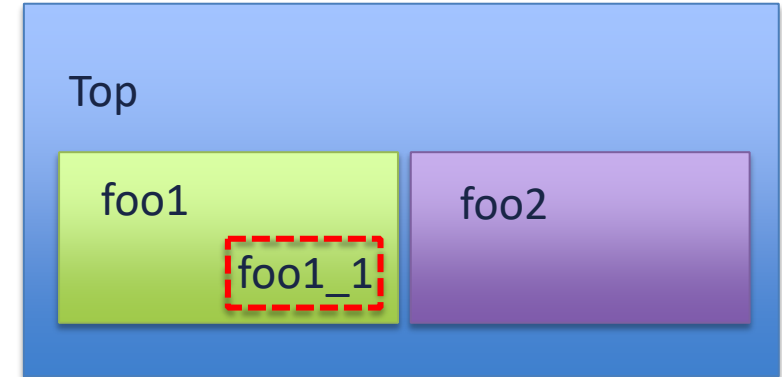
```
void foo1 (int a[100]) {  
    for(i=0; i<100; i++) {  
        a[i] = i;  
    }  
}  
void foo2 (int b[100]) {  
    for(j = 0; j < 100; j++) {  
        z += b[j];  
    }  
}  
void top (int a[100], int b[100]) {  
    foo1(a);  
    foo2(b);  
}
```



Functions: Inlining

- Inline functions to flatten the hierarchy
- Pros:
 - Enable more optimization opportunities
 - Creates more objects to schedule and optimize
 - Enable resource sharing
 - Eliminates function call overhead
- Cons:
 - Module instance is effectively copied for each inlined call
 - Inlining of multiple function calls increases area cost

```
void top(){int a, int b){  
    foo1(a);  
    foo2(b);  
}  
void foo1(int a){  
    foo1_1(a);  
}  
inline void foo1_1(int a){}
```



Inlining: Example


- Merge foo1 and foo2, and apply loop pipelining

```
int foo1(int a, int b, int c) {  
    z0 = a * b;        //s1_1  
    z1 = c * b;        //s1_2  
    return (z0+z1);  
}  
int foo2(int c, int d, int a){  
    z2 = c * d;        //s2_1  
    z3 = d * a;        //s2_2  
    return (z2+z3);  
}  
int top(int a, int b, int c, int d){  
    y1 = foo1(a, b, c);  
    y3 = foo2(c, d, a);  
    return (y3+y1);  
}
```



```
int top(int a, int b, int c, int d){  
    z0 = a * b;        //s1_1  
    z1 = c * b;        //s1_2  
    z2 = c * d;        //s2_1  
    z3 = d * a;        //s2_2  
    return (z0+z1+z2+z3);  
}
```

Assume each multiplication takes multiple cycles, and it can be pipelined



Also, enable resource sharing, e.g. for the multiplier

Loops: Bounds

- Not recommend using variable loop bounds, because:
 - HLS tools may not be able to estimate the performance
 - Limits the optimization opportunities
 - E.g., cannot deliver complete loop unrolling, thus prevent the pipelining
 - Solutions: Set the bounds to the maximum value, and conditionally execute the loop body

```
for (i=0; i < len; i++) {  
    A[i] += b;  
}
```

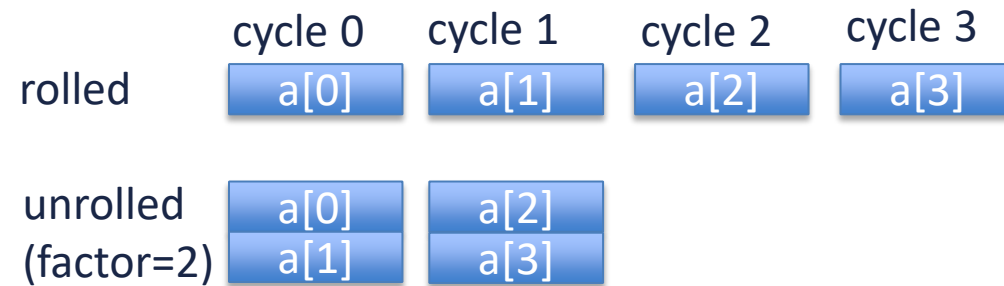


```
for (i=0; i < N-1; i++) {  
    if(i < len)  
        A[i] += b;  
}
```

Loop Unrolling

- Most HLS tools support automatic loop unrolling (by using HLS unroll pragmas)
 - But have strict limits, e.g., constant loop bounds
- May need manual loop unrolling to
 - Improve pipeline opportunity
 - Improve parallelism
- Things to notice
 - Need to consider memory port limit when accessing array
 - Expensive hardware replication
 - Need to balance the unroll level

```
void top(){
    for(i=0; i<=N; i++){
        a[i] = i;
        ...
    }
}
```



Loop Fusion/Merging

- Merge several loops into one loop
- Some tools support loop fusion
 - Strict bound check of loops, e.g., identical loop bounds
- Manually loop fusion can be used to
 - Improve data locality
 - Shorten latency
 - Share resources

```
void top(){  
    for(i=0; i<N; i++){  
        a[i] = i;  
    }  
    for(i=1; i<N; i++){  
        x+=a[i-1];  
    }  
}
```



```
void top(){  
    for(i=0; i<N; i++){  
        a[i] = i;  
        if(i>=1)  
            x+=a[i-1];  
    }  
}
```

After fusion, loop bodies are optimized together

Loop Perfectization

- Convert loops to perfect loop
 - All statements are in the innermost loop
 - An example: Loop Sinking

```
for (i=0, i<20; i++){  
    a = 0;  
    for(j=0; j<20; j++){  
        a += A[j] + j;  
    }  
    B[i] = a * 20;  
}
```



```
for (i=0, i<20; i++){  
    for(j=0; j<20; j++){  
        if(j==0) a =0;  
        a += A[j] + j;  
        if(j==19) B[i] = a*20; }  
}
```

HLS result (Use Vivado_HLS):

```
Loop_I:  
    *Trip count: 20  
    *Latency: 480  
Loop_J:  
    *Trip count: 20  
    *Latency: 21  
    *Pipeline II: 1  
    *Pipeline depth: 2
```



```
Loop_I_Loop_J:  
    *Trip count: 400  
    *Latency: 401  
    *Pipeline II: 1  
    *Pipeline depth: 2
```

The original code can only pipeline Loop_J. After perfection, two loops can be flattened and pipelined.

Why Loop Perfectization?

- Form very regular loop structure
 - Some tools can only flatten the perfect loops
 - Easy for HLS scheduling and discovering optimization opportunities
 - Improves performance by eliminating the loop entering/exiting/cost
- Different ways of loop perfection
 - Loop sinking
 - Loop unrolling
 - Loop splitting

Loop Perfectization: Code Sinking

```
for (i=0, i<20; i++){
    a = 0;
    for(j=0; j<100; j++){
        a += A[j] + j;
    }
    B[i] = a * 20;
}
```



```
for (i=0, i<20; i++){
    for(j=0; j<20; j++){
        if(j==0) a = 0;
        a += A[j] + j;
        if(j==19) B[i] = a*20;
    }
}
```

- Pros:
 - Easy to implement
 - No performance penalty for branch operations in pipeline
- Cons:
 - Expensive for hardware
 - sometimes the branches can be messy when code is complicated
 - Could sequentialize loop nest and affect performance
 - Change the data dependency

Loop Perfectization: Loop Unrolling

```
for (i=0; i<100; i++) {  
    z = a*b[i];  
    z+= c*d[i];  
    for(j=0; j<4; j++) {  
        z+=x[j] + j;  
    }  
}
```



```
for (i=0, i<100; i++) {  
    z = a*b[i];  
    z += c*d[i];  
    z += x[0] + 0;  
    z += x[1] + 1;  
    z += x[2] + 2;  
    z += x[3] + 3;  
}
```

- Pros:
 - Enable pipelining at outer loop
 - Enable resource sharing
 - More statements to schedule in the same loop level
 - Higher performance, good for expensive operations
- Cons:
 - Not always feasible
 - Expensive hardware, best if the innermost loop is small

Loop Perfectization: Loop Splitting

```
for (i=0; i<100; i++){  
    z = a*b[i];  
    z += c*d[i];  
    for(j=0; j<4; j++){  
        z+=x[i] + j;  
    }  
}
```



```
for(i=0; i<100; i++) {  
    z = a*b[i];  
    z += c*d[i];  
}  
for(i=0, i<100; i++) {  
    for(j=0; j<4; j++) {  
        z += x[0] + 0;  
        z+= x[1] + 1;  
        z+= x[2] + 2;  
        z+= x[3] + 3;  
    }  
}
```

- Pros:
 - Enable optimization per loop
 - Opportunity for dataflow streaming between two loops
- Cons:
 - Not always possible because of data dependency

Arrays

- By default, it is implemented as RAM/ROM
- Expensive and have limited ports
- Optimizations:
 - Array streaming
 - Replace the RAM with FIFO, if possible
 - Requires array access to be sequential
 - Requires the array access order to be the same
 - Transform the code to make array access sequential
 - Array partitioning
 - Eliminate the ports limit


Array Streaming: Example

```
for (i=0; i<100; i++)
    for(j=0; j<100; j++){
        A[i][j] = a * b; //row
        ...
    }
}
for (i=0; i<100; i++)
    for(j=0; j<100; j++) {
        C[i][j] = A[j][i]; //col
    }
}
```




```
hls::stream<int> A;
for (i=0, i<100; i++){
    for(j = 0; j<100; j++)
        A.write(a*b);
}
for(i=0, i<100; i++){
    for(j = 0; j<100; j++){
        A.read(C[j][i]);
    }
}
```

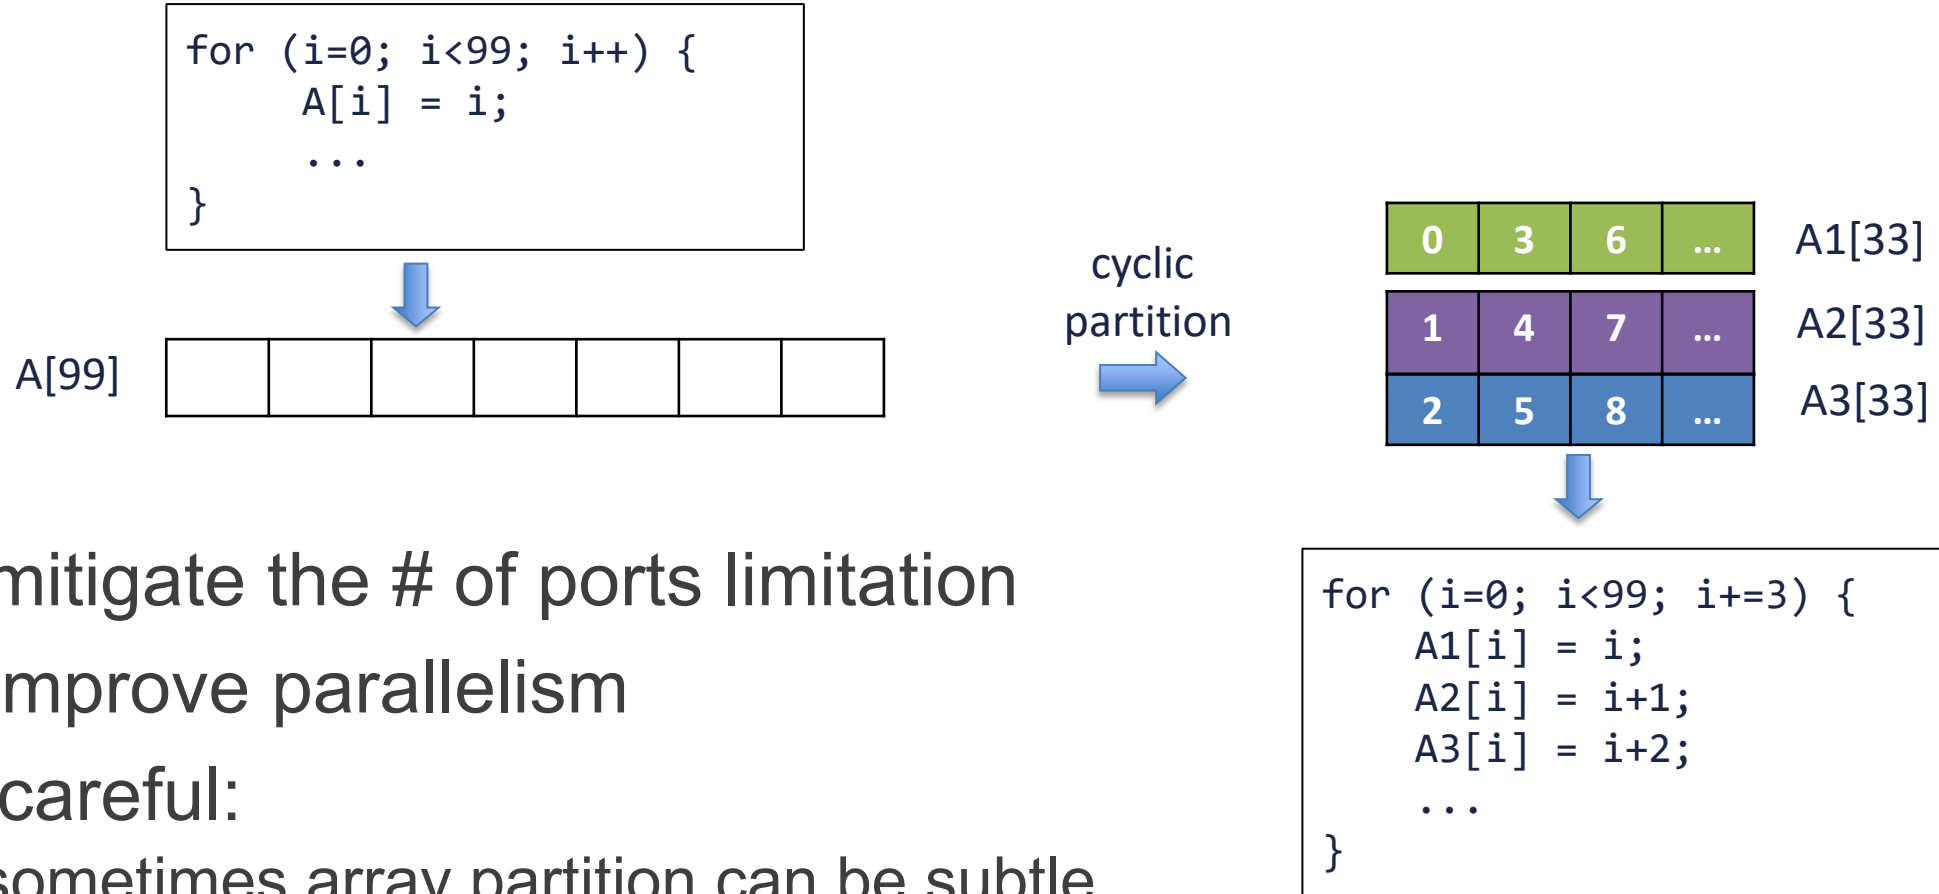
- Original array A access order:
 - Loop1: row; Loop2: column
 - No dataflow pipeline
 - No streaming



- New array A access order:
 - Loop1: row; Loop2: row
 - Dataflow pipeline
 - Streaming



Array Partitioning



- To mitigate the # of ports limitation
- To improve parallelism
- Be careful:
 - sometimes array partition can be subtle
 - can easily introduce bugs

Some Tips for C-based HLS Design

- Understand the HLS tool behavior and the optimization goal
 - No unique coding style is “useful” or “good” for all optimizations
- Reuse data and minimize array access
 - Once the data has been read into a block or buffer and reused, it can improve the parallelism.
 - Array ports are limited. Array partitioning can be an alternative, but it's not free.
- Loop transformation to enable parallelism
 - Pipelining/Unrolling/Dataflow
- Bit-width selection for efficient hardware
- The optimization goals for CPU and Hardware are different
 - E.g., branch mis-prediction
 - But many software code optimization techniques are useful for hardware design too

EECS 298: System-on-Chip Design

Sitao Huang, sitaoh@uci.edu

