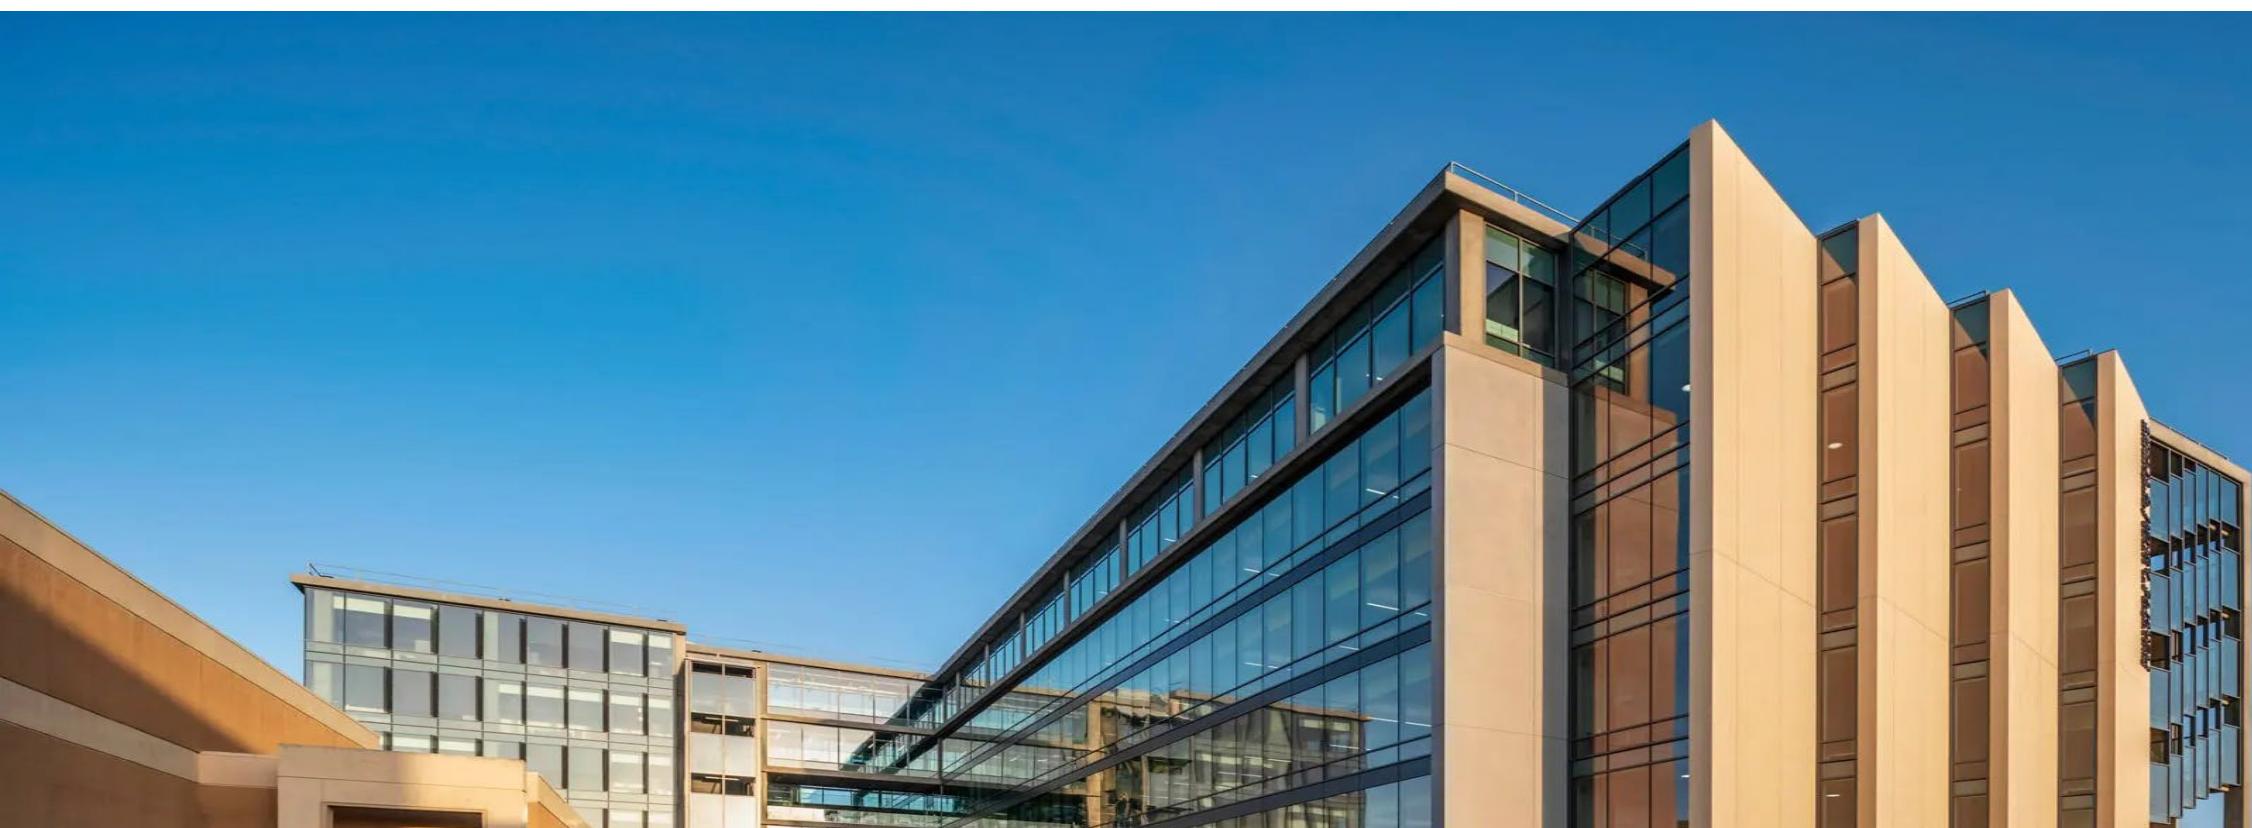


Lecture 4:

System-on-Chip FPGAs

Sitao Huang, sitaoh@uci.edu

October 18, 2023



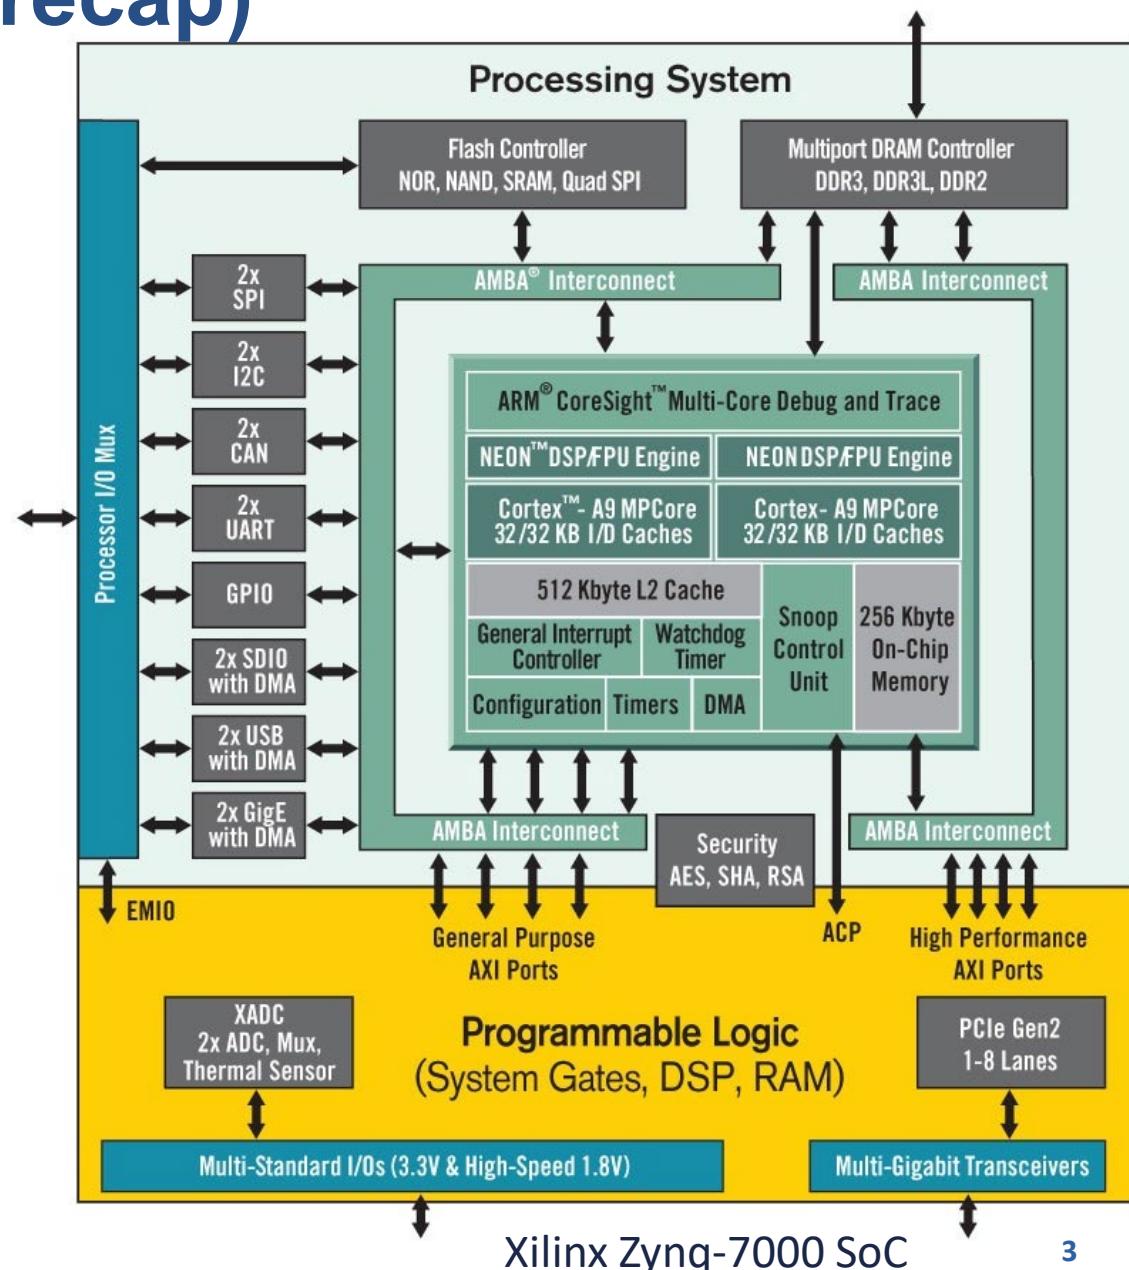
AMD Xilinx FPGA Design Tools

Vitis Unified Software Platform

- **Option 1:** UCI Engineering Remote Labs:
 - Remote Windows machines
 - <https://laptops.eng.uci.edu/computer-labs/remote-labs>
- **Option 2:** UCI EECS Instructional Servers
 - Servers: laguna.eecs.uci.edu, bondi.eecs.uci.edu, crystalcove.eecs.uci.edu
 - Setup X11 display server (Xming for Windows and XQuartz for macOS)
 - SSH to any of these servers with your UCI credentials
 - `ssh -X yourUCInetID@laguna.uci.edu`
 - Xilinx tools installed under: /ecelib/eceware/xilinx_2022.2/
 - Initialization: `source /ecelib/eceware/xilinx_2022.2/Vitis/2022.2/settings64.csh`
 - Start Vitis HLS: `vitis_hls &`
- **Option 3:** Install Vitis tools on your own computer (>100 GB)
 - <https://www.xilinx.com/support/download.html>

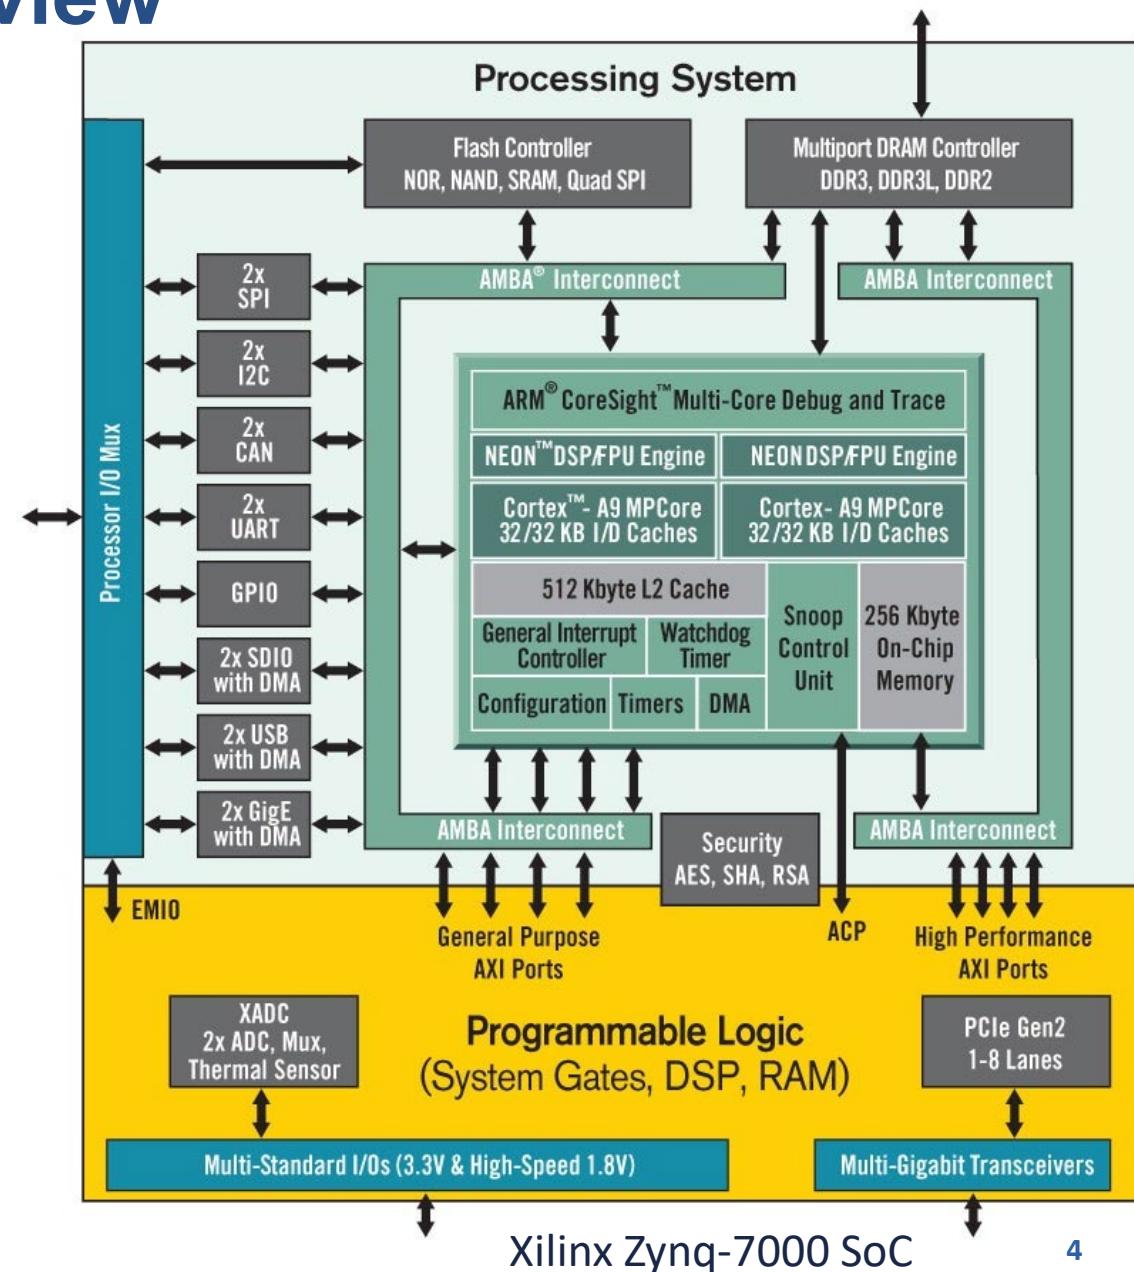
AMD Xilinx Zynq SoC Overview (recap)

- Can be divided into two parts:
 - **PS**: Processing System (ARM CPU)
 - **PL**: Programmable Logic (FPGA)



AMD Xilinx Zynq-7000 SoC Overview

- PS: Dual Core ARM A9
 - 1-2 GOPS
- Floating point support
 - ARM NEON support as well
- Up to 1GHz Clock
- L2 Cache
 - Unified 512KB
- AXI High performance interconnects
- 256KB on-chip memory scratchpad
- PL: FPGA
 - 85K Logic cells + 220 DSP slices
 - 10-100 GOPS!



Zynq SoC Boot Flow

Multi-Stage

1. Run from ROM

- a) Copy FSBL from boot device to OCM (on-chip memory)

2. First Stage Boot Loader (FSBL)

- a) Load Uboot from boot device to DDR
- b) Program PL
- c) Initiate PS boot

3. Uboot

- a) Load linux kernel to DDR
- b) Device tree init
- c) FPGA init

4. OS Boot

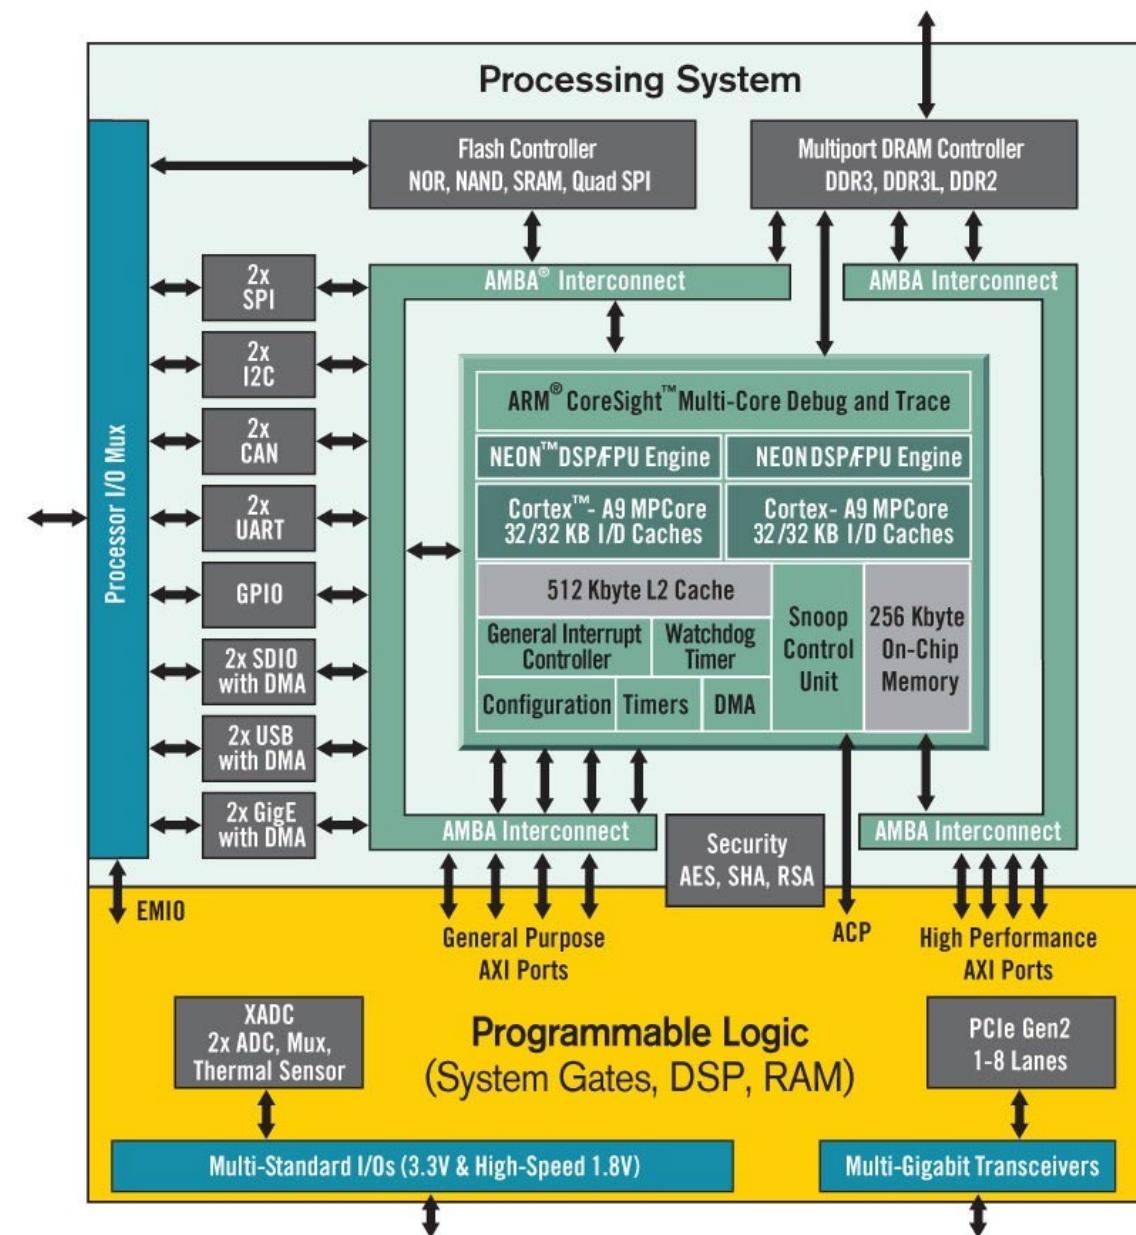
- a) Linux Boot

More Info:

- <http://www.wiki.xilinx.com/Getting+Started>
- <http://xillybus.com/tutorials/device-tree-zynq-1>
- <https://xilinx.github.io/Embedded-Design-Tutorials/docs/2023.1/build/html/docs/Introduction/ZynqMPSoC-EDT/ZynqMPSoC-EDT.html>

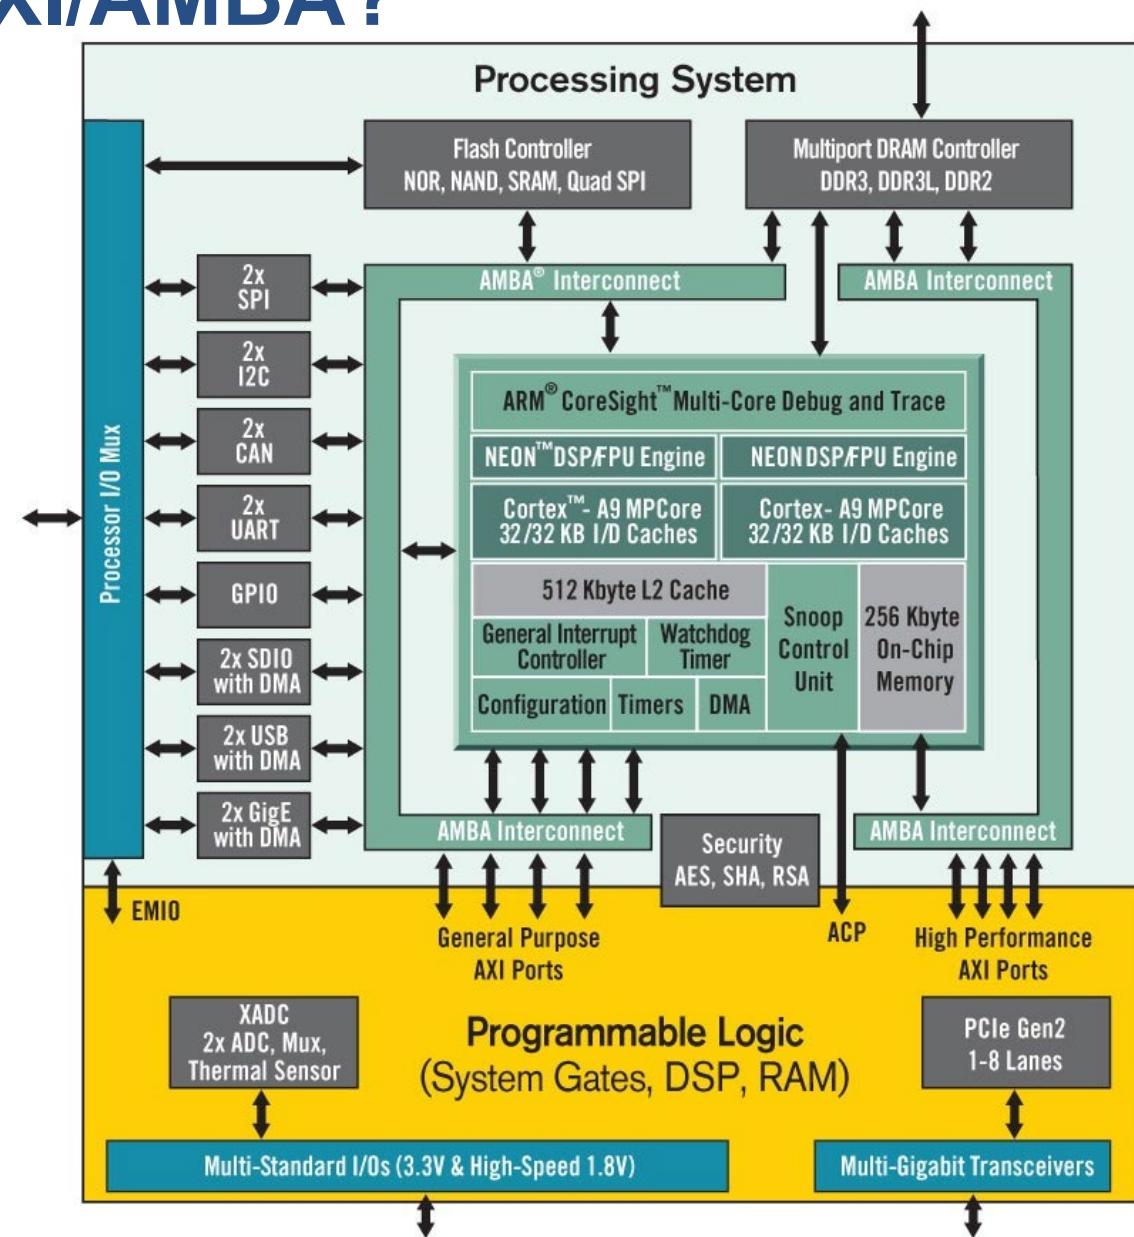
AMD Xilinx Zynq SoC: Interrupts

- Several interrupts available for PS-PL interface
- 16 peripheral interrupts available from PL to PS
 - Use for accelerator dev
- PS to PL interrupts exists as well
 - Read up via manual



AMD Xilinx Zynq SoC: What is AXI/AMBA?

- AMBA: Advanced Microcontroller Bus Architecture
 - Protocol
 - Open standard, on-chip interconnect for SoC
- AXI:
 - Advanced eXtensible Interface
 - Very common AMBA interface
- Why:
 - Flexible
 - IP reuse
 - Ease of use



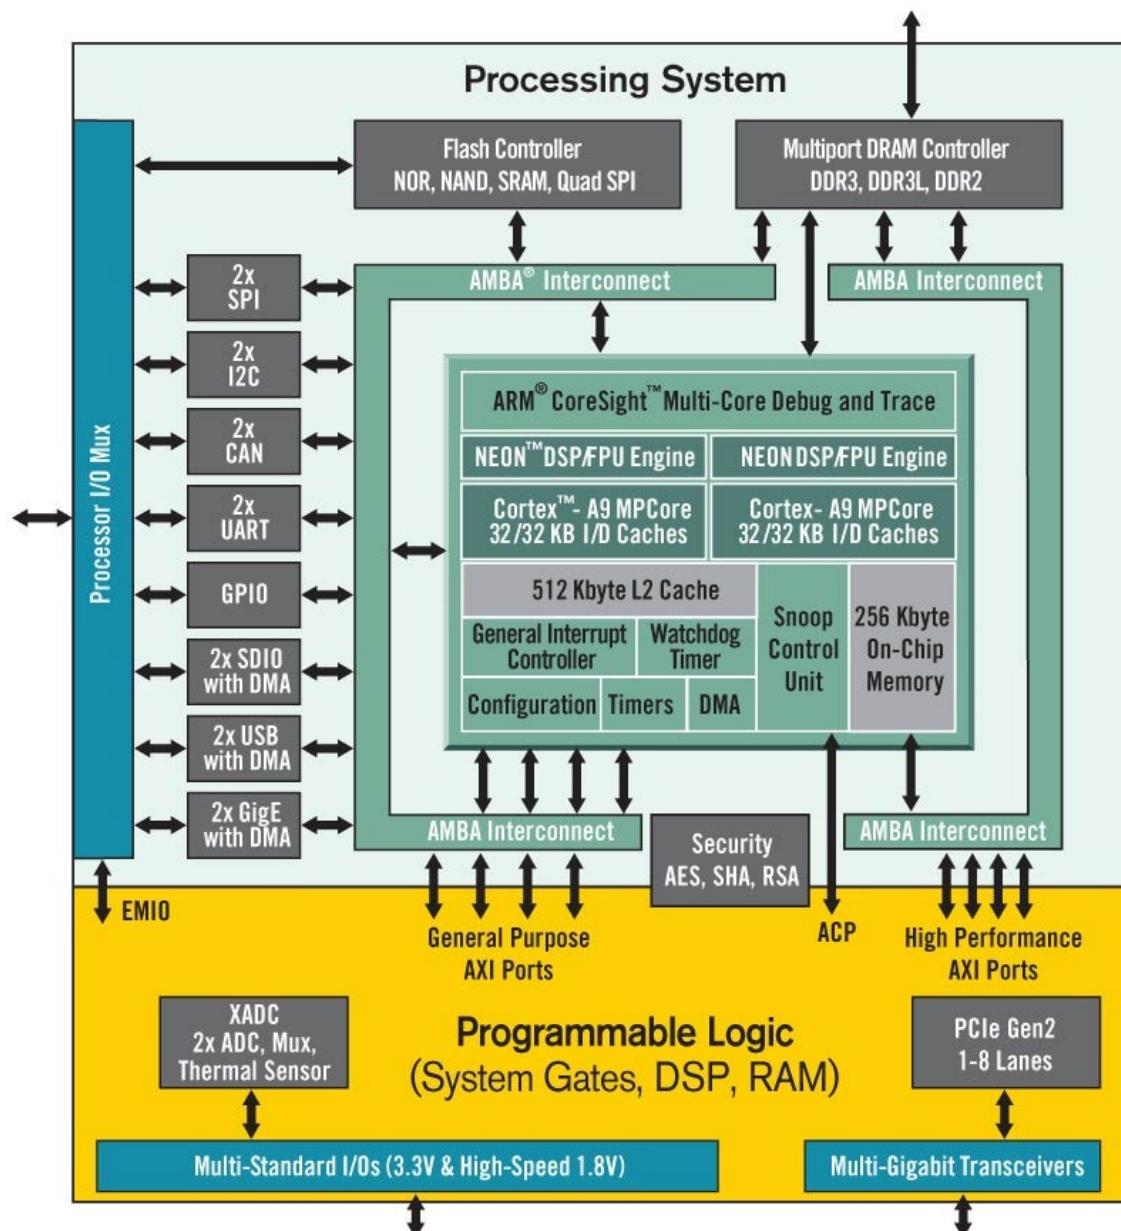
AMD Xilinx Zynq SoC: AXI

AXI Interfaces:

- AXI4 *Mainly for Data Movement*
 - High performance
 - Memory mapped
- AXI4-Lite *Mainly for Control and Status*
 - Low throughput
 - Memory mapped
- AXI4-Stream *Mainly for Data Movement*
 - High performance
 - Streaming Data

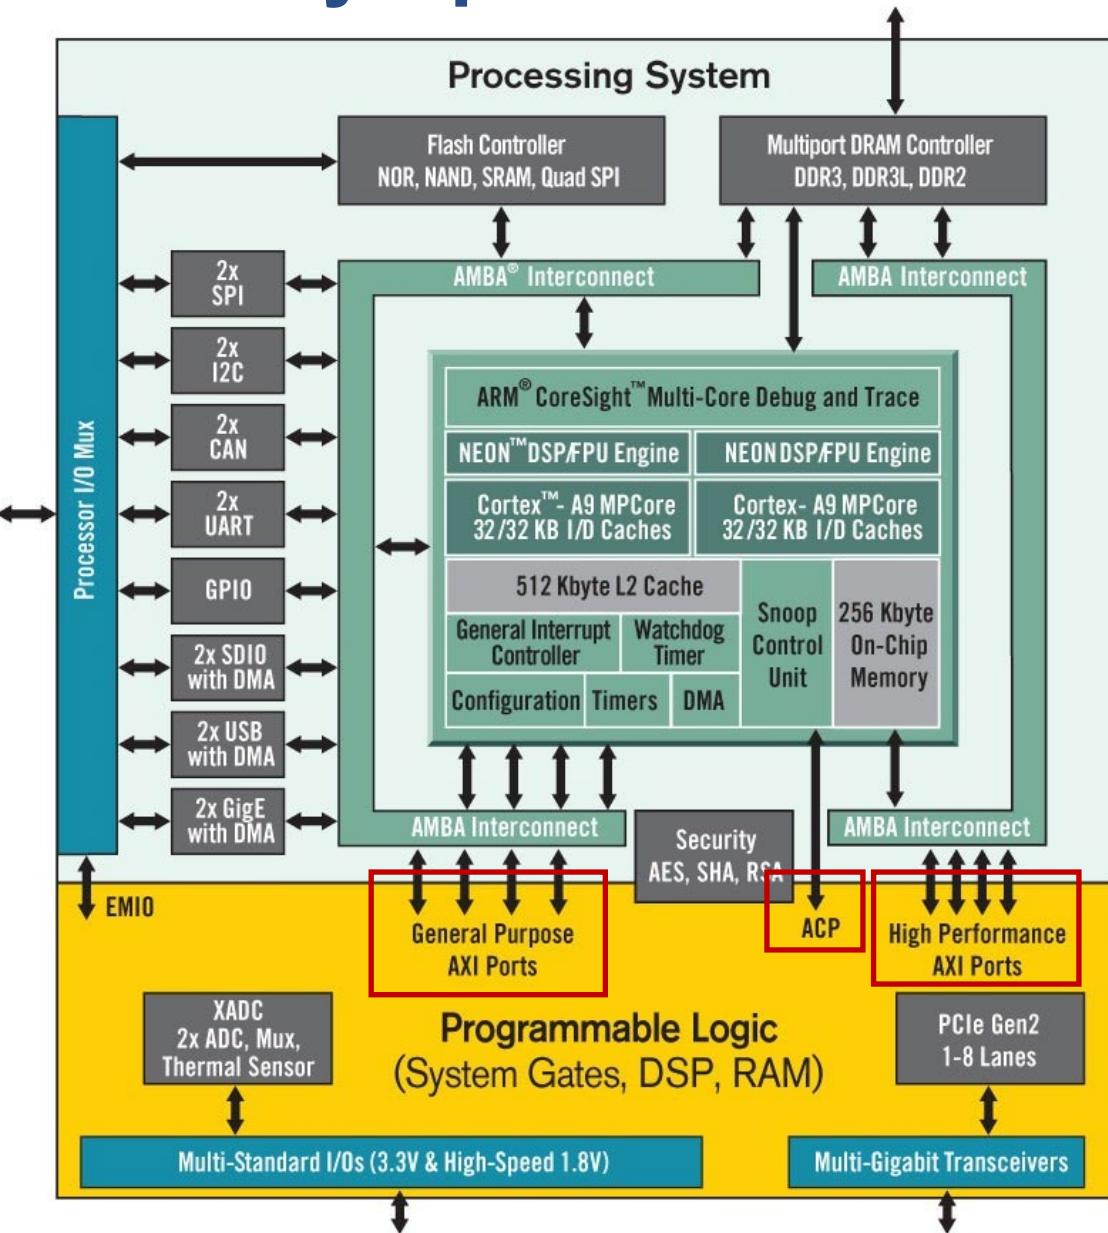
More Info:

<https://docs.xilinx.com/v/u/en-US/ug1037-vivado-axi-reference-guide>



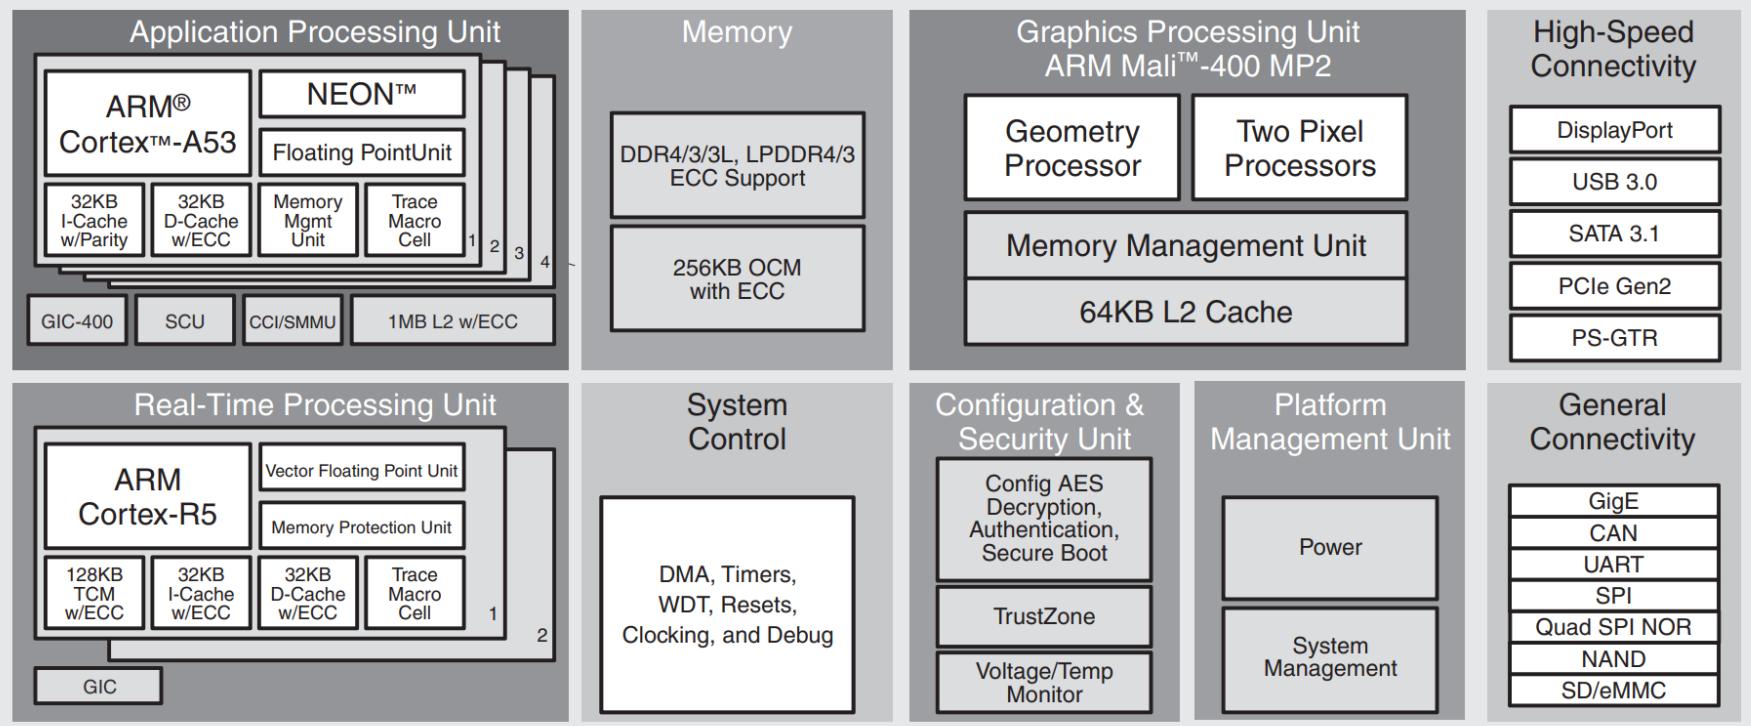
AMD Xilinx Zynq SoC: AXI Interfaces on Zynq

- HP (High Performance): **Mainly for burst data movement**
 - 4x64bit Slave
 - **High bandwidth** access to external memory
 - @150MHz, bandwidth = 9.6GB/s
 - Large Data bursts
- GP (General Purpose): **Mainly for Control and Status**
 - 2x32bit Slave
 - PL to PS peripherals
 - 2x32bit Master
 - PS to PL access
- ACP (Accelerator Coherency Port): **for cache-coherent data access**
 - 1x64bit Slave
 - PL accesses processor L2 cache
 - Hardware coherence

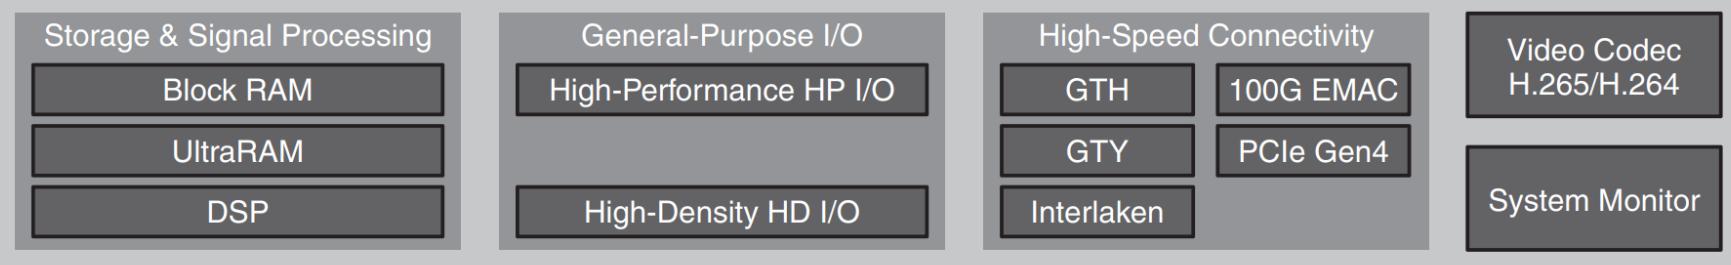


AMD Xilinx Zynq UltraScale+ MPSoC

Zynq UltraScale+ MPSoC Processing System



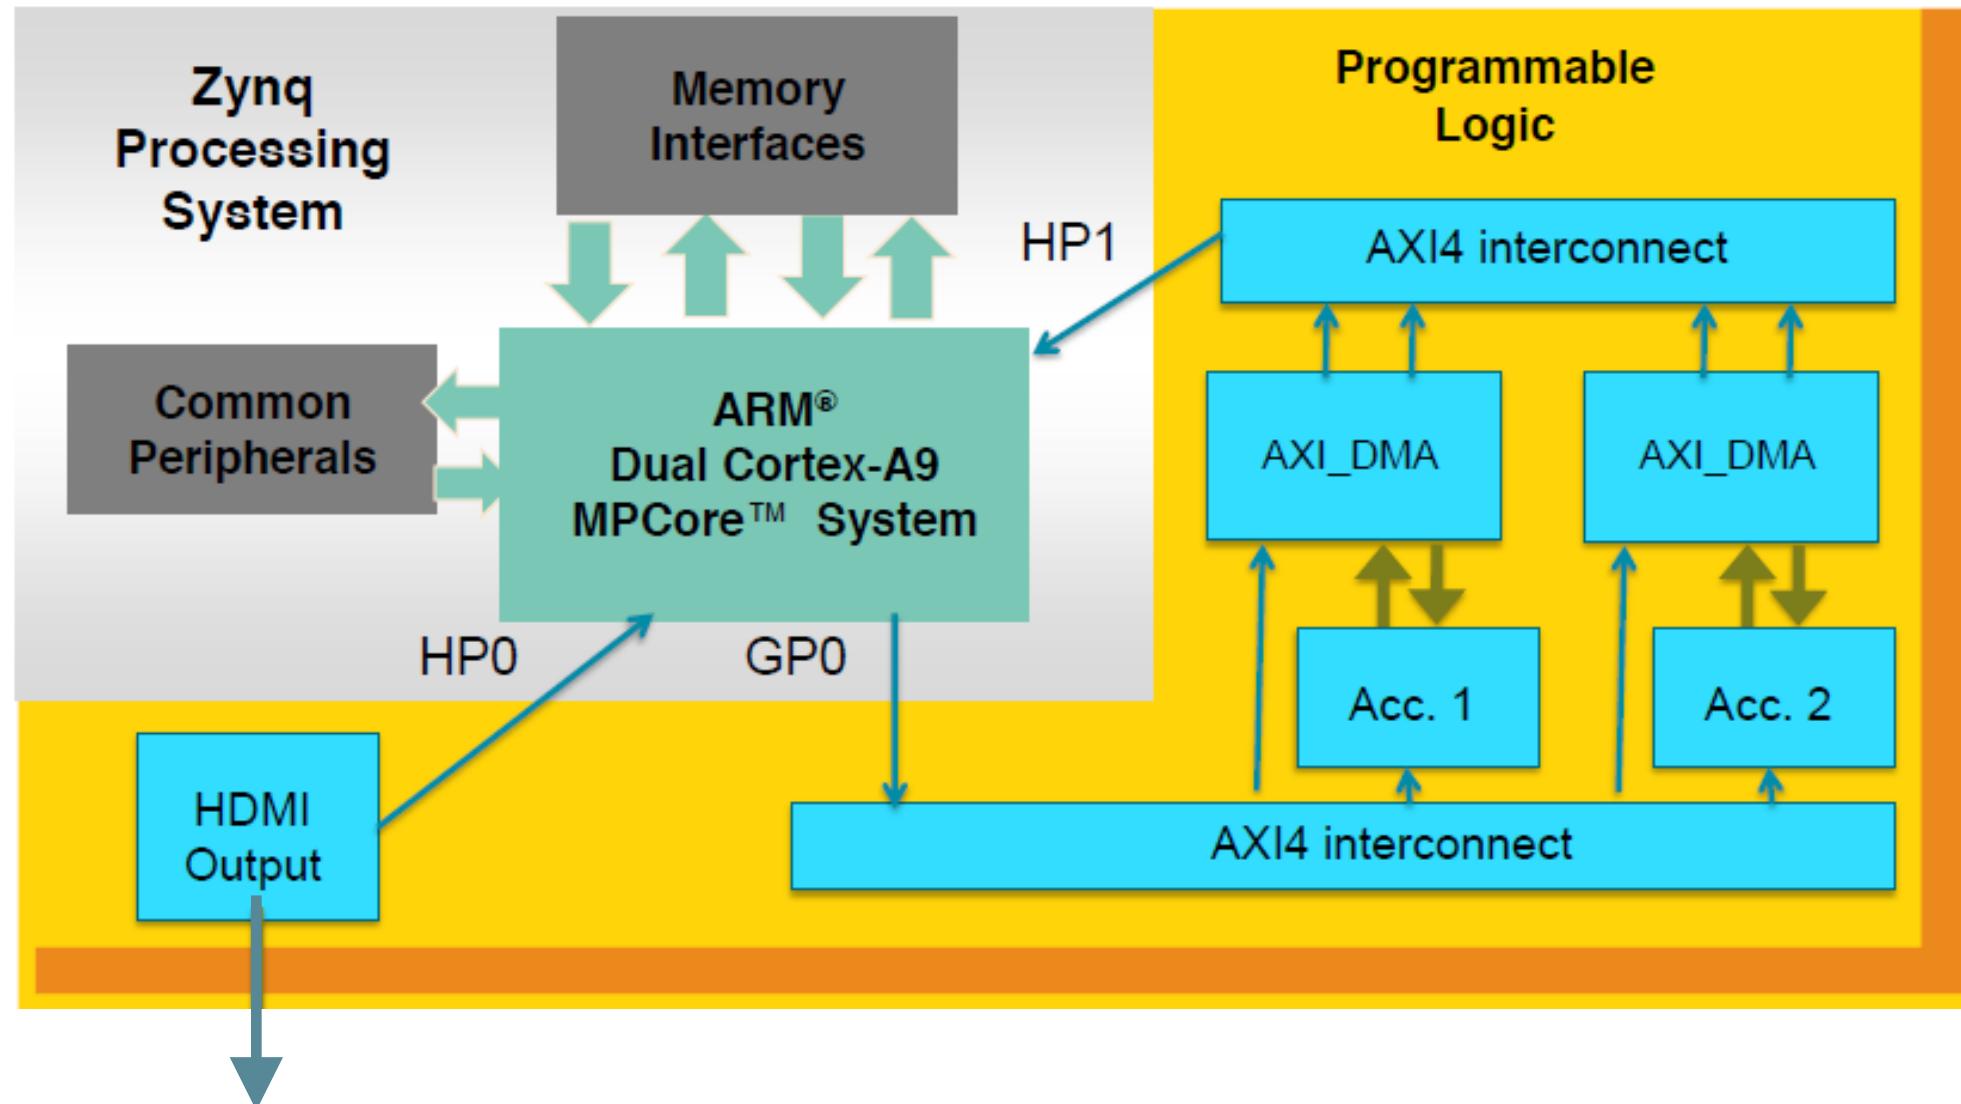
Zynq UltraScale+ MPSoC Programmable Logic



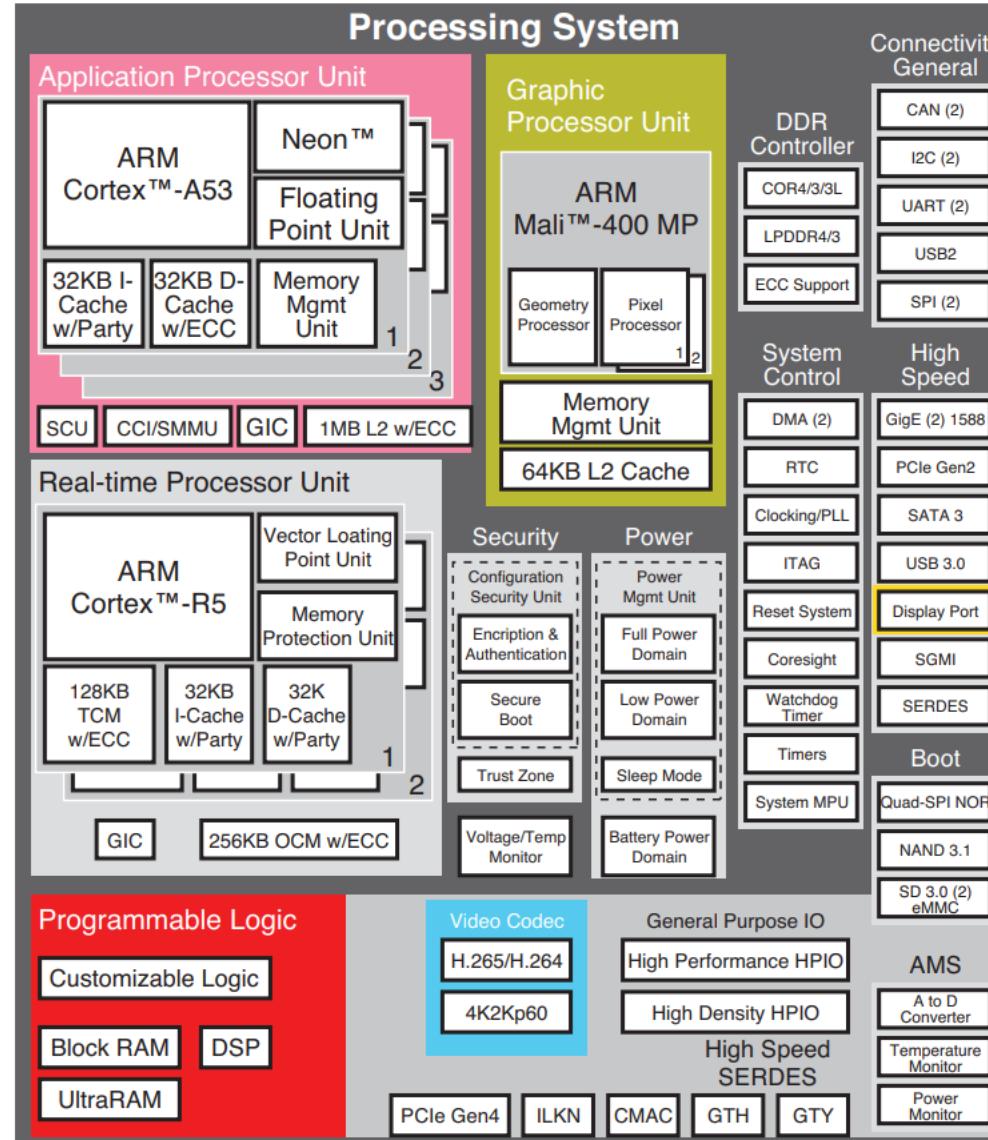
Zynq UltraScale+ EV

- Quad Arm Cortex-A53
- Dual Arm Cortex-R5F
- 16nm FinFET+ Programmable Logic
- Arm Mali-400MP2
- H.264/H.265 Video Codec

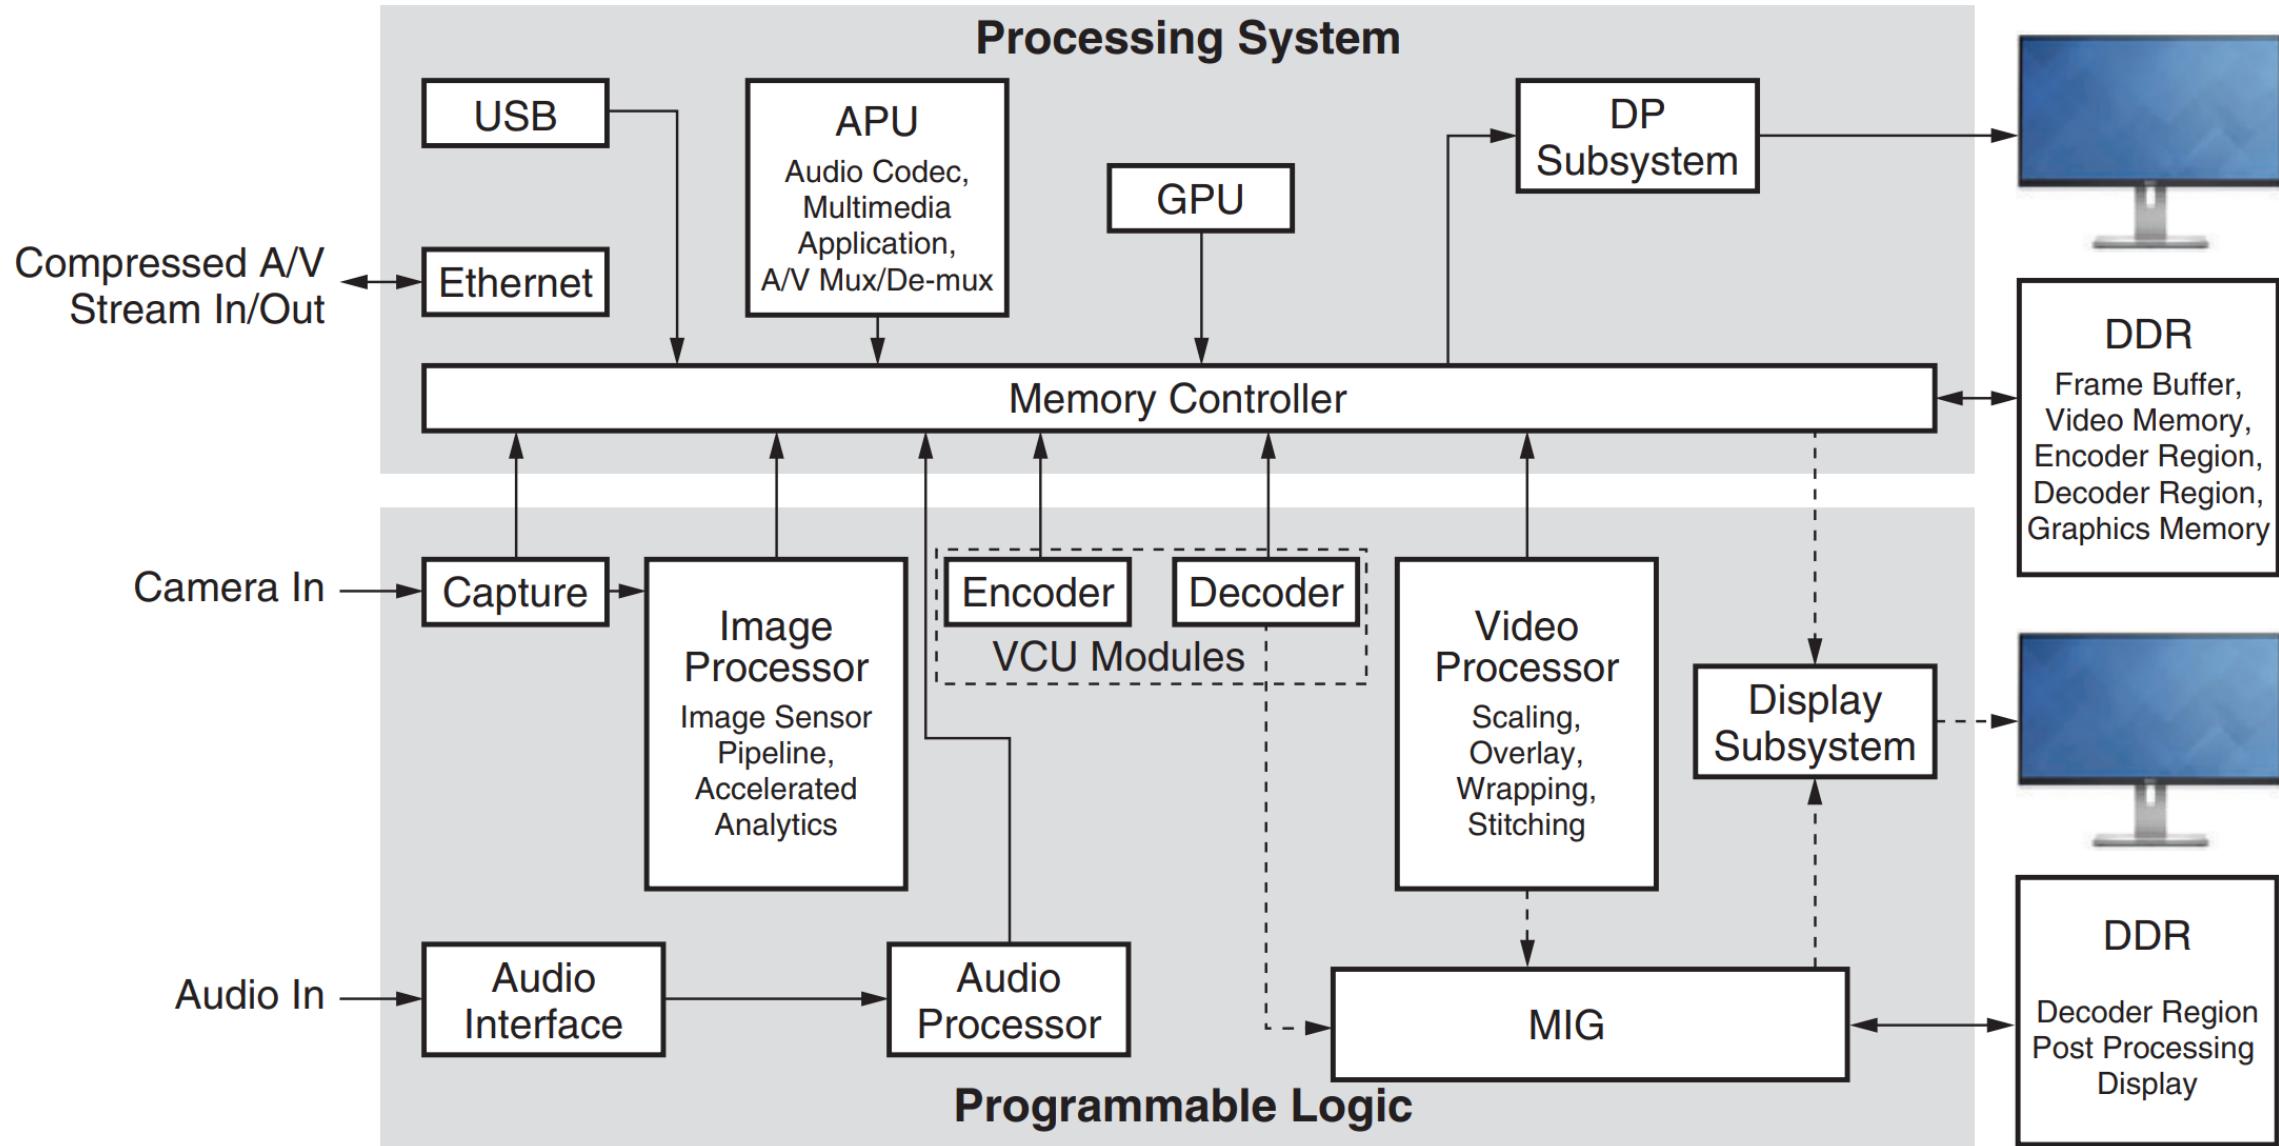
AMD Xilinx Zynq SoC: Example



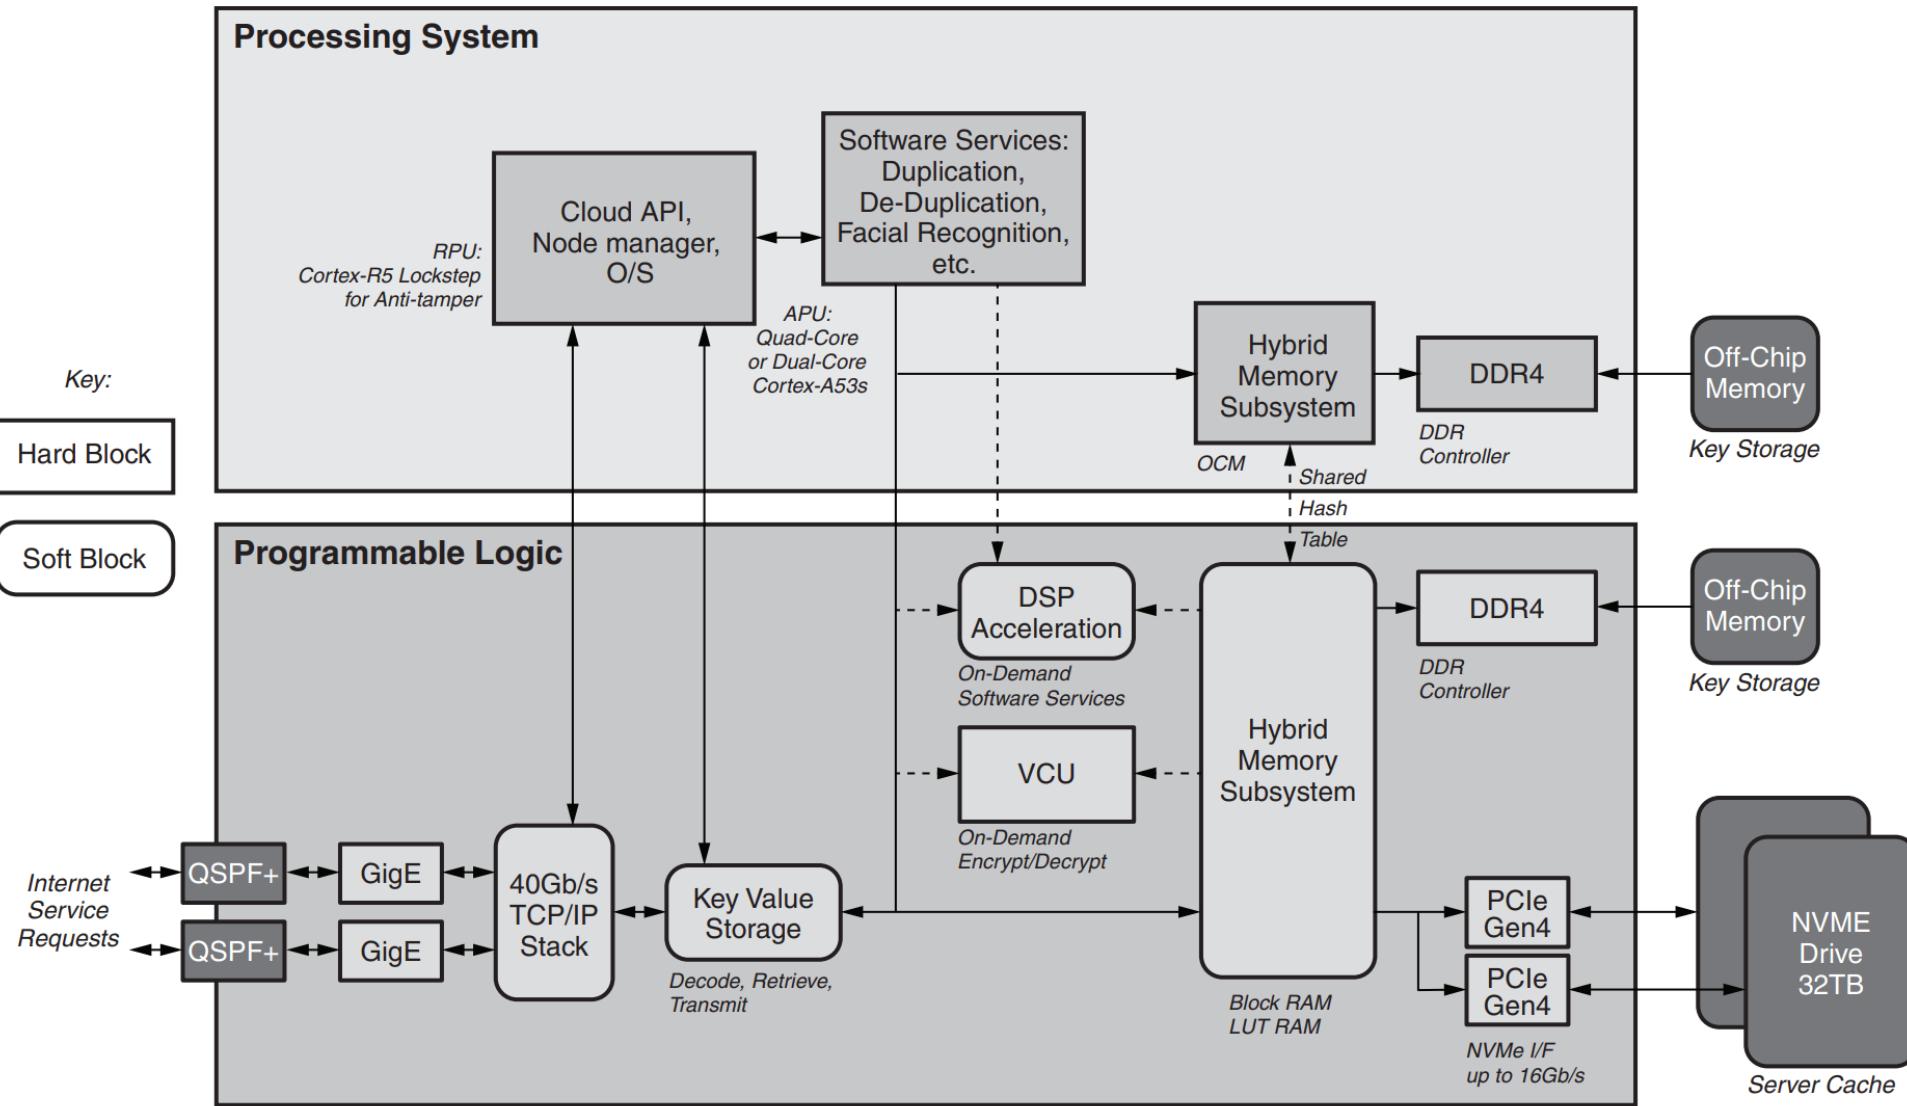
Example: Video Conferencing Application



Example: Video Conferencing Application



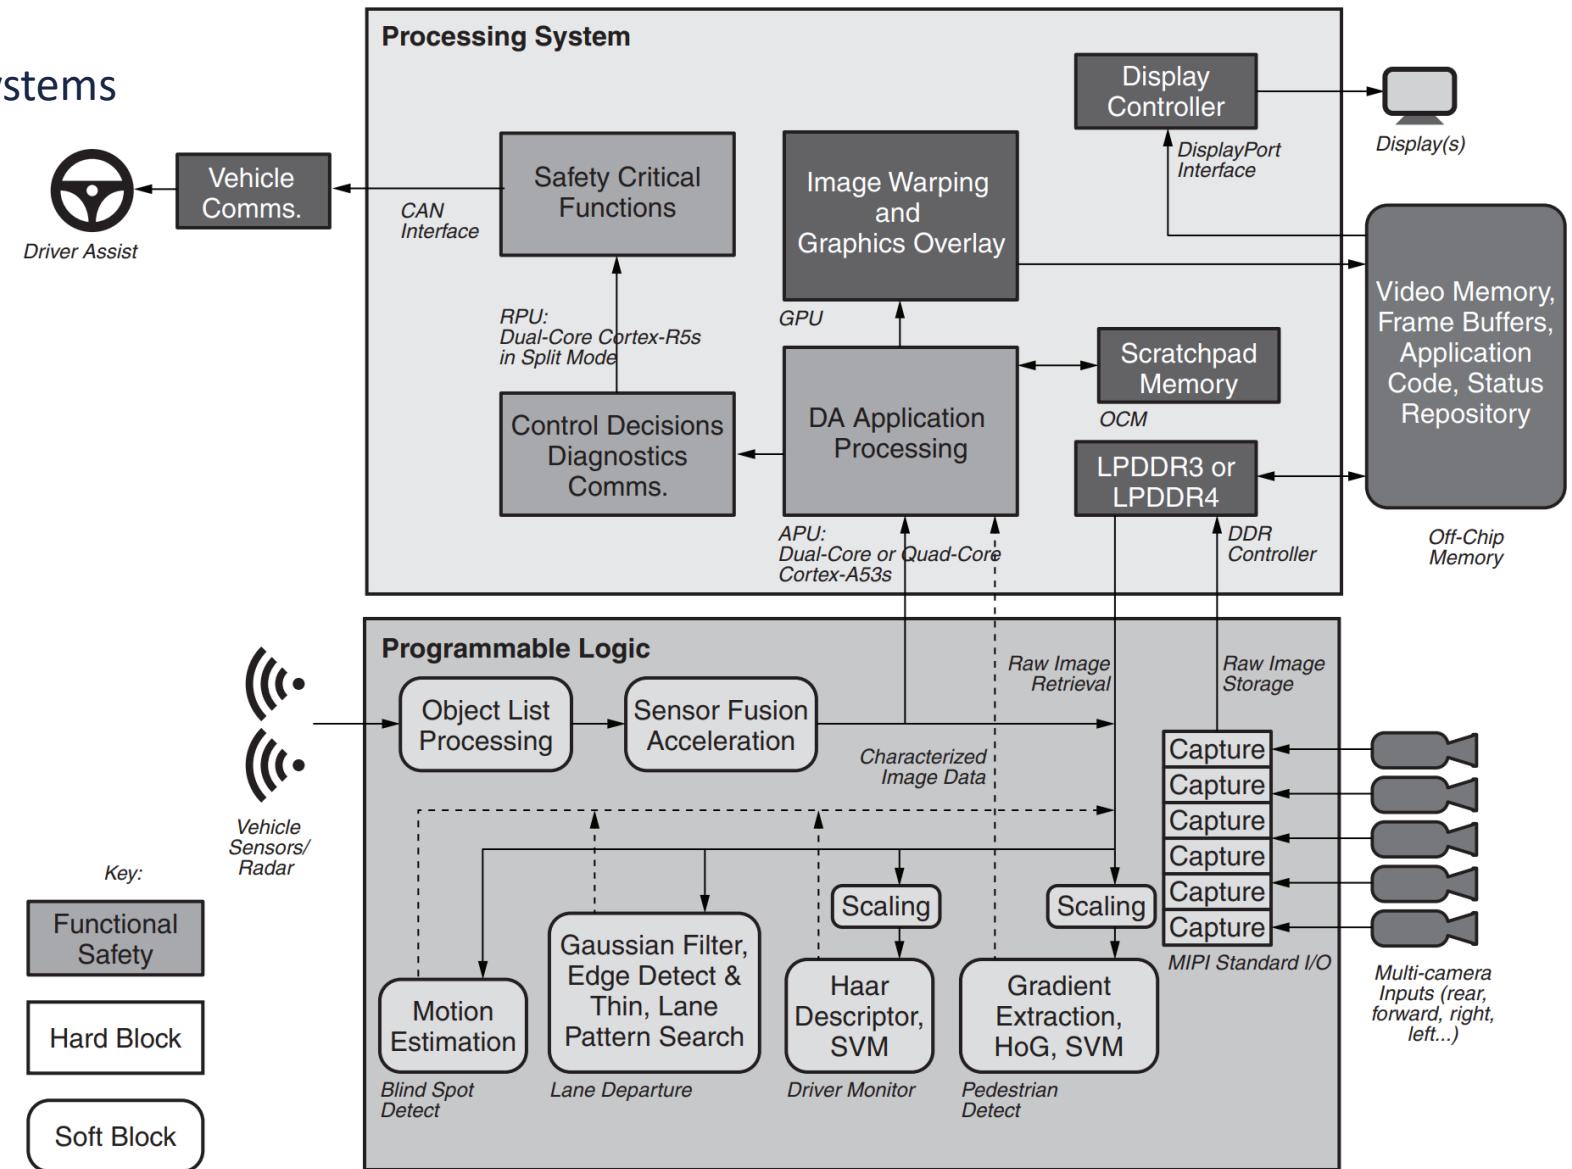
Example: Data Center Networked Storage/Service Platform



WP470_03_050916

Example: Central ADAS Module

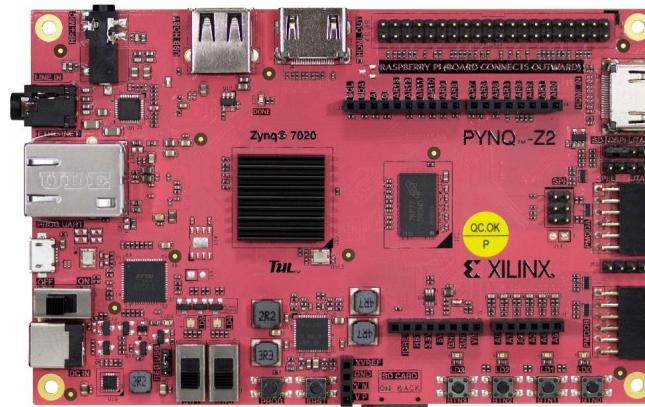
ADAS: Advanced Driver Assistance Systems



WP470_04_050916

PYNQ Boards

- PYNQ-Z2 contains a Zynq-7000 SoC
 - Zynq SoC: ZYNQ XC7Z020-1CLG400C
 - 650MHz dual-core Cortex-A9 processor
 - Programmable logic equivalent to Artix-7 FPGA
 - 13,300 logic slices
 - 630 KB of fast block RAM
 - 220 DSP slices
- Memory
 - 512MB DDR3 with 16-bit bus @ 1050Mbps
 - 16MB Quad-SPI Flash
 - MicroSD Slot
- USB and Ethernet
 - Gigabit Ethernet PHY
 - Micro USB-JTAG Programming circuitry
 - Micro USB-UART bridge
 - USB 2.0 OTG PHY (supports host only)
- Audio and Video (input/output)
- Switches, Push-buttons and LEDs
- Expansion Connectors
 - Two standard Pmod ports
 - 16 total FPGA I/O
 - Arduino connector
 - Raspberry Pi connector



Supports PYNQ programming environment

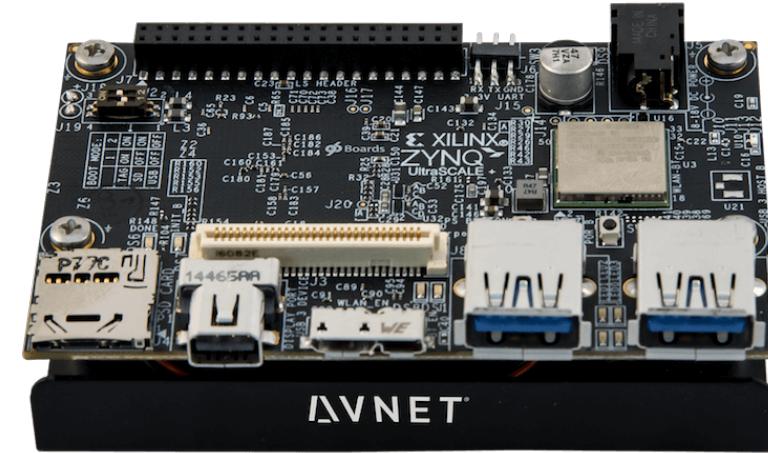
PYNQ-Z2 specs: https://www.tul.com.tw/images/PYNQ-Z2_PA_v2_pp_20201209_STD.pdf

PYNQ-Z2 user manual: https://dpoauwgwqsy2x.cloudfront.net/Download/PYNQ_Z2_User_Manual_v1.1.pdf

PYNQ-Z2 setup guide: https://pynq.readthedocs.io/en/latest/getting_started/pynq_z2_setup.html

Ultra96 Boards

- Ultra96 contains a Zynq UltraScale+ MPSoC
 - MPSoC: Xilinx Zynq UltraScale+ MPSoC ZU3EG A484
 - Memory: Micron 2GB LPDDR4 memory
 - Storage: 16GB microSD card + adaptor
 - Wireless: 802.11b/g/n Wi-Fi and Bluetooth 4.2
 - USB and Ethernet
 - 1x USB 3.0 Type Micro-B upstream port
 - 2x USB 3.0, 1x USB 2.0 Type A downstream ports
 - Display: Mini DisplayPort (MiniDP or mDP)
 - Expansion interface:
 - 40-pin 96Boards Low-speed expansion header
 - 60-pin 96Boards High speed expansion header

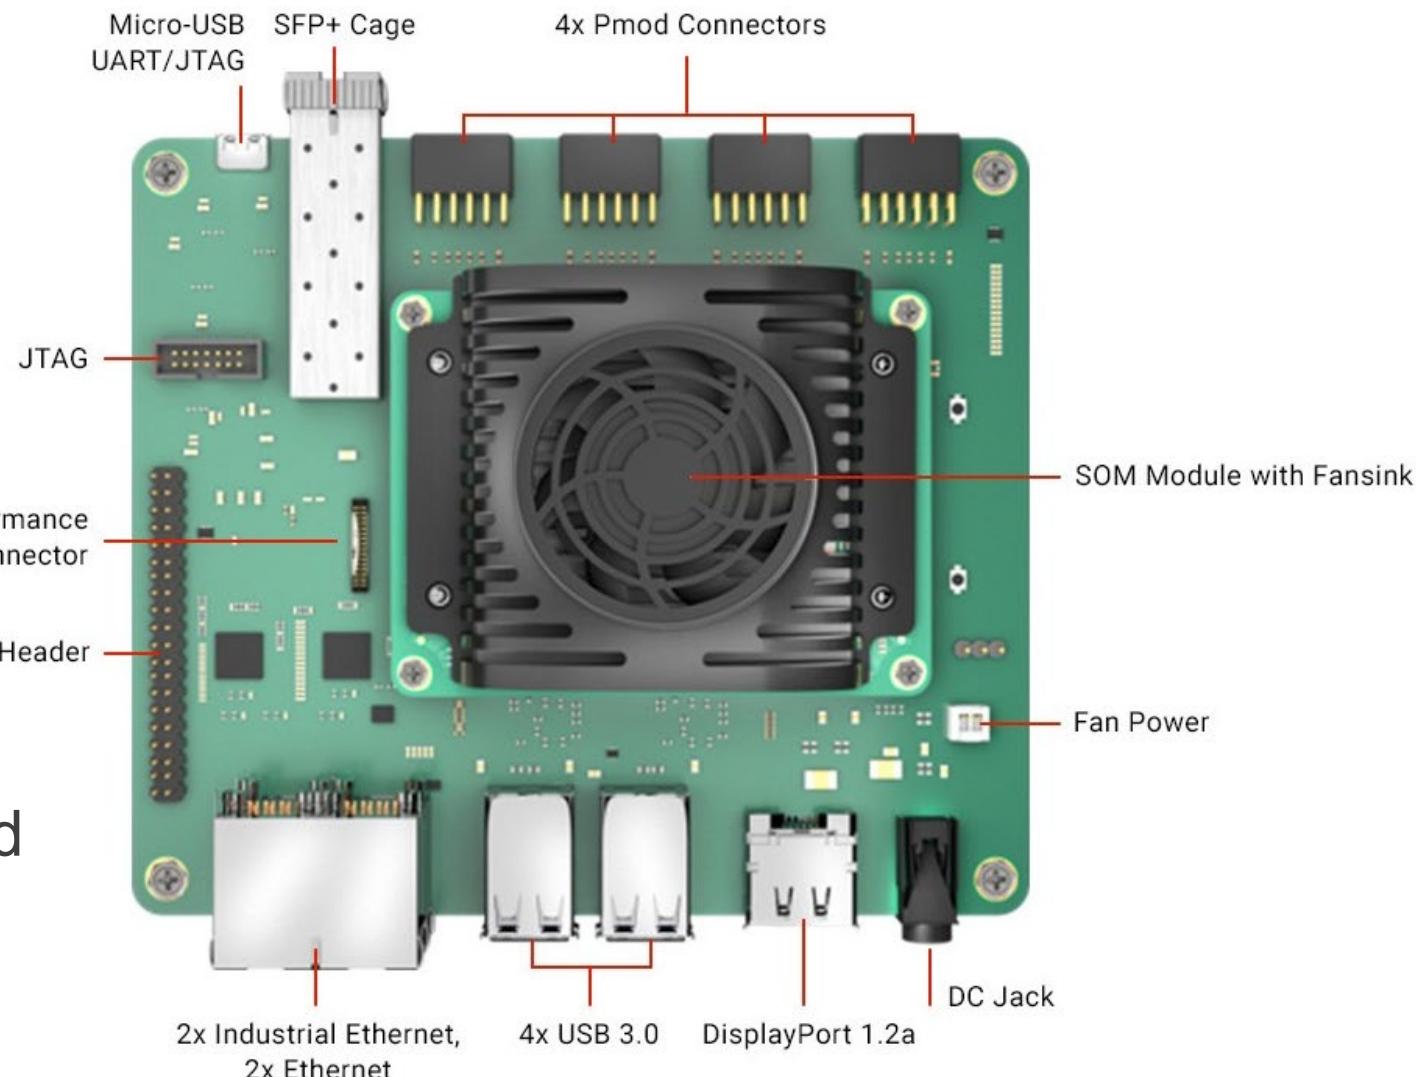


Supports PYNQ programming environment

- Ultra96-v2 user's guide: https://www.avnet.com/wps/wcm/connect/onesite/b85b9556-0b2a-42b3-ad6a-8dcf3eac1ff9/Ultra96-V2-HW-User-Guide-v1_3.pdf?MOD=AJPERES&CACHEID=ROOTWORKSPACE.Z18_NA5A1I41L0ICD0ABNDMDDG0000-b85b9556-0b2a-42b3-ad6a-8dcf3eac1ff9-nDBi1I6
- Ultra96-v2 PYNQ setup guide: https://ultra96-pynq.readthedocs.io/en/latest/getting_started.html

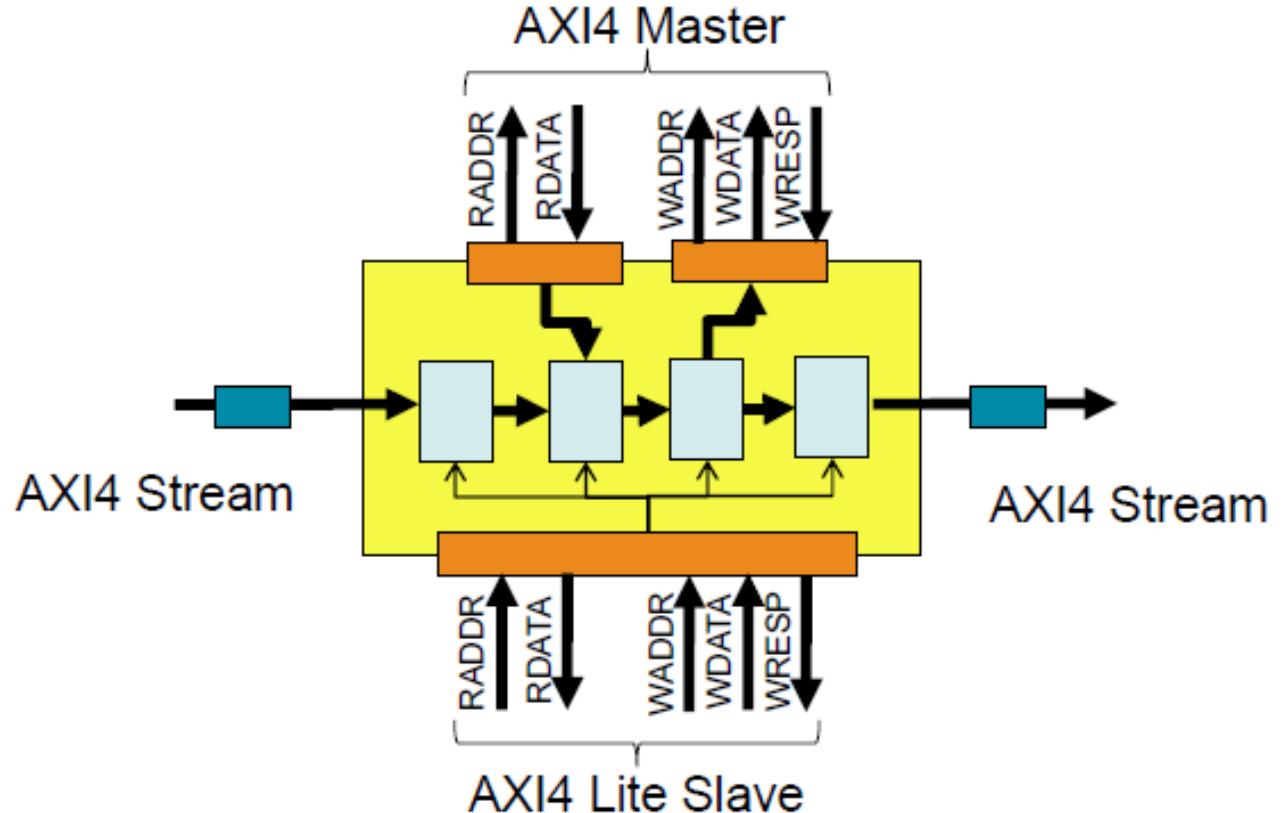
AMD Kria KR260 Robotics Starter Kit

- Device: Zynq™ UltraScale+™ MPSoC EV (XCK26)
- System logic cells: 256K
- 36Kb Block RAM blocks: 144
- 288Kb UltraRAM blocks: 64
- DSP slices: 1.2K
- DDR memory: 4GB DDR4
- Primary boot memory: 512 Mb QSPI
- Secondary boot memory: SDHC card



AMD Xilinx Zynq SoC: Accelerator Development Example

- Example:
 - Design a vector add accelerator to perform $c[i] = (a[i] + b[i]) \quad (0 \leq i \leq N)$ for given vectors
 - Variable size arrays
 - Performance and power
 - Arrays to initialized in software
 - N is constant
- Choose Interface:
 - How should I move data to FPGA?
 - Stream or memory-mapped?
- Latency calculations
 - Time to process/move data
- Resource Usage
 - How much parallelism is available?
- On-FPGA memory
 - Can all vectors be stored on FPGA?
- HW-SW co-design



AMD Xilinx Zynq SoC: Accelerator Development Example

- Simplified Example (sequential Vector Add)

```
for(i = 0; i < N; i++)  
    C[i] = A[i] + B[i];
```

- Basic Performance modeling/estimation:

- Break down into smaller operations
- Compute time per operation
- E.g., read latency = r , Write latency = w , floating point add latency = c

- Vector Add
 - Repeat N times

- LOAD A[i]
- LOAD B[i]
- ADD C[i], A[i], B[i]
- STORE C[i]

$$\text{Total time} = N(2r + c + w)$$

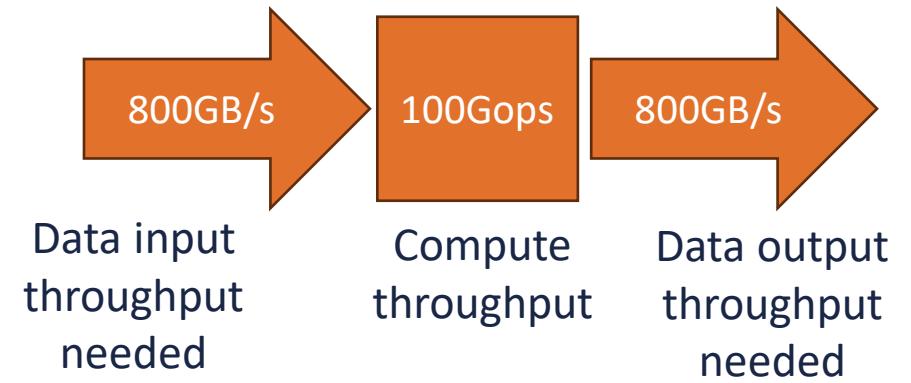
AMD Xilinx Zynq SoC: Accelerator Development Example

$$\text{Total time} = N(2r + c + w)$$

- What does this tell us?
 - How large does N need to be before it makes sense to accelerate this?
 - For each compute operation, 12 bytes of data is moved
 - How do I bring data in?
 - MEMORY < --- > PS < --- > PL
 - MEMORY < --- > PL
- Bandwidth vs Latency:
 - Bandwidth:
How much data can you bring in per unit time.
 - Latency:
How long does it take data to arrive.
- To perform K operations in parallel, you need K*12bytes
 - B/W needed = K*12*Frequency
 - Choose K wisely

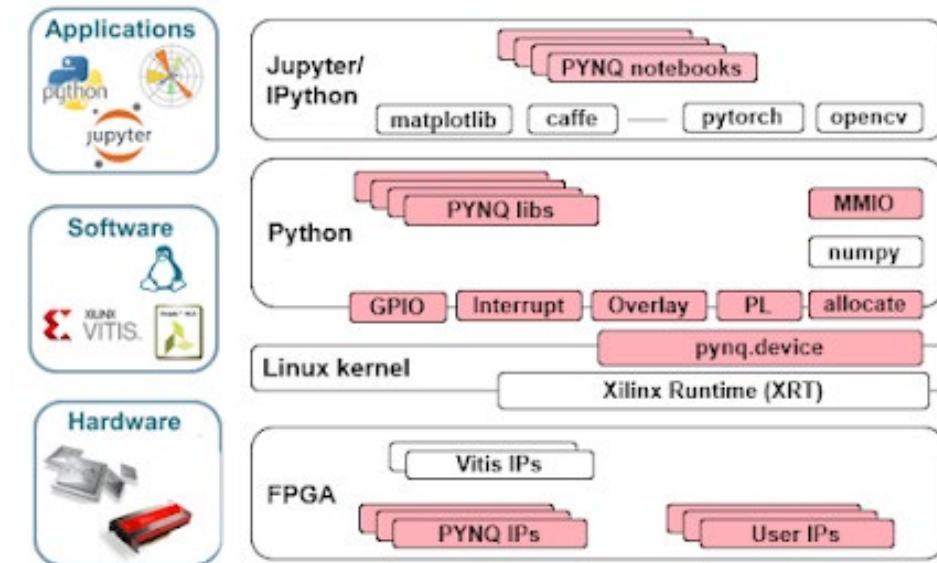
Zynq SoC: Accelerator Development Example

- Is 100 GOPS really possible?
 - Yes and No
- Assumption: Each op is floating point
 - 8 bytes per operation
 - Total bandwidth needed 800 Gb/s?
- Look to Data Reuse
 - On-Chip memory can be a limit
- Streaming/Systolic Design:
 - Smaller blocks, feeding into each other
- Be smart about bandwidth
 - Dedicated read and dedicated write channels may not be smart
 - Think about communication patterns
 - Do you need concurrent read and write?



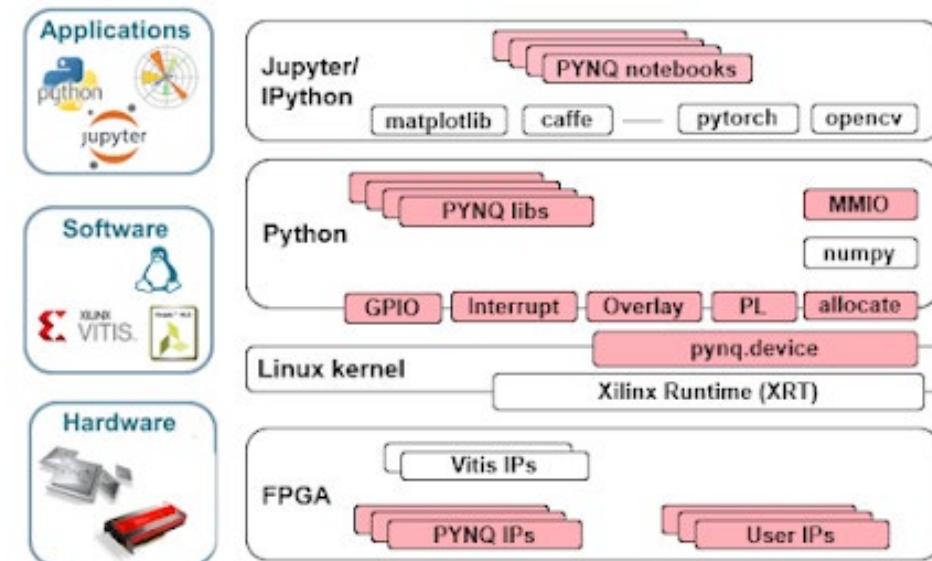
What is PYNQ?

- An open-source project from Xilinx® that makes it easier to use Xilinx platforms
- **Python**-based APIs and libraries
- Simplifies **host** programming
- Supports wide range of Xilinx devices:
 - Zynq, Zynq UltraScale+, Zynq RFSoC, MPSoC, Alveo, AWS-F1, etc.



What is PYNQ?

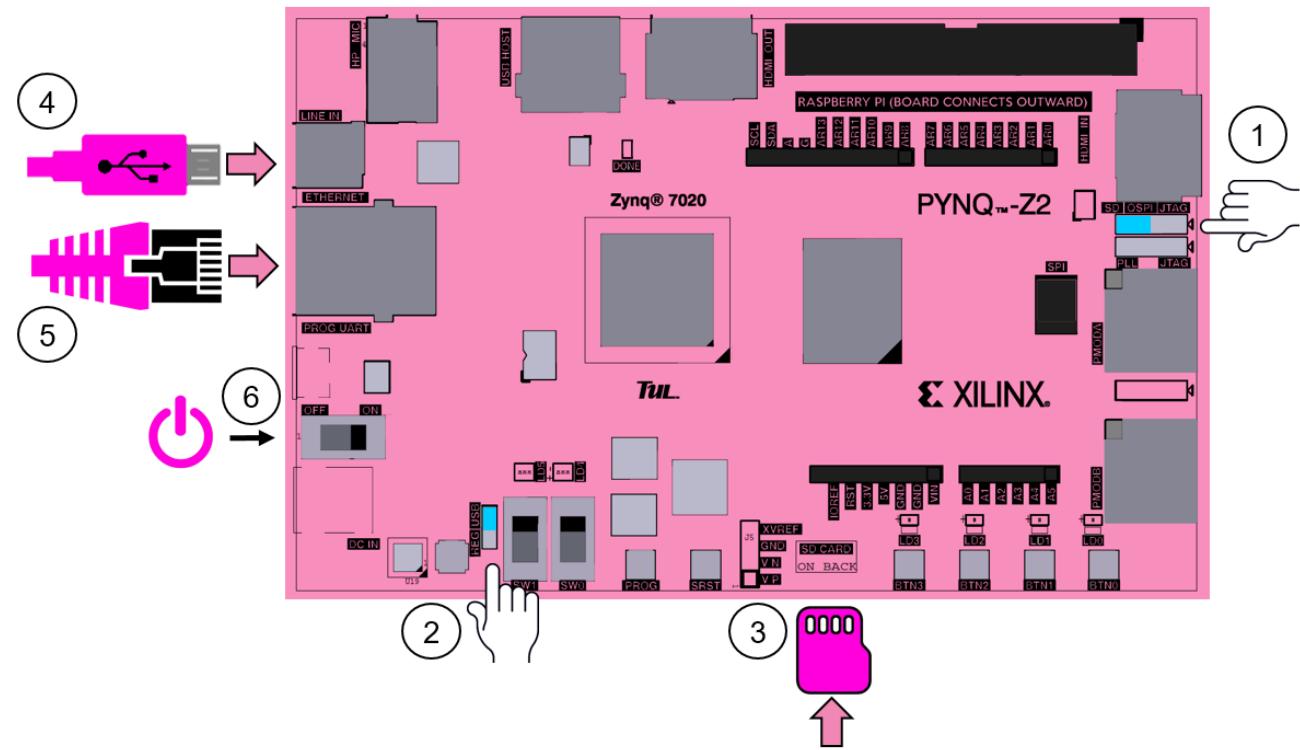
- Jupyter Notebook: a browser based interactive computing environment
- PYNQ enabled FPGA boards can be programmed in Jupyter notebook using Python (host programming)
- PYNQ is delivered in two forms:
 - Bootable Linux image for Zynq boards
 - Open-source Python package for Alveo and AWS-F1



Getting Started with PYNQ

Using PYNQ-Z2 board as example

1. Set the **Boot** jumper to the **SD** position (boot from the Micro-SD card)
2. To power the board from the micro-USB cable, set the **Power** jumper to the **USB** position. (You can also power the board from an external 12V power regulator by setting the jumper to **REG**.)
3. Insert the Micro SD card loaded with the PYNQ-Z2 image into the **Micro SD** card slot underneath the board
4. Connect the USB cable to your PC/Laptop, and to the **PROG - UART** Micro-USB port on the board
5. Connect the Ethernet port by following the instructions below
6. Turn on the PYNQ-Z2

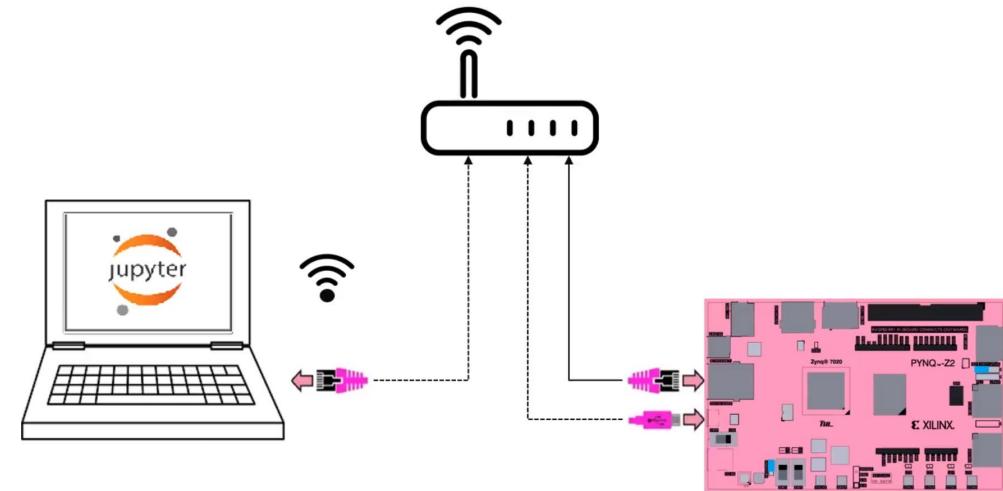


Getting Started with PYNQ

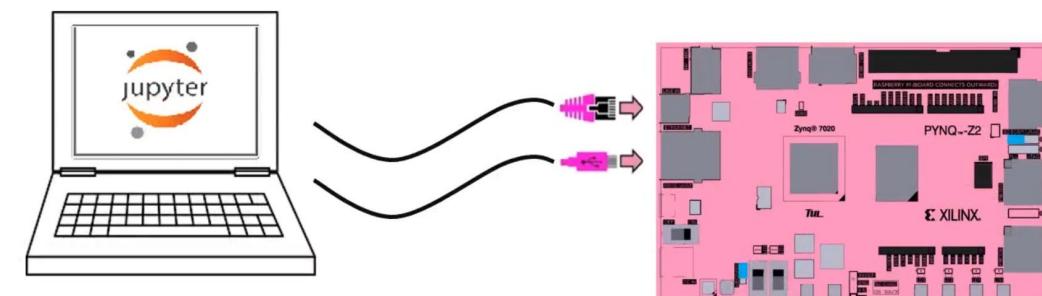
Network connection

Once your board is setup, you need to connect to it to start using Jupyter notebook

- Option A: Connect to a Network Router
 - Connect the Ethernet port on your board to a router/switch
 - Connect your computer to Ethernet or Wi-Fi on the router/switch
 - Browse to <http://<board IP address>>
- Option B: Connect to a Computer
 - Assign your computer a static IP address
 - Connect the board to your computer's Ethernet port
 - Browse to <http://192.168.2.99>



Option A: Connect to a Network Router



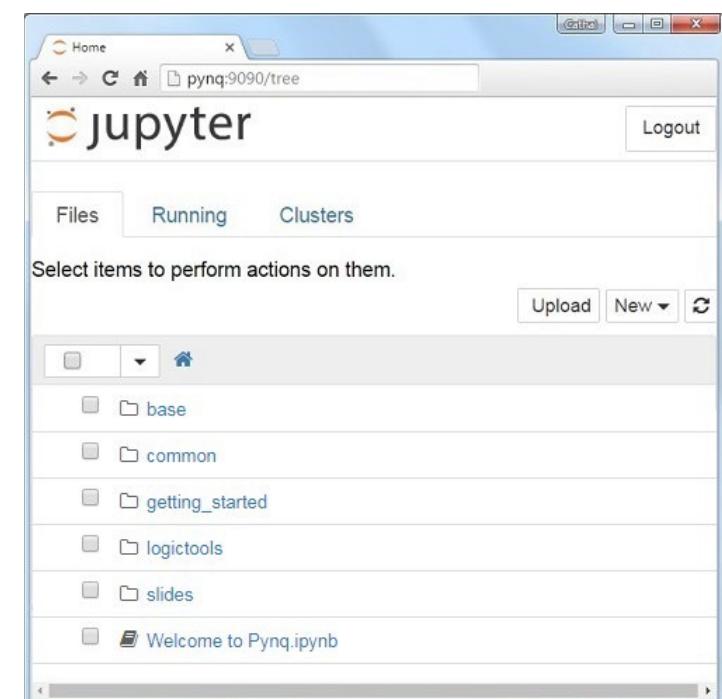
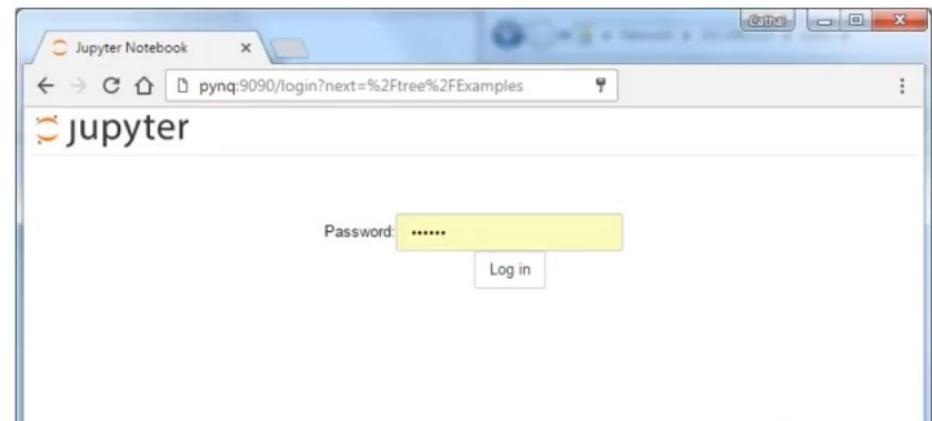
Option B: Connect to a Computer

Getting Started with PYNQ

Connecting to Jupyter portal

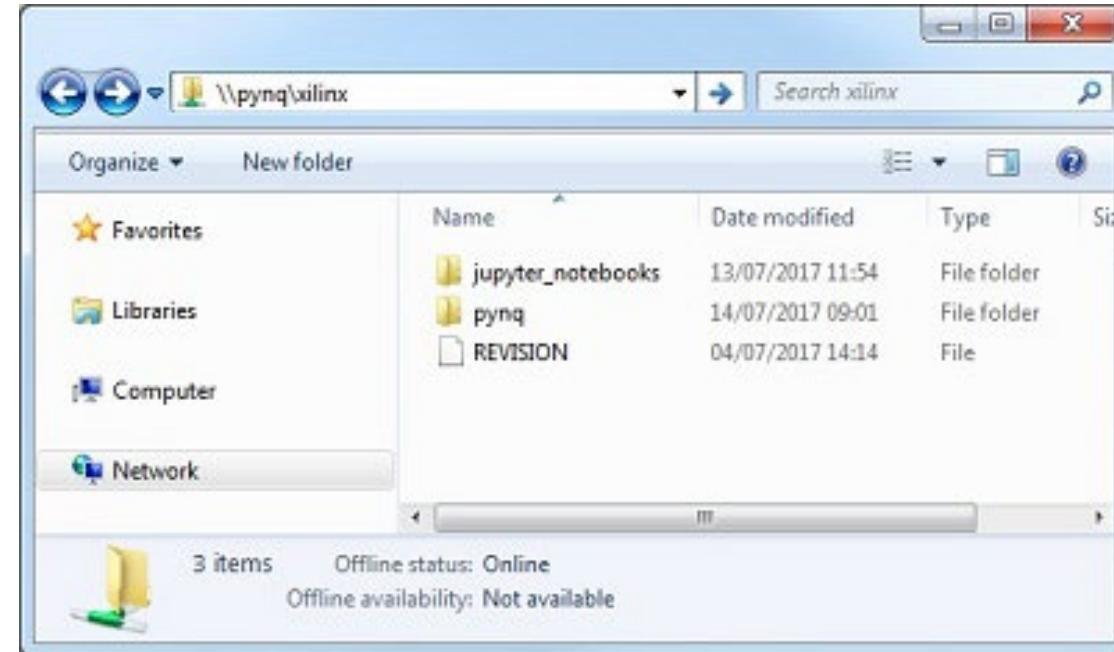
After network connection is established, browse to board's corresponding IP address, then you can start to use the board using Jupyter notebook.

- For boards connected to network:
 - Browse to <http://pynq:9090>
- For boards connected to computer:
 - Browse to <http://192.168.2.99:9090>
- Default password: xilinx



Getting Started with PYNQ

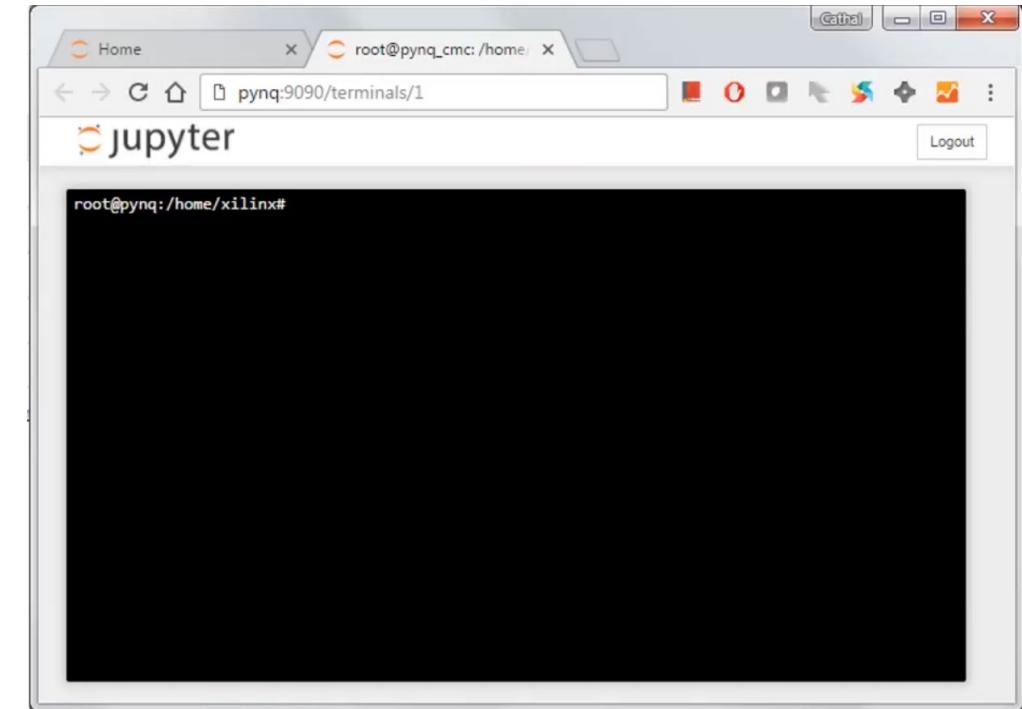
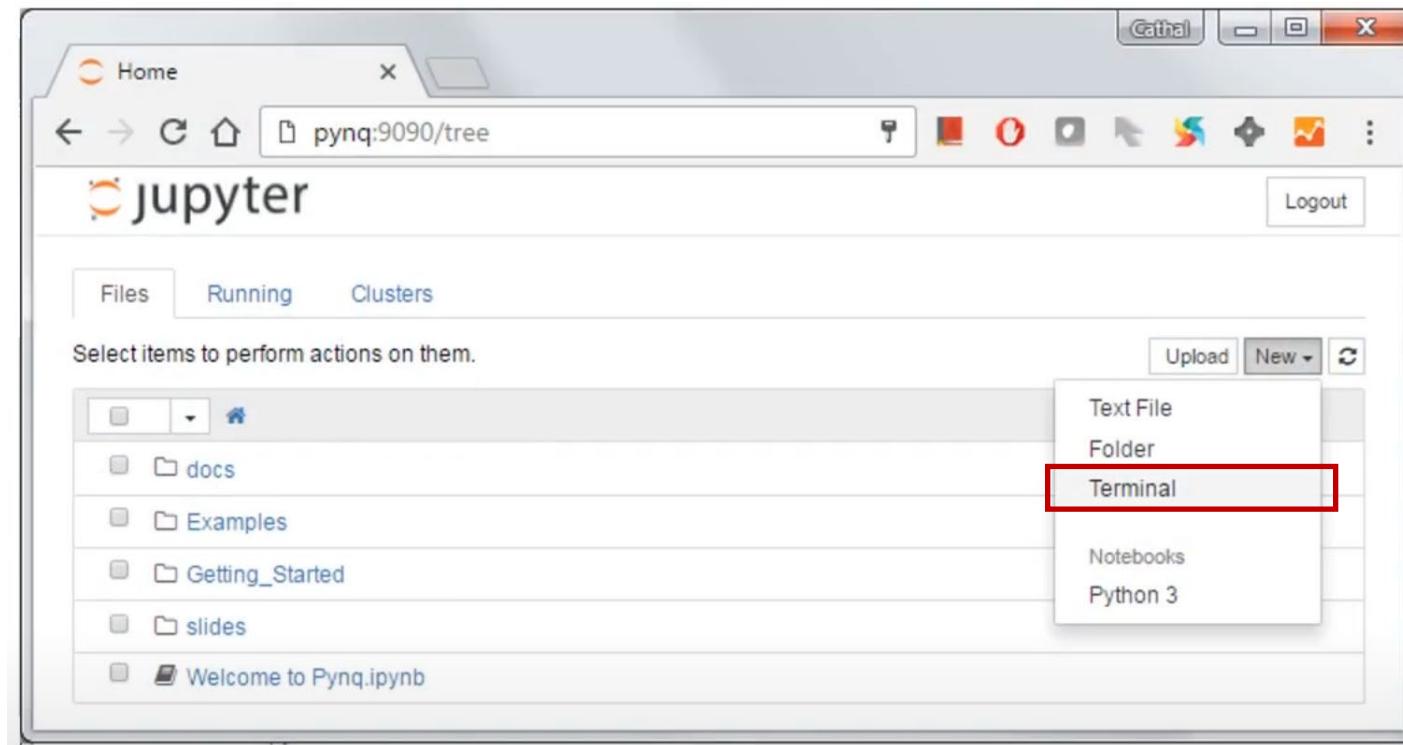
Connect via Samba



- Windows: <\\pynq\xilinx>
- Mac or Linux: <smb://pynq/xilinx>

Getting Started with PYNQ

Jupyter terminal



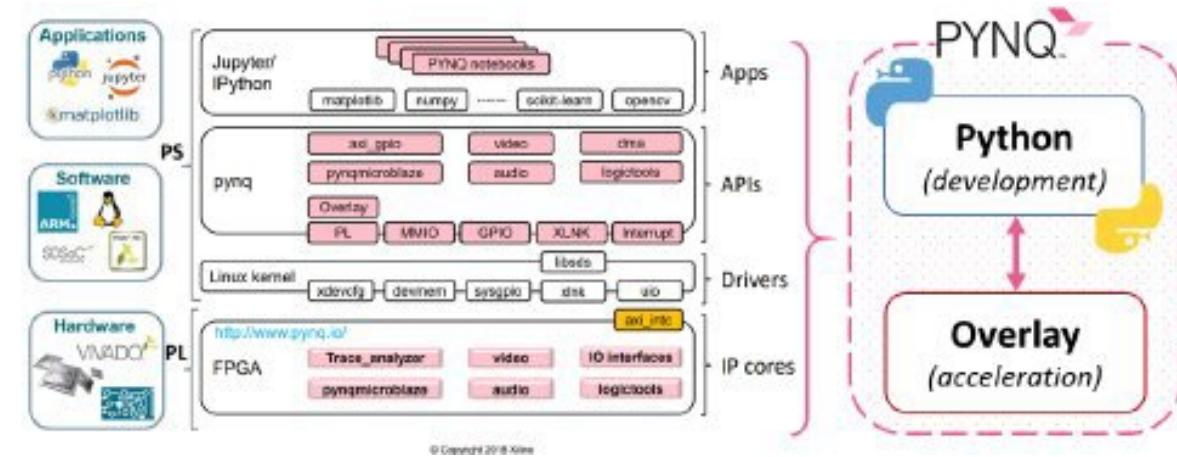
The PS runs a full Linux OS, so you can also *ssh* to the system:

ssh pynq@192.168.2.99

PYNQ Overlays

What are overlays?

- “hardware libraries”, where IP instances become objects
- programmable/configurable FPGA designs
- Can be used in a similar way to a software library to run some functions on the FPGA fabric
- Can be loaded to FPGA dynamically, just like a software library



PYNQ Overlays:

- Provide a Python interface for controlling PL from Python running in the PS
- Created by hardware designers and wrapped with PYNQ Python API
- Each entry (*IPs* or *ports*) in the hardware becomes an **object** and has specific attributes and methods (e.g. LED IO ports can be accessed with `overlay.leds`)

An overlay usually includes:

- A *bitstream* to configure FPGA fabric
- A Vivado design *Tcl file* to determine the available IPs
- Python API that exposes the IPs as attributes (PYNQ library)

Roughly, overlay is:

bitstream + block design structure file + APIs

PYNQ Overlays

Loading an Overlay:

- The PYNQ **Overlay** class can be used to load an overlay

```
from pynq import Overlay  
overlay = Overlay("base.bit")
```

- An overlay can be instantiated by specifying the bitstream file
- Overlay instantiation also downloads the bitstream to FPGA

Inspecting an Overlay:

- Once overlay is instantiated, **help()** method can be used to discover what is in an overlay

```
help(overlay)
```

- help()** can also be used to get more information about a specific object in the overlay

```
help(overlay.leds)
```

- An example output from calling **help(base_overlay)**:

```
Help on BaseOverlay in module pynq.overlays.base object:
```

```
class BaseOverlay(pynq.overlay.Overlay)  
    The Base overlay for the Pynq-Z1  
  
    This overlay is designed to interact with all of the on board peripherals  
    and external interfaces of the Pynq-Z1 board. It exposes the following  
    attributes:  
  
Attributes  
-----  
leds : AxiGPIO  
    4-bit output GPIO for interacting with the green LEDs LD0-3  
buttons : AxiGPIO  
    4-bit input GPIO for interacting with the buttons BTN0-3  
switches : AxiGPIO  
    2-bit input GPIO for interacting with the switches SW0 and SW1  
rgbleds : [pynq.board.RGBLED]  
    Wrapper for GPIO for LD4 and LD5 multicolour LEDs  
video : pynq.lib.video.HDMIWrapper  
    HDMI input and output interfaces  
audio : pynq.lib.audio.Audio  
    Headphone jack and on-board microphone
```

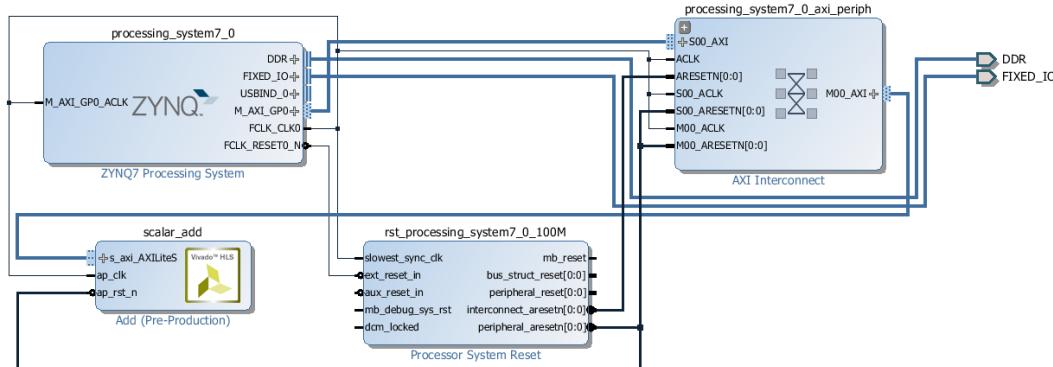
- An API can be used to control the object. For example, turning on LED0 on the board:

```
base_overlay.leds[0].toggle()
```

PYNQ Overlays

Example of Creating and Using Overlay

Assume we create an IP `scalar_add`, and create a block diagram (BD):



This block diagram consists of the IP and glue logic to connect to PS IP.

After synthesis, we get the *bitstream file* (say, `hw.bit`) and the *BD structure file* (`tcl` file or `hwh` file).

With these two files, we can wrap them with PYNQ Overlay class to create a PYNQ overlay (both files should be in the same directory):

```
from pynq import Overlay
overlay = Overlay('hw.bit')
```

Now we get the PYNQ overlay object “`overlay`”.

(note that we get this overlay by bitstream + BD structure + PYNQ API)

Creating the overlay will automatically download the bitstream to FPGA

After creating the overlay, we can inspect the overlay. In Jupyter notebook, we can use a question mark to find out what is inside:

```
overlay?
```

All the entries in the overlay are accessible via attributes on the overlay class. For example, we can access `scalar_add` IP:

```
add_ip = overlay.scalar_add
```

We can also expose the register map associated with IP:

```
add_ip.register_map
```

It prints:

```
RegisterMap {
    a = Register(a=0),
    b = Register(b=0),
    c = Register(c=0),
    c_ctrl = Register(c_ap_vld=1, RESERVED=0)
}
```

We can also interact with the IP using the register map:

```
add_ip.register_map.a = 3
add_ip.register_map.b = 4
Print(add_ip.register_map.c)
```

Alternatively, by reading the driver source code generated by HLS we can determine the offsets we need to write the two arguments (they are all in the same memory space):

```
add_ip.write(0x10, 4)
add_ip.write(0x18, 5)
add_ip.read(0x20)
```

PYNQ Overlays – Prepare Input/Output Buffers

Allocate

- The `pynq.allocate` function is used to allocate memory that will be used by IP in the PL
- `pynq.allocate` function returns a `pynq.Buffer` object that is a sub-class of NumPy's `ndarray` with additional properties and methods suited for use with the programmable logic
 - `device_address` is the address that should be passed to the programmable logic to access the buffer
 - `coherent` is True if the buffer is cache-coherent between the PS and PL
 - `flush` flushes a non-coherent or mirrored buffer ensuring that any changes by the PS are visible to the PL
 - `invalidate` invalidates a non-coherent or mirrored buffer ensuring any changes by the PL are visible to the PS
 - `sync_to_device` is an alias to `flush`
 - `sync_from_device` is an alias to `invalidate`

Example: Create a contiguous array of 5 32-bit unsigned integers

```
from pynq import allocate
input_buffer = allocate(shape=(5,), dtype='u4')
input_buffer[:] = range(5)
input_buffer.flush()
```

PYNQ Overlays – Running Accelerators

After creating the overlay and have data buffers ready, we can start the accelerator.

Running Accelerators

Start the kernel synchronously:

```
ol.my_kernel.call(input_buf, output_buf)
```

The call function has the same function signature as the top function in original HLS source code.

Alternatively, start the kernel in a non-blocking way:

```
handle = ol.my_kernel.start(input_buf, output_buf)  
handle.wait()
```

Freeing designs (overlays):

```
ol.free()
```

Execution results can be collected by accessing output buffers.

NOTE: call, start, and wait are newly added since PYNQ 2.5.

For older version of PYNQ, to invoke the accelerator, we need to manually set the *ap_start* bit of the IP, and then wait for the *ap_start* signal.

The address of these bits can be found from the driver file generated by Vivado HLS.

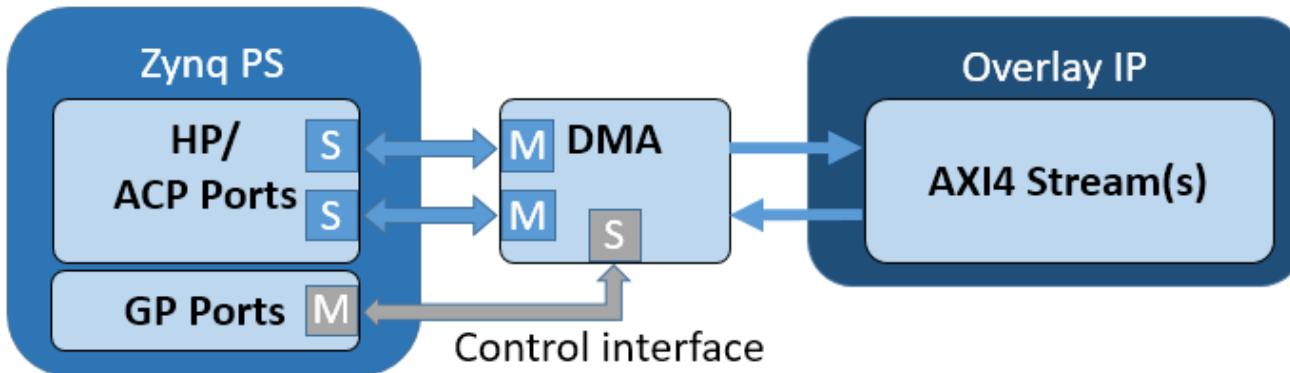
PYNQ Libraries

PYNQ provides a Python API for common peripherals and PL control

- Audio/Video
- GPIO devices (buttons, switches, LEDs, etc.)
- Headers and IO pins (e.g. Raspberry Pi header)
- PynqMicroBlaze subsystem
- Low-level PL control e.g. memory-mapped IO, memory allocation, overlay control, etc.

PYNQ Libraries - DMA

- DMA (direct memory access) can be used for high performance burst transfers between PS DRAM and the PL.
- PYNQ supports the AXI central DMA IP with the PYNQ DMA class.



```
overlay = Overlay('example.bit')
dma = overlay.axi_dma
# allocate arrays
input_buffer = allocate(shape=(5,), dtype=np.uint32)
output_buffer = allocate(shape=(5,), dtype=np.uint32)
# write some data to input array
for i in range(5):
    input_buffer[i] = i
```

```
# actual compute
...
# transfer data using DMA
dma.sendchannel.transfer(input_buffer)
dma.recvchannel.transfer(output_buffer)
dma.sendchannel.wait()
dma.recvchannel.wait()
```

PYNQ Libraries – PYNQ MicroBlaze Subsystem

The PYNQ MicroBlaze subsystem allows loading of programs from Python, controlling executing by triggering the processor reset signal, reading and writing to shared data memory, and managing interrupts received from the subsystem.

- Each PYNQ MicroBlaze subsystem is contained within an IO Processor (IOP)
- An IOP defines a set of communication and behavioral controllers that are controlled by Python.
- Supported IOPs: Arduino, Grove, Pmod, Raspberry Pi

Example:

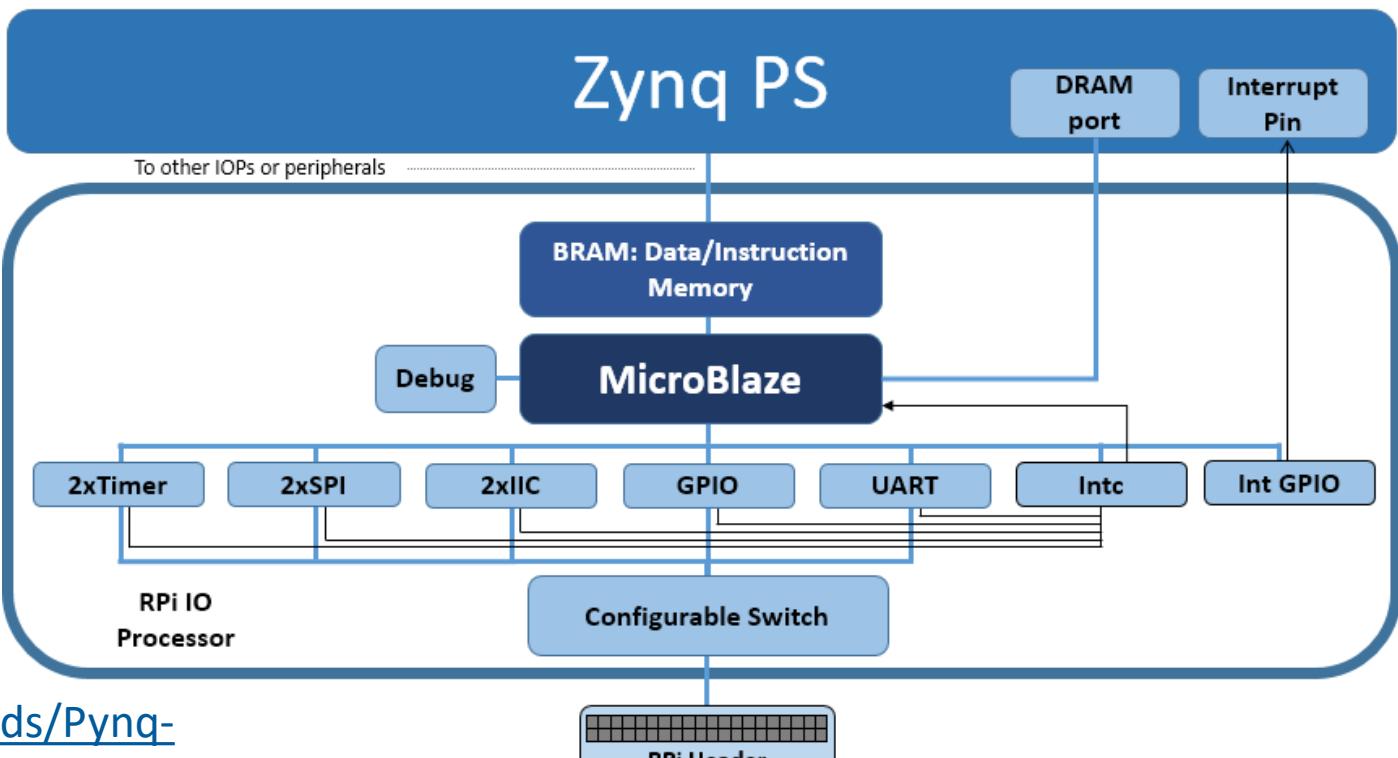
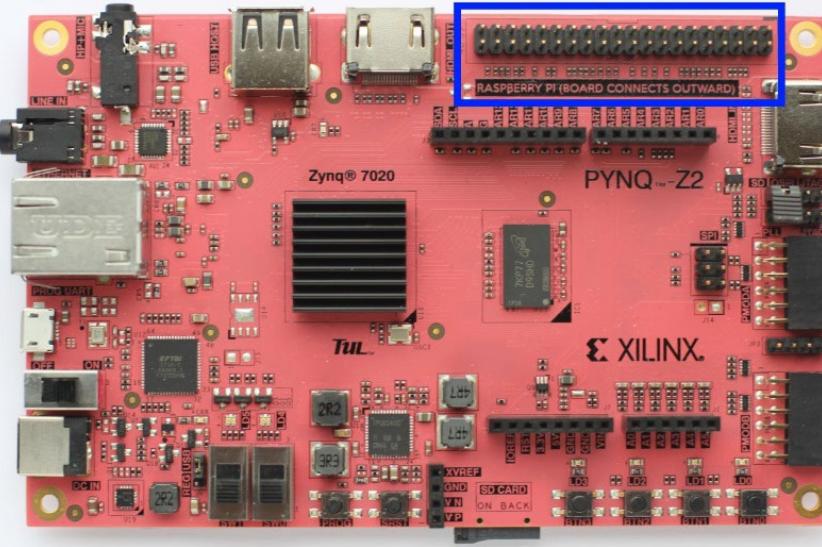
```
from pynq.overlay import BaseOverlay
from pynq.lib import PynqMicroblaze

base = BaseOverlay('base.bit')

mb = PynqMicroblaze(base.iop1.mb_info, # iop1 corresponds to Pmod A connector
                     "/home/xilinx/pynq/lib/pmod/pmod_timer.bin")
mb.reset()
```

PYNQ Libraries – Raspberry Pi Header

- The rpi subpackage is a collection of drivers for controlling peripherals attached to a RPi (Raspberry Pi) interface.
- The RPi PYNQ MicroBlaze is available to control the RPi interface
- RPi PYNQ MicroBlaze has a PYNQ MicroBlaze Subsystem, a configurable switch, and the following AXI controllers:
 - 2x AXI I2C
 - 2x AXI SPI
 - 1x AXI GPIO
 - 2x AXI Timer
 - 1x AXI UART
 - AXI Interrupt controller
 - Interrupt GPIO
 - Configurable Switch



Example: Reading Values from Touch Keypad:

https://github.com/Xilinx/PYNQ/blob/master/boards/Pynq-Z2/base/notebooks/rpi/rpi_touchpad.ipynb

PYNQ on XRT Platforms

Besides Xilinx Zynq platforms, PYNQ also support XRT-based platforms such as Amazon's AWS F1 and Alveo for cloud and on-premise deployment.

(XRT: Xilinx Runtime Library)

Running Accelerators

Start the kernel synchronously:

```
ol.my_kernel.call(input_buf, output_buf)
```

The call function has the same function signature as the top function in original HLS source code.

Alternatively, start the kernel in a non-blocking way:

```
handle = ol.my_kernel.start(input_buf, output_buf)  
handle.wait()
```

Freeing designs (overlays):

```
ol.free()
```

Efficient Scheduling of Multiple Kernels

start and **call** have an optional keyword parameter **waitfor** that can be used to create a dependency graph which is executed in the hardware.

```
handle = ol.vadd_1.start(input1, input2, output)  
ol.vadd_1.call(input3, output, output, waitfor=(handle,))
```

Multiple FPGA Cards

PYNQ supports multiple accelerator cards in one server. It provides a **Device** class to designate which card should be used for given operations.

```
> for i in range(len(pynq.Device.devices)):  
>   print("{} {}".format(i, pynq.Device.devices[i].name))  
0) xilinx_u200_xdma_201830_2  
1) xilinx_u250_xdma_201830_2  
2) xilinx_u250_xdma_201830_2  
3) xilinx_u250_xdma_201830_2
```

Summary

Embedded FPGAs

- Embedded FPGA platforms are typically built with SoC (System-on-Chip)
- SoC consists of PL (programmable logic, FPGA) and PS (processing system, CPU)
- There are multiple connection interfaces between PL, PS, memory, and peripherals
- AXI interfaces are the common interconnection interface in Xilinx SoCs
- Embedded FPGAs have very limited resources and require careful design considerations

PYNQ Environment

- PYNQ is a set of open-source Python-based APIs and libraries that simplifies (host) programming of Xilinx FPGAs and SoCs
- PYNQ supports both low-power Zynq SoC platforms, as well as high-performance XRT-based FPGA platforms, e.g., Alveo boards and AWS-F1 FPGA
- The PYNQ overlay APIs and libraries allow people to use and interact with programmable/configurable FPGA designs in the form of “hardware libraries”

References

- Xilinx Zynq-7000 website: <https://www.xilinx.com/products/silicon-devices/soc/zynq-7000.html>
- PYNQ-Z2 board: <https://www.tul.com.tw/ProductsPYNQ-Z2.html>
- Ultra96-v2 board: <https://www.avnet.com/wps/portal/us/products/new-product-introductions/npi/aes-ultra96-v2/>
- Kria KR260 board: <https://www.xilinx.com/products/som/kria/kr260-robotics-starter-kit.html>
- Xilinx Vivado Design Suite – AXI Reference Guide:
https://www.xilinx.com/support/documentation/ip_documentation/axi_ref_guide/latest/ug1037-vivado-axi-reference-guide.pdf
- PYNQ website: <http://www.pynq.io/>
- PYNQ documents: <https://pynq.readthedocs.io/en/latest/index.html>
- PYNQ GitHub repo: <https://github.com/Xilinx/PYNQ>
- PYNQ tutorial: https://github.com/Xilinx/PYNQ_Workshop
- PYNQ with Alveo examples: <https://github.com/Xilinx/Alveo-PYNQ>
- PYNQ Overlay: https://pynq.readthedocs.io/en/latest/overlay_design_methodology.html

EECS 298: System-on-Chip Design

Sitao Huang, sitaoh@uci.edu

